

Part I

VLFeat

An Open and Portable Library of Computer
Vision Algorithms

Andrea Vedaldi

Visual Geometry Group

Oxford

Brian Fulkerson

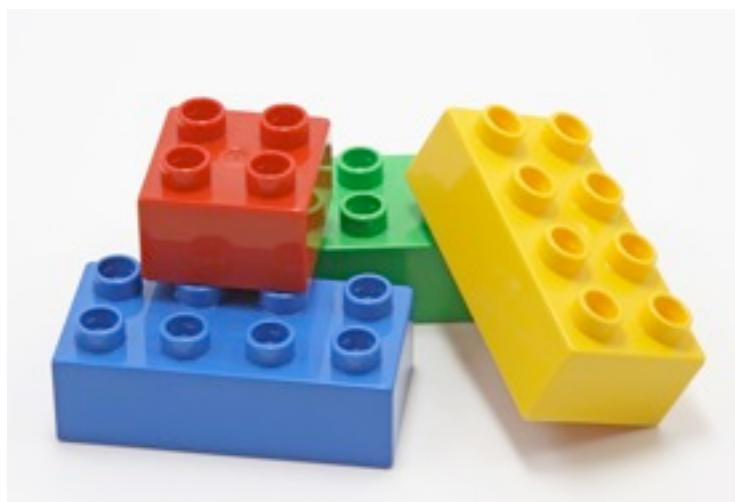
VisionLab

UCLA



Plan

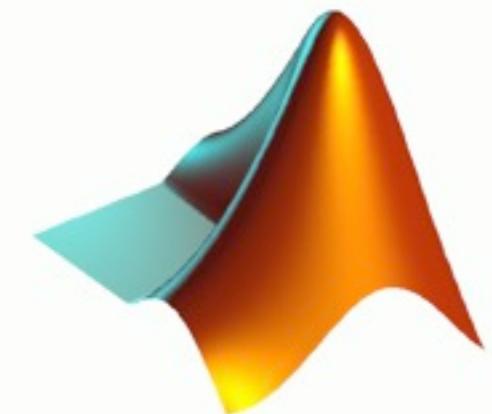
- The VLFeat library
 - SIFT example (`vl_sift`)
- Caltech-101 running example
- Visual descriptors
 - PHOW feature (fast dense SIFT, `vl_phow`)
 - Vector Quantization (Elkan, `vl_kmeans`, `vl_kdtreebuild`,
`vl_kdtreequery`)
 - Spatial histograms (`vl_binsum`, `vl_binsearch`)
- Learning and classification
 - Fast linear SVMs
 - PEGASOS (`vl_pegasos`)
 - Fast non-linear SVMs
 - Homogeneous kernel maps (`vl_homkermapper`)
- Other VLFeat features



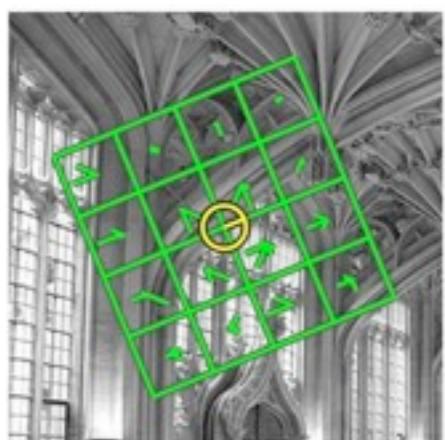
computer vision
building blocks



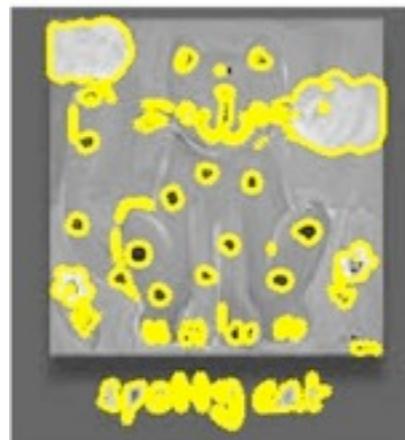
GPL



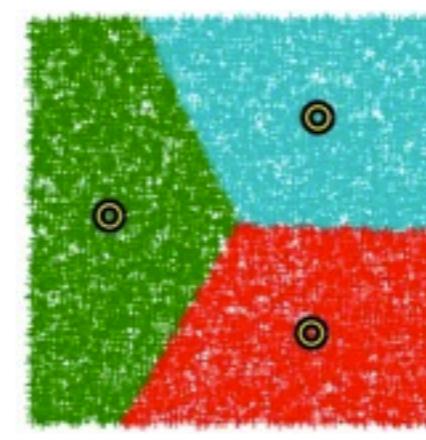
MATLAB



features

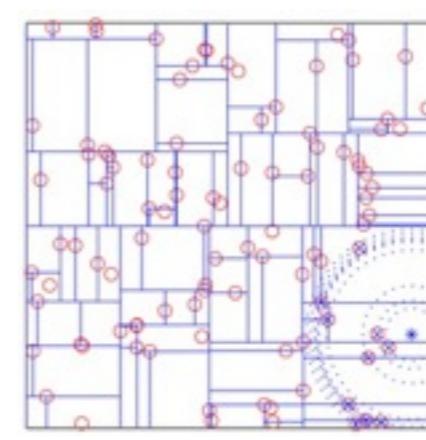


,



clustering

,



matching

...

Running VLFeat



Mac



Linux

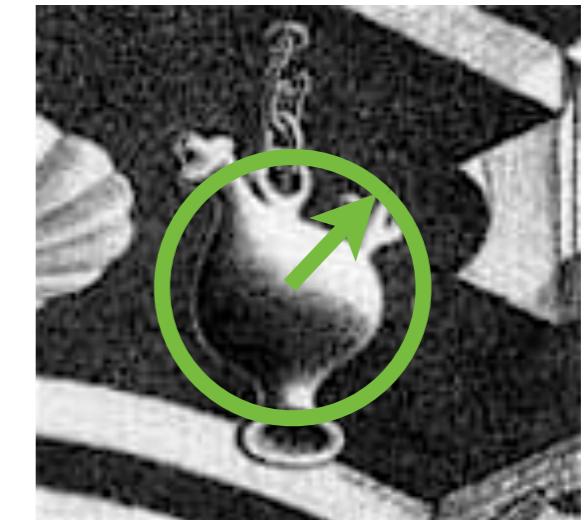


Windows

- **VLFeat.org**
- Quick start
 1. Get **binary package**
`wget http://www.vlfeat.org/download/vlfeat-0.9.9-bin.tar.gz`
 2. Unpack
`tar xzf vlfeat-0.9.9-bin.tar.gz`
 3. Setup MATLAB
`run('VLROOT/toolbox/vl_setup')`

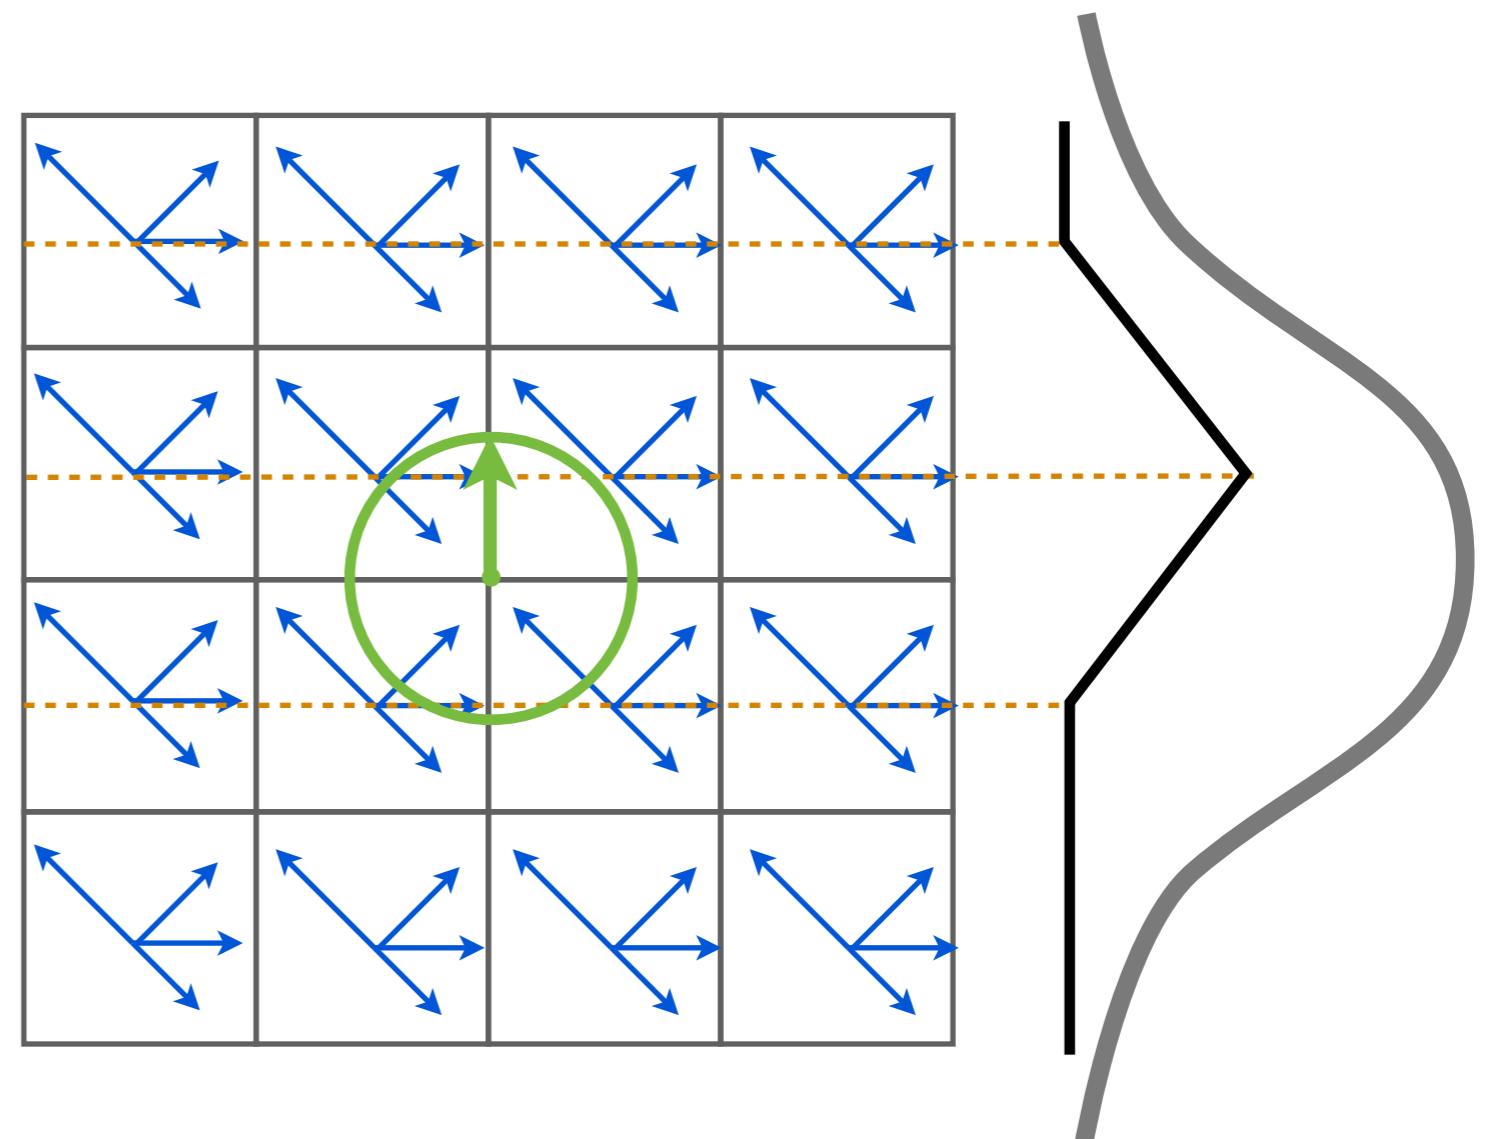
Demo: SIFT features

- **SIFT keypoint [Lowe 2004]**
 - extrema of the DoG scale space
 - blob-like image structure
 - oriented



- **SIFT descriptor**

- 4 x 4 spatial histogram of gradient orientations
- linear interpolation
- Gaussian weighting



Demo: Computing SIFT features

- Output-equivalent to D. Lowe's implementation
- Example

1. load an image

```
imPath = fullfile(vl_root, 'data', 'a.jpg') ;
im = imread(imPath) ;
```

2. convert into float array

```
im = im2single(rgb2gray(im)) ;
```

3. run SIFT

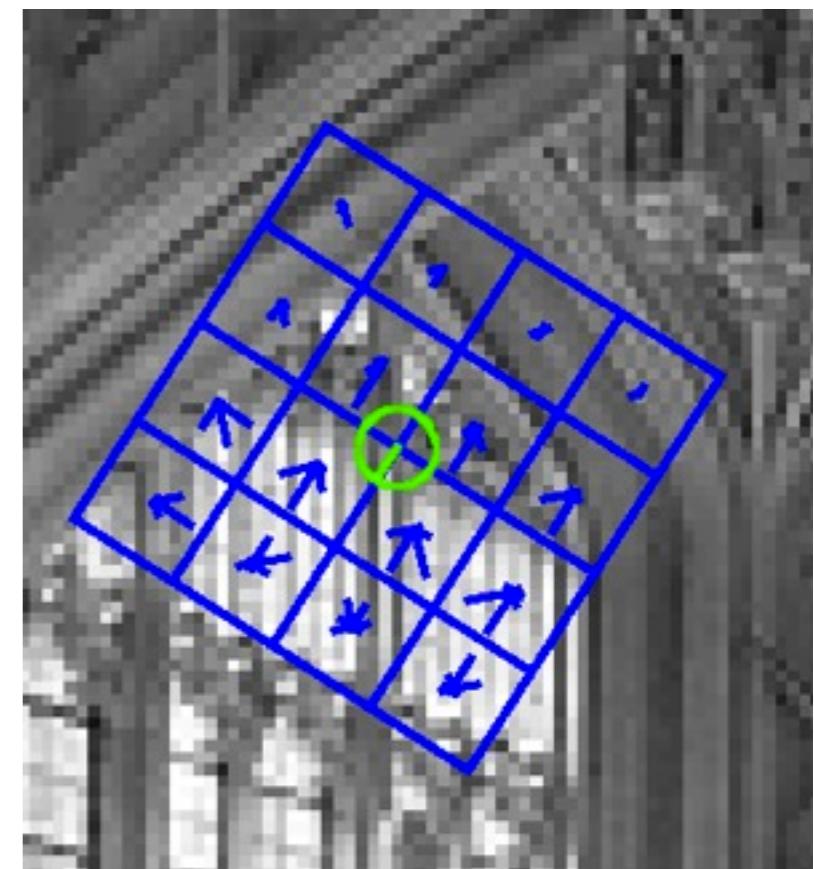
```
[frames, descrs] = vl_sift(im) ;
```

4. visualize keypoints

```
imagesc(im) ; colormap gray; hold on ;
vl_plotframe(frames) ;
```

5. visualize descriptors

```
vl_plotsiftdescriptor(...  
descrs(:,432), frames(:,432)) ;
```

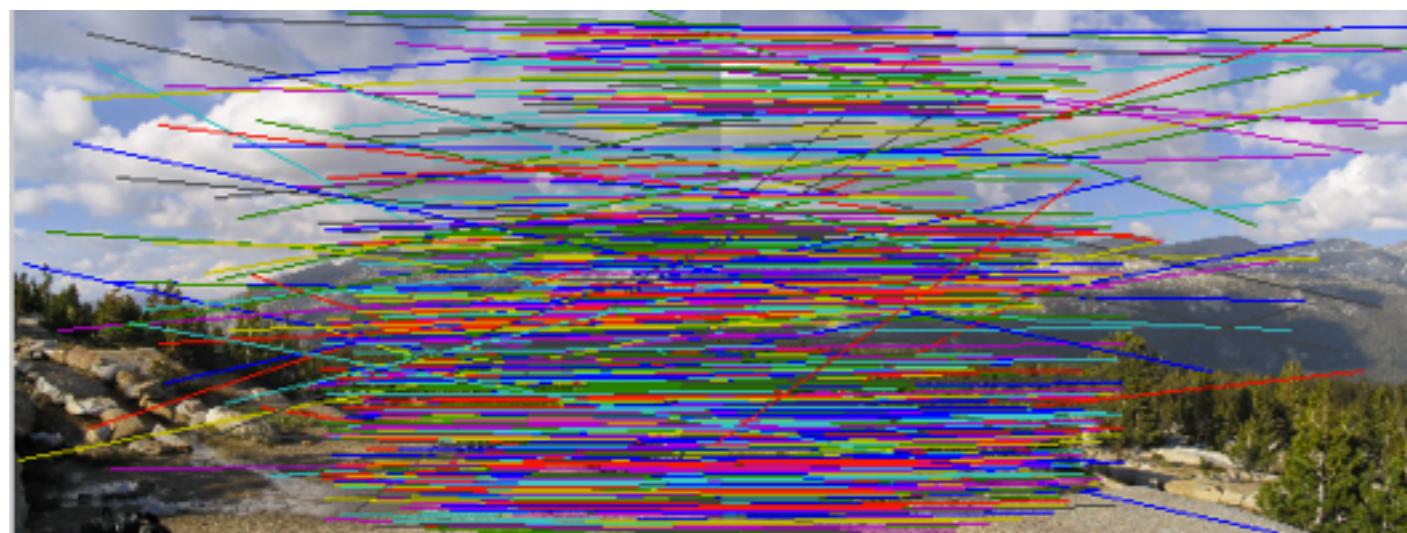




- Matching code

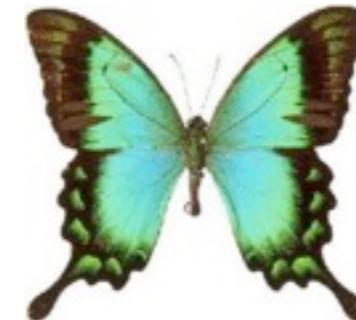
```
[f1,d1] = vl_sift(im2single(rgb2gray(im1))) ;  
[f2,d2] = vl_sift(im2single(rgb2gray(im2))) ;  
[matches, scores] = vl_ubcmatch(d1,d2) ;
```

684 tentative matches

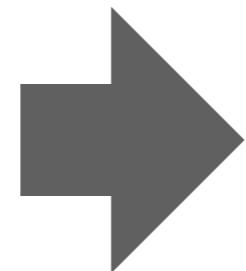


Running example: Caltech-101

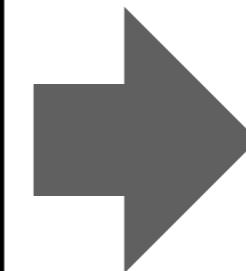
102 semantic labels



[Fei-Fei et al. 2003]



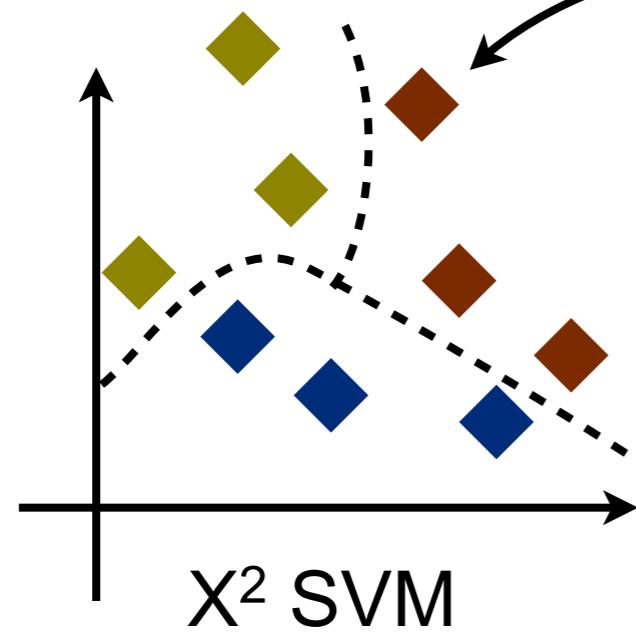
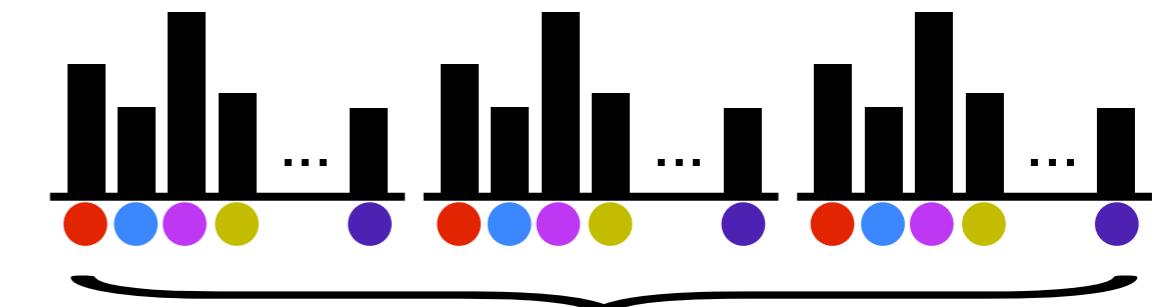
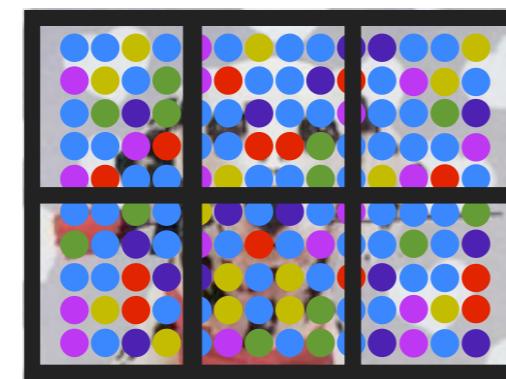
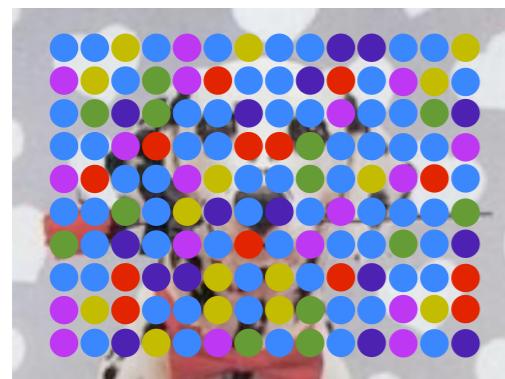
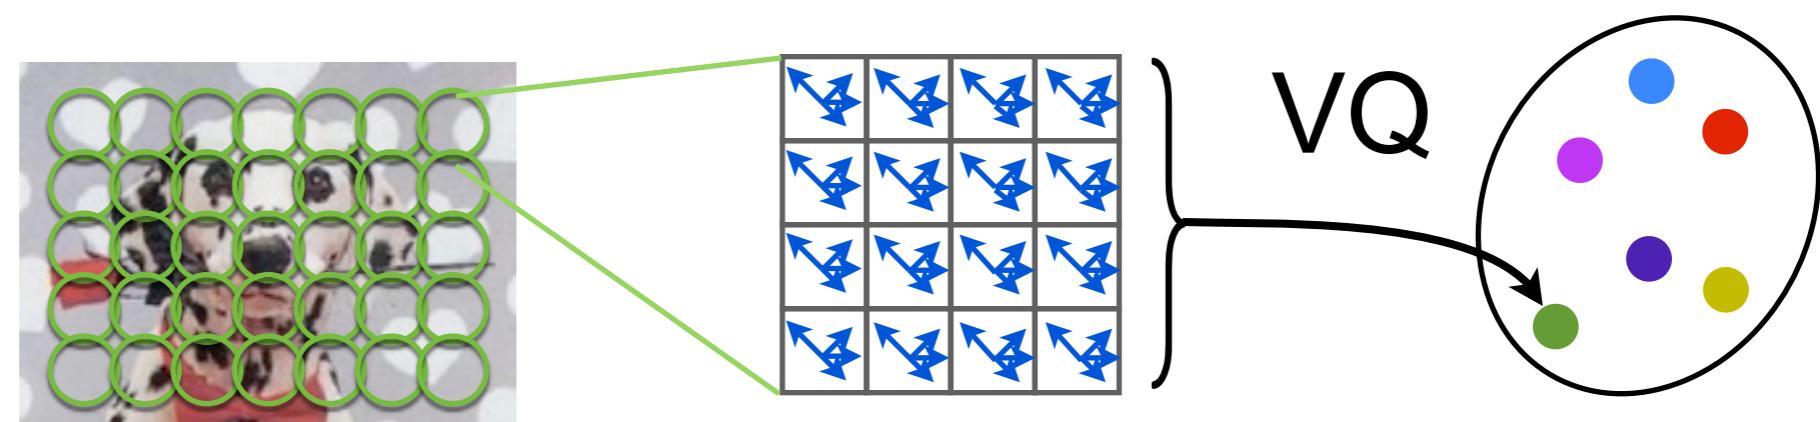
computer
vision



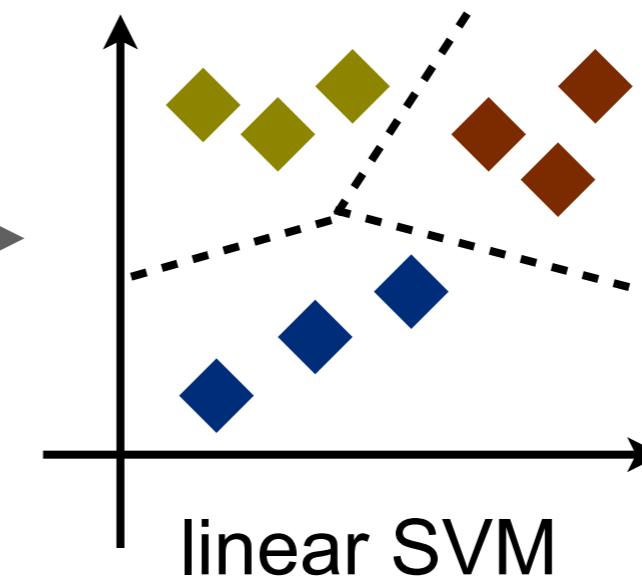
Dalmatian

Running example: System

13



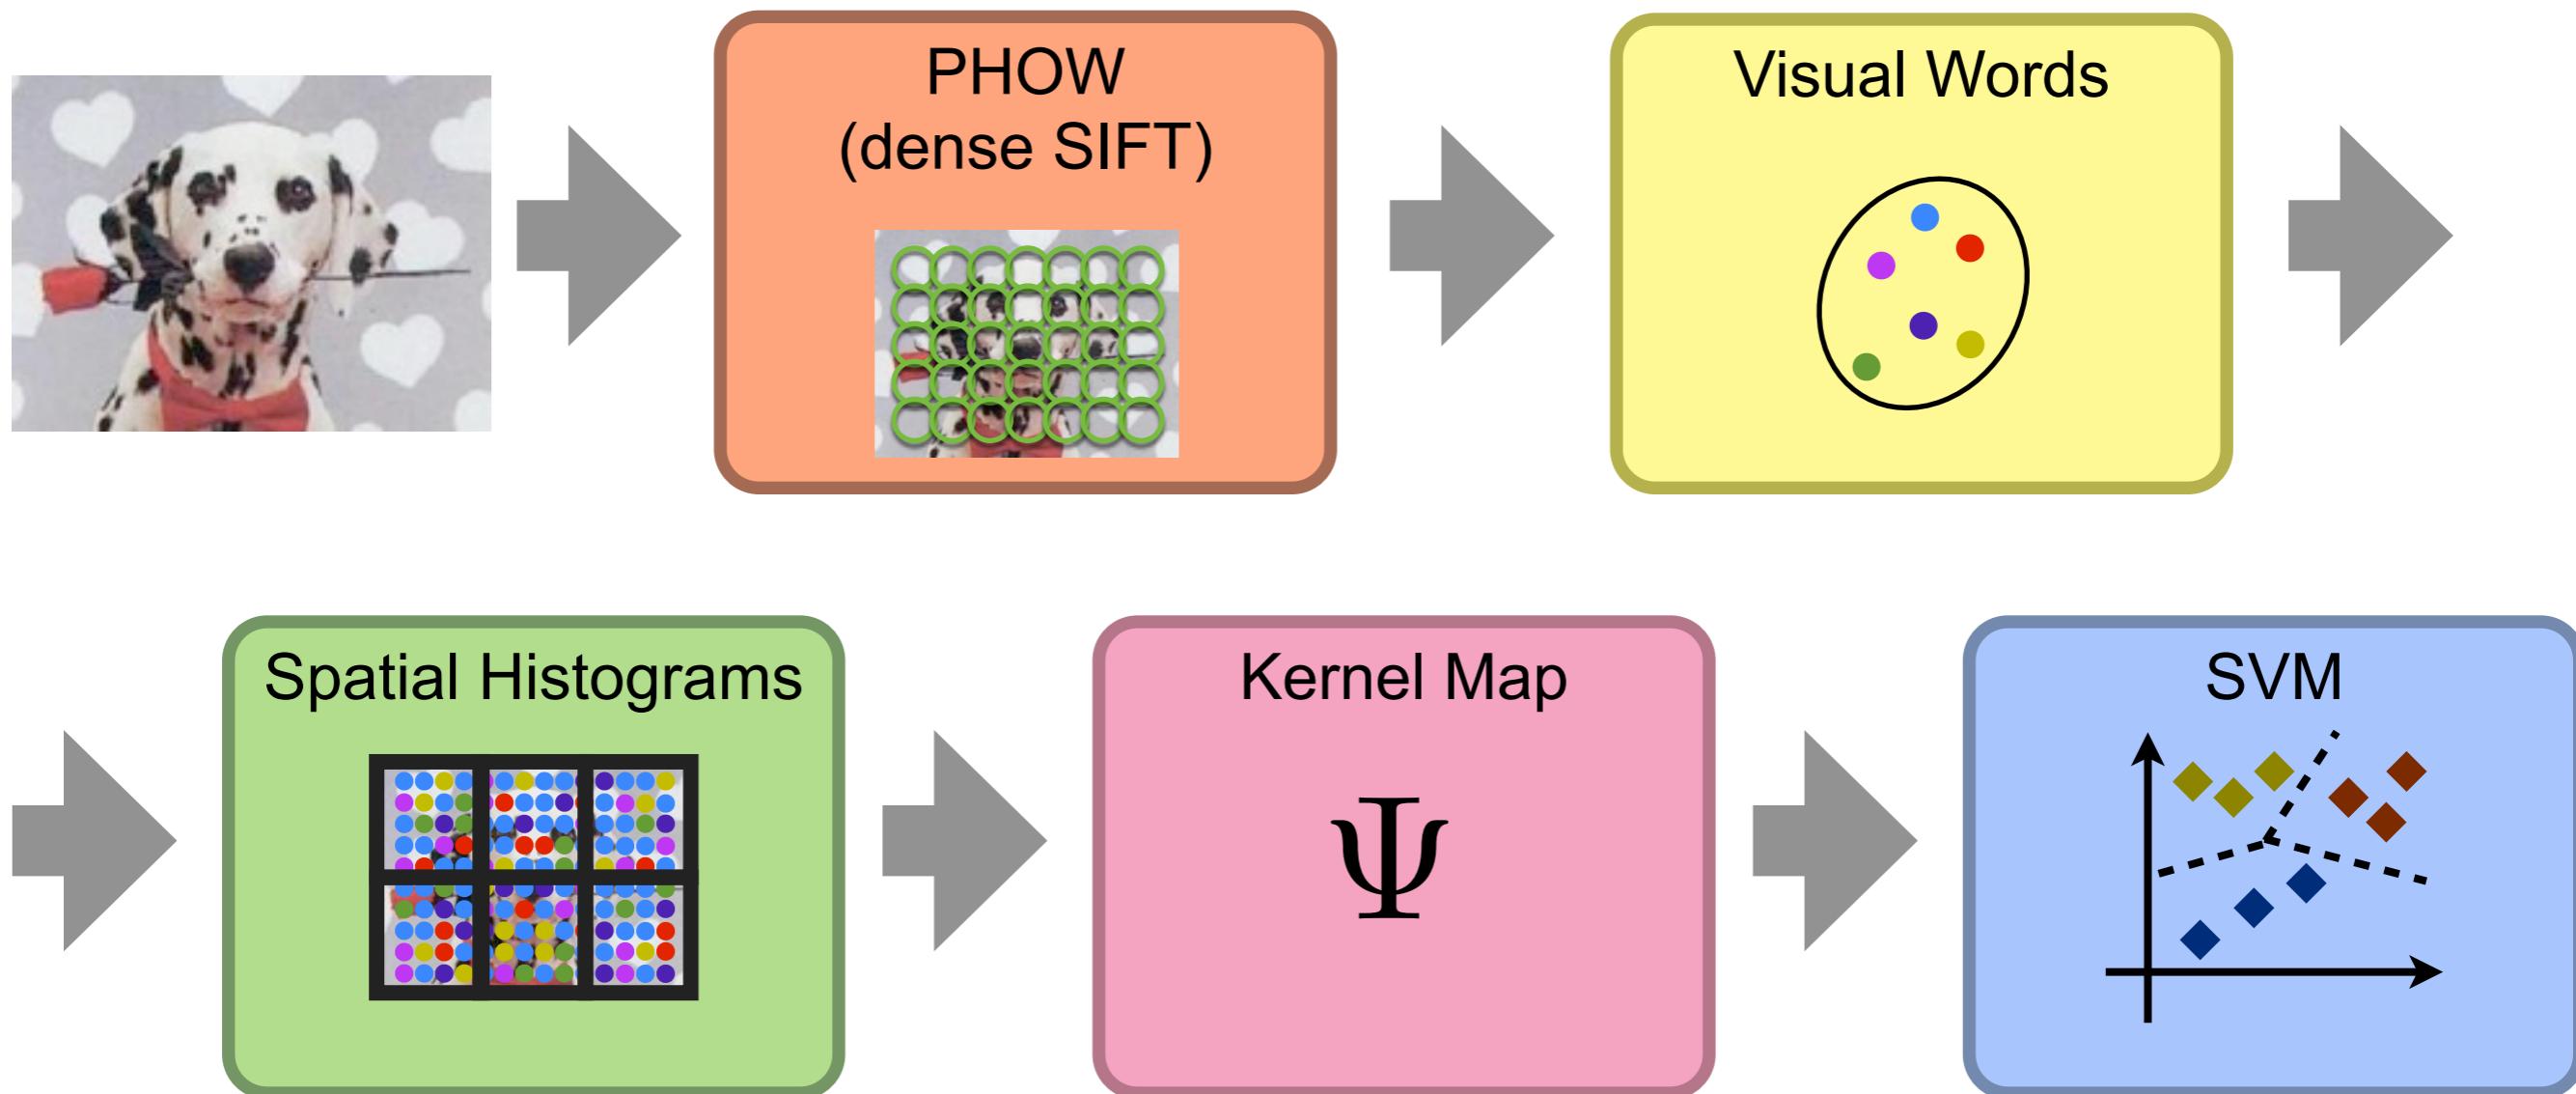
Ψ



linear SVM

Dalmatian

Running example: Components



- Complete source code available
 - `phow_caltech101.m`
 - <http://www.vlfeat.org/applications/apps.html#apps.caltech-101>

Running example: Complete code

- single file
 - ~ 240 lines
 - all inclusive

Running example: Complete code

```

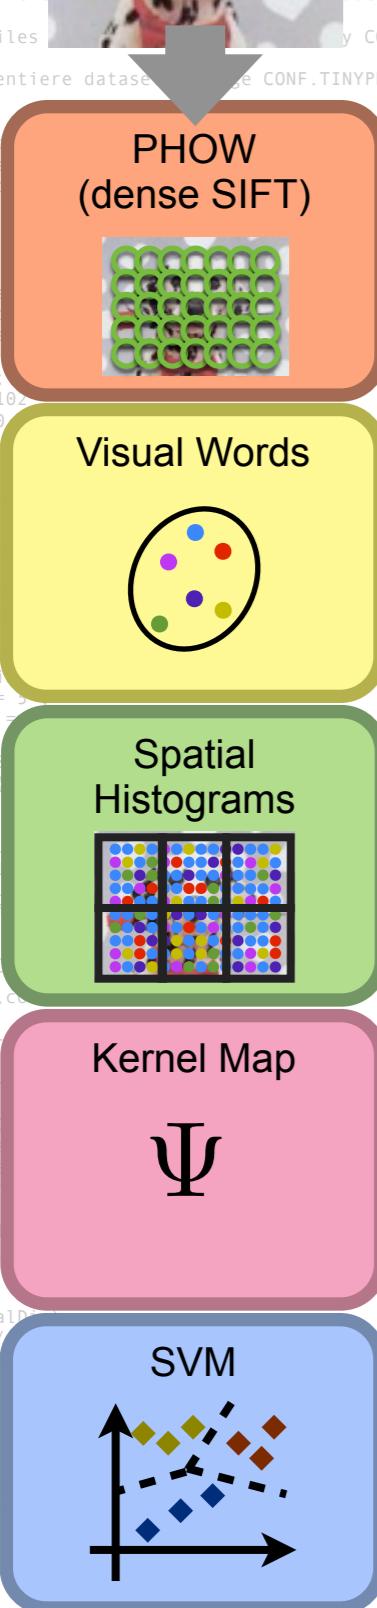
function phow_caltech101
% PHOW_CALTECH101 Image
% A quick test should
% from MATLAB (provide
%
% The program automatically
% interred and decomp
% 'data/caltech-101'.
% instance to point to
%
% Intermediate files
%
% To run on the entire dataset, set CONF.TINYPROBLEM to
% FALSE.
%
% To run with a smaller
% CONF.NUMTRAIN and
% % used, which should
% % remarkable for
%
% AUTORIGHTS
%
conf.calDir = 'data';
conf.dataDir = 'data';
conf.autoDownloadData = true;
conf.numTrain = 15;
conf.numTest = 15;
conf.numClasses = 102;
conf.numWords = 800;
conf.numSpatialX = 10;
conf.numSpatialY = 10;
conf.svm.C = 10;
conf.svm.solver = 'liblinear';
conf.phowOpts = {};
conf.clobber = false;
conf.tinyProblem = false;
conf.prefix = 'base';
conf.randSeed = 1;
if conf.tinyProblem
    conf.prefix = 'tiny';
    conf.numClasses = 5;
    conf.numSpatialX = 5;
    conf.numSpatialY = 5;
    conf.numWords = 3;
    conf.phowOpts = {};
end
conf.vocabPath = fullfile(conf.calDir, 'vocab.mat');
conf.histPath = fullfile(conf.calDir, 'hists.mat');
conf.modelPath = fullfile(conf.calDir, 'model.mat');
conf.resultPath = fullfile(conf.calDir, 'result.mat');

randn('state',conf.randSeed);
rand('state',conf.randSeed);
vl_twister('state',conf.randSeed);

% -----
% -----
% -----
if ~exist(conf.calDir, 'dir')
    (~exist(fullfile(conf.dataDir, 'Caltech-101')) && ...
    ~exist(fullfile(conf.dataDir, 'airplanes'))) ...
    error('... Caltech-101 or airplanes datasets are required');
    ['Caltech-101', 'airplanes'];
    'Set conf.calDir to the directory containing the Caltech-101 ...
    data.');
end
vl_xmkdir(conf.calDir);
calUrl = ['http://www.vlfeat.org/edu/datasets/ ...'];
printf('Download datasets from %s. This will take a ...
while.', conf.calDir);
untar(calUrl, conf.calDir);
end

if ~exist(fullfile(conf.calDir, 'classes'));
    conf.calDir = fullfile(conf.dataDir, 'Caltech-101');
    classes = dir(conf.calDir);
    classes = classes([classes.isdir]);
end

```



```

    classes = {classes(3:conf.numClasses+2-1).name} ;
    images = {} ;
    imageClass = {} ;
    for ci = 1:length(classes)
        ims = dir(fullfile(conf.calDir, classes{ci}, '*.jpg'))' ;
        ims = vl_colsubset(ims, conf.numTrain + conf.numTest) ;
        ims = cellfun(@(x)fullfile(classes{ci},x), ...
        {ims.name}, 'UniformOutput',false) ;
        images = {images{:}, ims{:}} ;
        imageClass{end+1} = ci * ones(1,length(ims)) ;
    end
    selTrain = find(mod(0:length(images)-1, conf.numTrain+conf.numTest) <
    conf.numTrain) ;
    selTest = setdiff(1:length(images), selTrain) ;
    imageClass = cat(2, imageClass{:}) ;

    % -----
    % -----
    % ----- Train vocabulary
    if ~exist(conf.vocabPath) || conf.clobber
        % Get some PHOW descriptors to train the dictionary
        selTrainFeats = vl_colsubset(selTrain, 30) ;
        descrs = {} ;
        for ii = 1:length(selTrainFeats)
            im = imread(fullfile(conf.calDir, images{ii})) ;
            im = imageStandarize(im) ;
            [drop, descrs{ii}] = vl_phow(im, conf.phowOpts{:}) ;
        end
        descrs = vl_colsubset(cat(2, descrs{:}), 10e4) ;
        descrs = single(descrs) ;

        % Quantize the descriptors to get the visual words
        words = vl_kmeans(descrs, conf.numWords, 'verbose', 'algorithm',
        'elkan') ;
        save(conf.vocabPath, 'words') ;
    else
        load(conf.vocabPath) ;
    end

    % -----
    % -----
    % ----- Compute spatial histograms
    if ~exist(conf.histPath) || conf.clobber
        hists = {} ;
        % par
        for ii = 1:length(images)
            fprintf('Processing %s (%.2f %%)\n', images{ii}, 100 * ii / length
            (images)) ;
            im = imread(fullfile(conf.calDir, images{ii})) ;
            im = imageStandarize(im) ;
            [frames, descrs] = vl_phow(im, conf.phowOpts{:}) ;

            % quantize appearance
            [drop, binsa] = min(vl_alldist(words, single(descrs)), [], 1) ;

            % quantize location
            width = size(im, 2) ;
            height = size(im, 1) ;
            binsx = vl_binsearch(linspace(1,width,conf.numSpatialX+1), frames
            (1,:)) ;
            binsy = vl_binsearch(linspace(1,height,conf.numSpatialY+1), frames
            (2,:)) ;

            % combined quantization
            bins = sub2ind([conf.numSpatialY, conf.numSpatialX,
            conf.numWords], ... , binsy,binsx,binsa) ;
            hist = zeros(conf.numSpatialY * conf.numSpatialX * conf.numWords,
            1) ;
            hist = vl_binsum(hist, ones(size(bins)), bins) ;
            hist = single(hist / sum(hist)) ;

            hists{ii} = hist ;
        end
        hists = cat(2, hists{:}) ;
        save(conf.histPath, 'hists') ;
    else
        load(conf.histPath) ;
    end

    % -----
    % -----
    % ----- Compute feature map
    psix = vl_homkermap(hists, 1, .7, 'kchi2') ;

    % -----
    % -----
    % ----- Train SVM
    biasMultiplier = 1 ;
    if ~exist(conf.modelPath) || conf.clobber
        switch conf.svm.solver
            case 'pegasos'
                lambda = 1 / (conf.svm.C * length(selTrain)) ;
                models = [] ;
                for ci = 1:length(classes)
                    perm = randperm(length(selTrain)) ;
                    fprintf('Training model for class %s\n', classes{ci}) ;
                    y = 2 * (imageClass(selTrain) == ci) - 1 ;
                    models(:,ci) = vl_pegasos(psix(:,selTrain(perm)), ...
                    int8(y(perm)), lambda, ...
                    'NumIterations', 20/lambda, ...
                    'BiasMultiplier', biasMultiplier) ;
                end
            case 'liblinear'
                model = train(imageClass(selTrain)', ...
                sparse(double(psix(:,selTrain))), ... ,
                sprintf(' -s 3 -B 1 -c %f', conf.svm.C), ...
                'col') ;
                models = model.w' ;
            end
        save(conf.modelPath, 'models') ;
    else
        load(conf.modelPath) ;
    end

    % -----
    % -----
    % ----- Test SVM and evaluate
    % Estimate the class of the test images
    scores = [] ;
    for ci = 1:length(classes)
        scores(ci, :) = models(1:end-1,ci)' * psix + models(end,ci) * ...
        biasMultiplier ;
    end
    [drop, imageEstClass] = max(scores, [], 1) ;

    % Compute the confusion matrix
    idx = sub2ind([length(classes), length(classes)], ...
    imageClass(selTest), imageEstClass(selTest)) ;
    confus = zeros(length(classes)) ;
    confus = vl_binsum(confus, ones(size(idx)), idx) ;

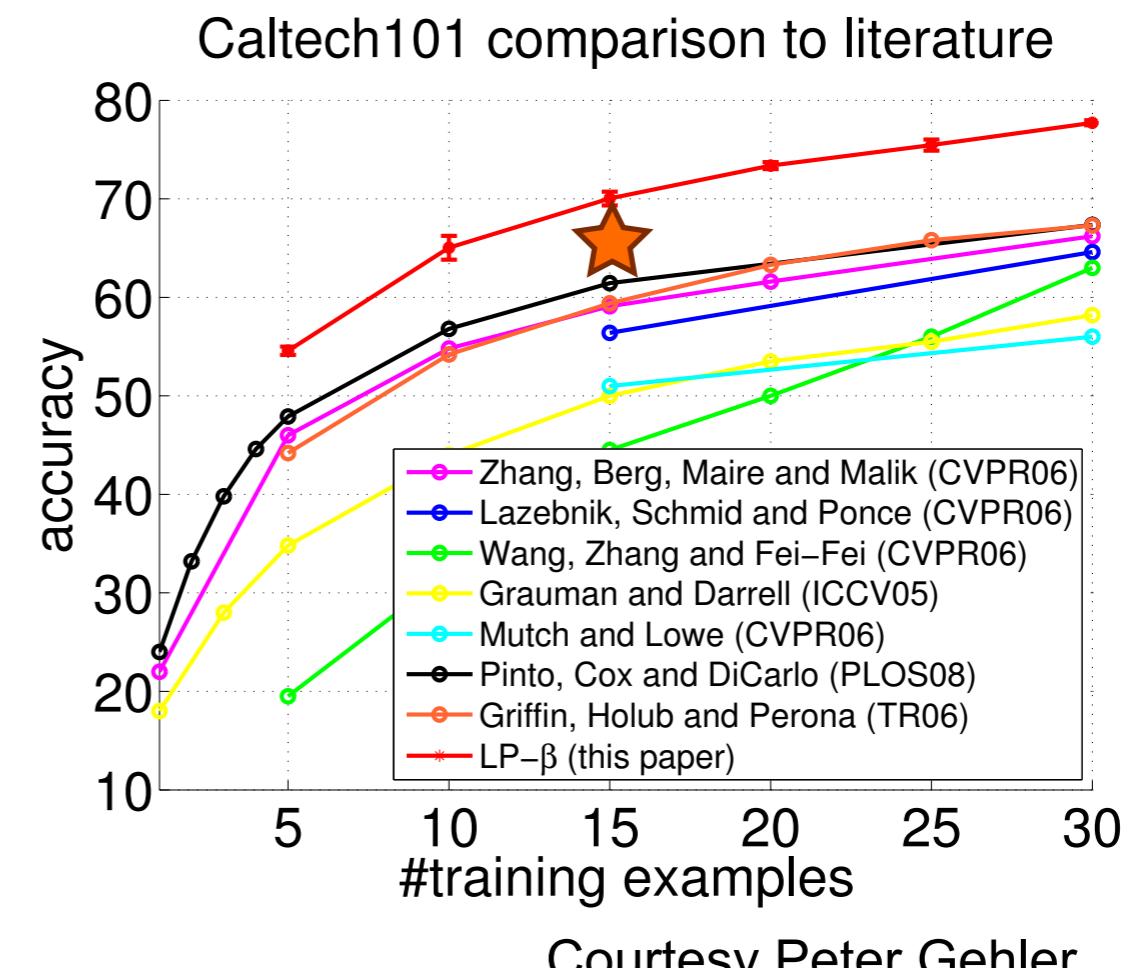
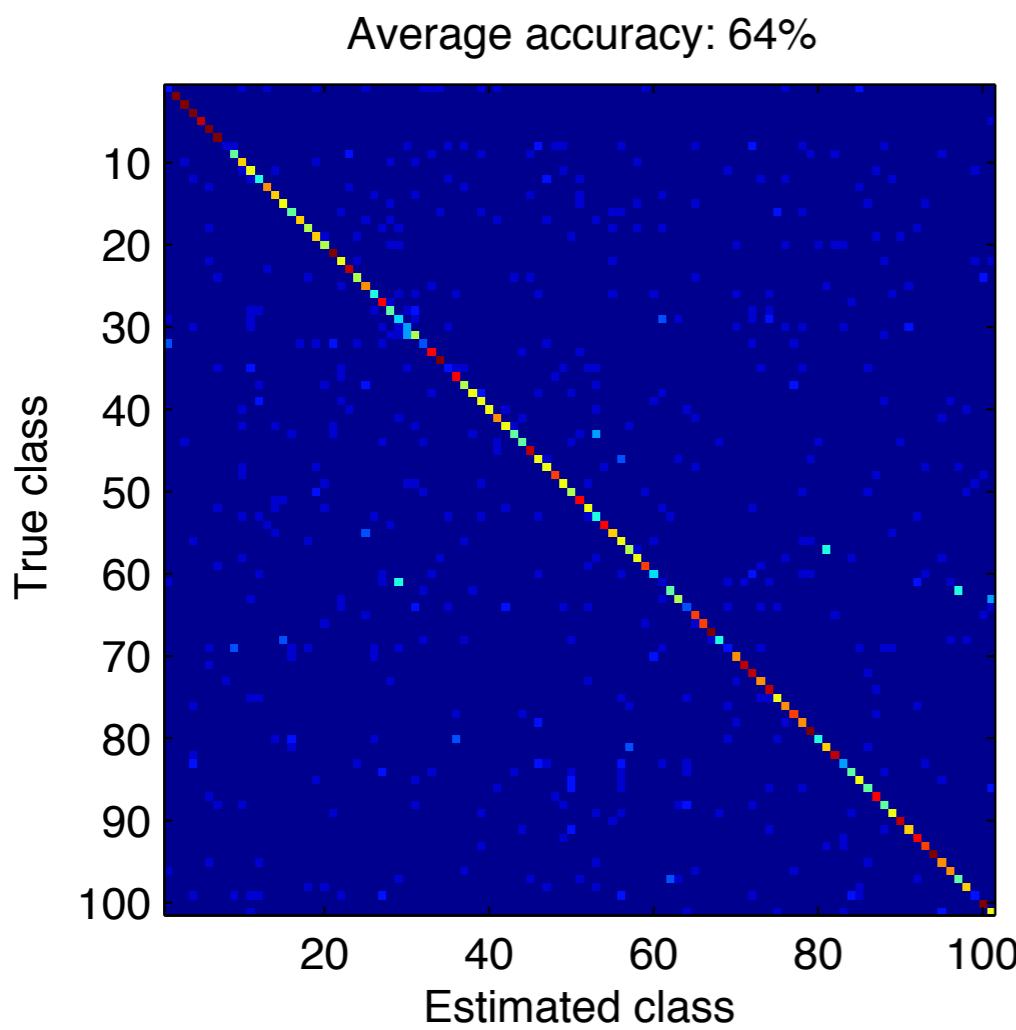
    % Plots
    figure(1) ; clf;
    subplot(1,2,1) ;
    imagesc(scores(:,[selTrain selTest])) ; title('Scores') ;
    set(gca, 'ytick', 1:length(classes), 'yticklabel', classes) ;
    subplot(1,2,2) ;
    imagesc(confus) ;
    title(sprintf('Confusion matrix (%.2f %% accuracy)', ...
    100 * mean(diag(confus)/conf.numTest))) ;
    print('-depsc2', [conf.resultPath '.ps']) ;
    save([conf.resultPath '.mat'], 'confus', 'conf') ;

    % -----
    % -----
    % ----- function im = imageStandarize(im)
    im = im2single(im) ;
    if size(im,1) > 480, im = imresize(im, [480 NaN]) ; end

```

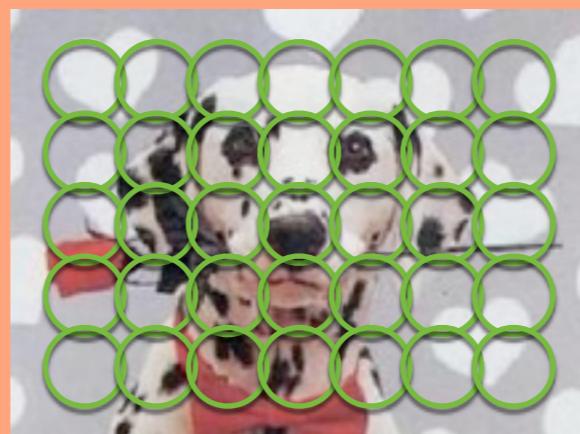
Running example: Results

- Among the best single-feature methods on Caltech-101
 - PHOW (dense SIFT)
 - 15 training - 15 testing images
 - 64-66% average accuracy
- Fast training and testing



PHOW

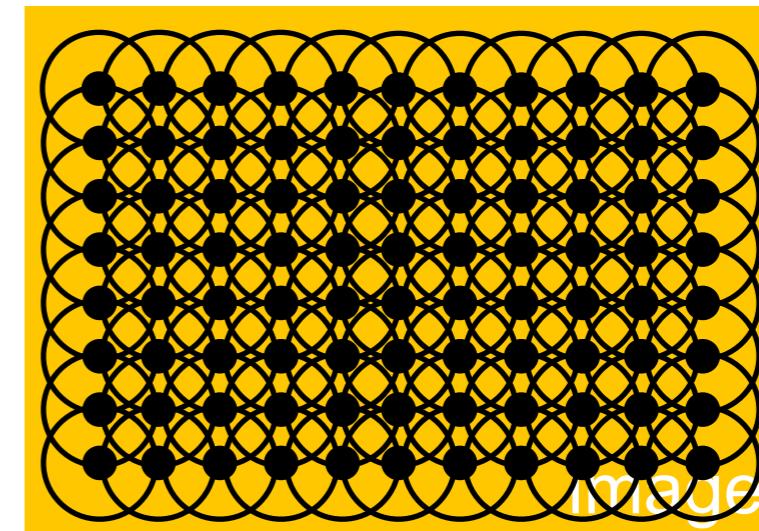
(dense SIFT)



Dense SIFT: PHOW features

- **PHOW features**

- [Bosch et al. 2006], [Bosch et al. 2007]
- dense multiscale SIFT
- uniform spacing (e.g. 5 pixels)
- four scales (e.g. 5, 7, 10, 12 pixels)



- **Direct implementation**

- for each scale
 - create a list of keypoints
 - call **vl_sift**



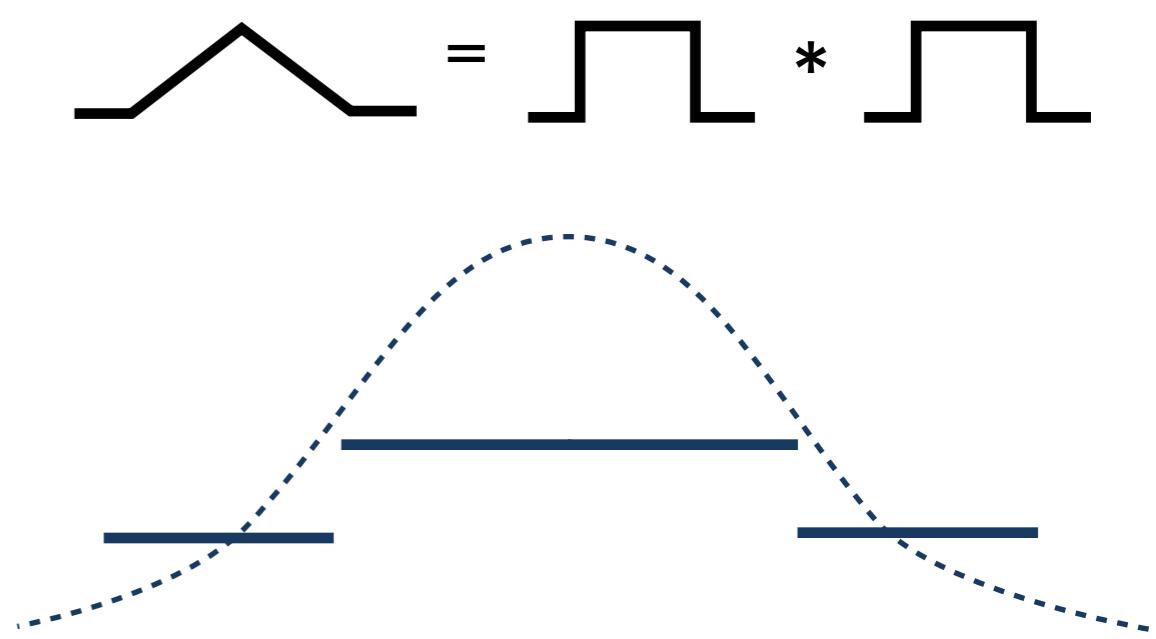
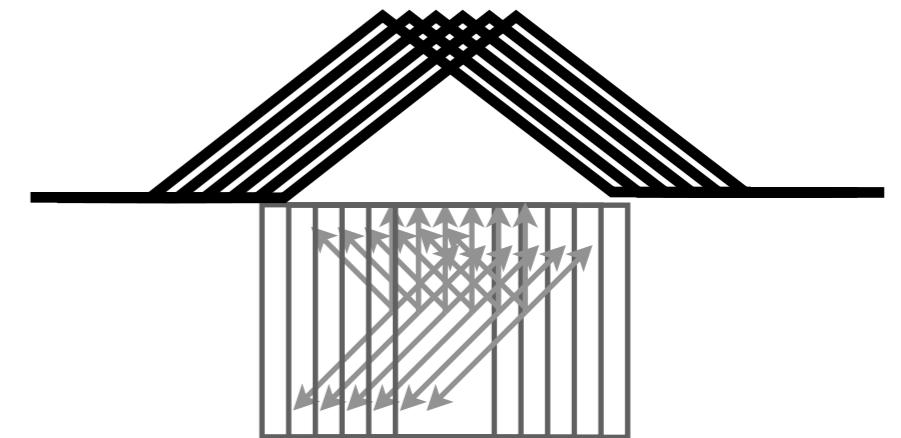
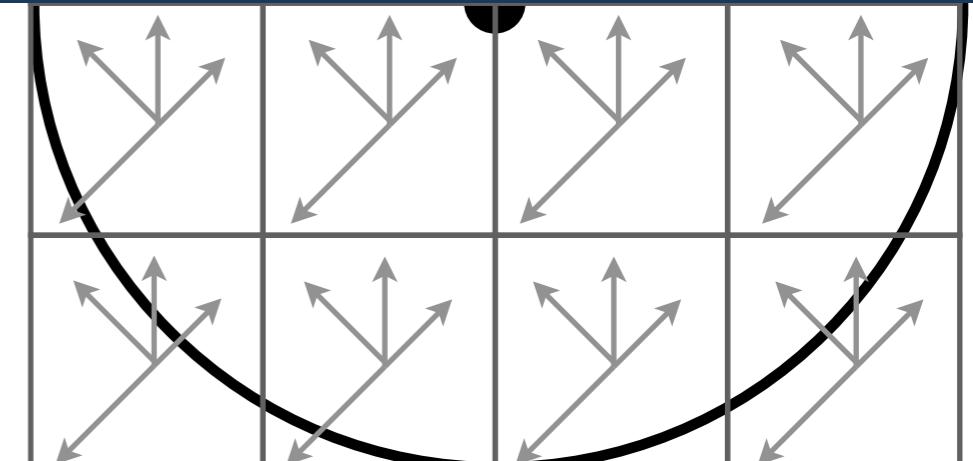
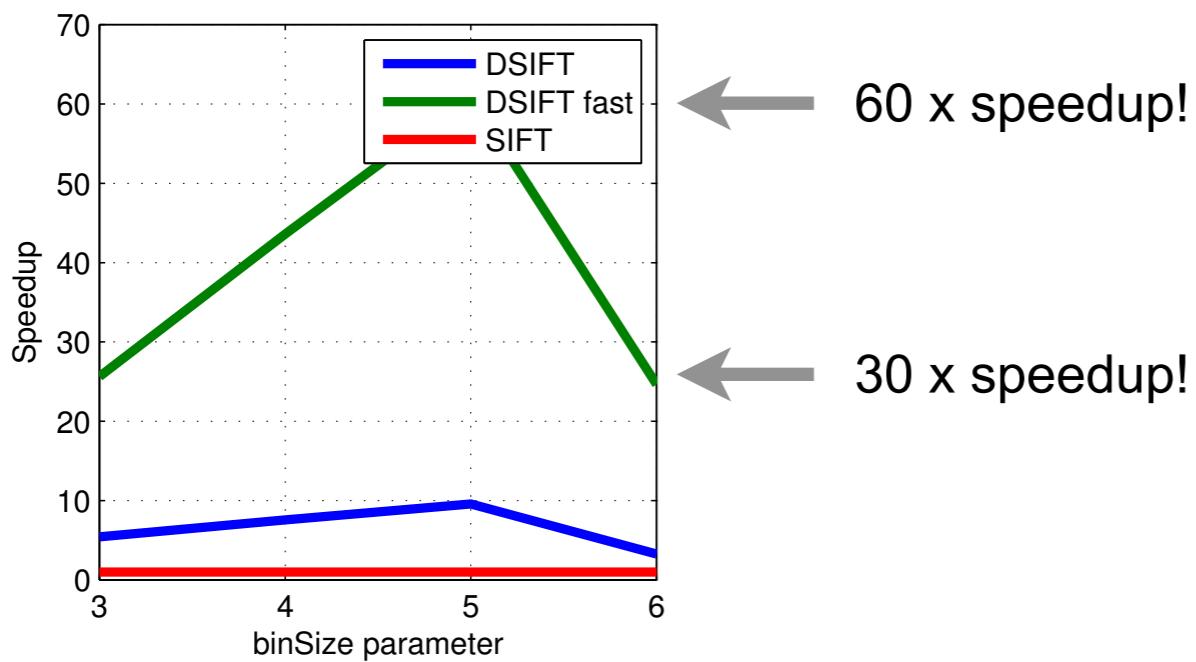
```

step = 5 ;
for s = [5, 7, 10, 12]
  [x, y] = meshgrid(1:step:width, 1:step:height) ;
  frames = [x(:)' ; y(:)'] ;
  frames(3,:) = s / 3 ;
  frames(4,:) = 0 ;
  [frames, descrs] = vl_sift(im, 'Frames', frames) ;
end

```

Accelerating dense SIFT

- Large **speedup** by exploiting
 - same scale and orientation
 - uniform sampling
- **Dense SIFT**
 - linear interpolation = convolution
 - convolution by integral images
 - piecewise approximation of Gaussian window
- Implemented by ***vl_dsift***



Dense SIFT: The fast version

- PHOW = Fast dense SIFT at multiple scales

Use `vl_imsmooth` compute the scale space.

```
for i = 1:4
    ims = vl_imsmooth(im, scales(i) / 3) ;
    [frames{s}, descrs{s}] = ...
        vl_dsift(ims, 'Fast', ...
            'Step', step, 'Size', scales(i)) ;
end
```

- Remark

- for demonstration only!
- in practice use the provided wrapper

```
[frames, descrs] = vl_phow(im) ;
```

Accelerating dense SIFT: related methods

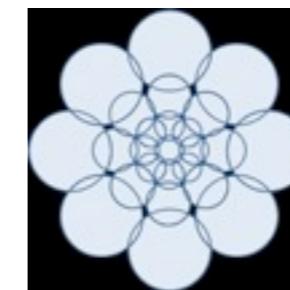
- **SURF**

- [Bay et al. 2008]
- Original closed source
<http://www.vision.ee.ethz.ch/~surf/index.html>
- OpenSURF
<http://www.chrisevansdev.com/computer-vision-opensurf.html>



- **Daisy**

- [Tola et al. 2010]
- <http://cvlab.epfl.ch/~tola/daisy.html>



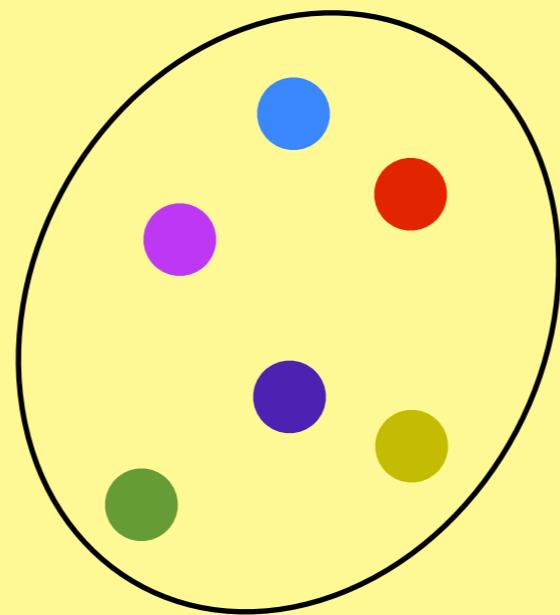
- **GPU-SIFT**

- [S. N. Sinha et al. 2010]
- http://cs.unc.edu/~ssinha/Research/GPU_KLT/

- **SiftGPU**

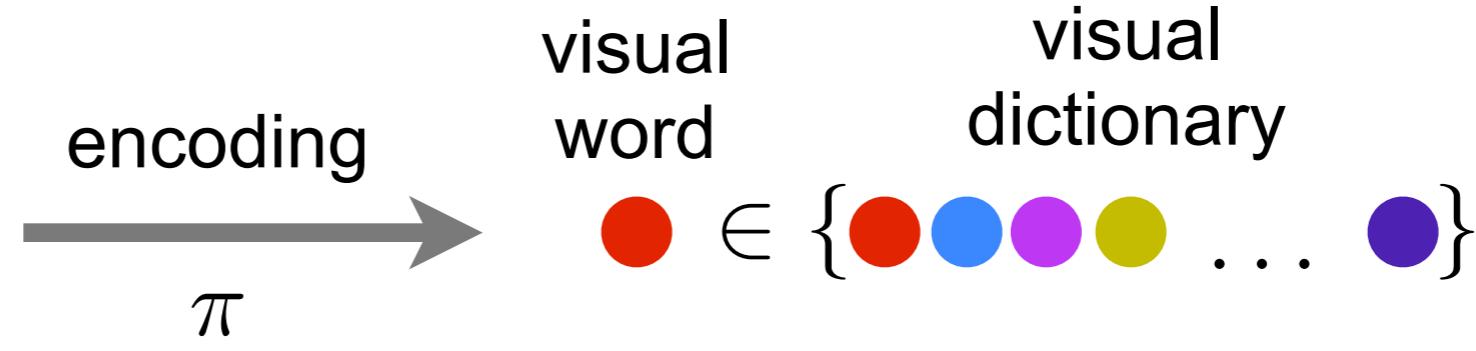
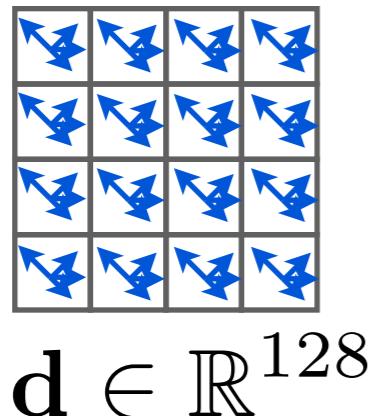
- [C. Wu 2010]
- <http://www.cs.unc.edu/~ccwu/siftgpu/>

Visual Words



Visual words

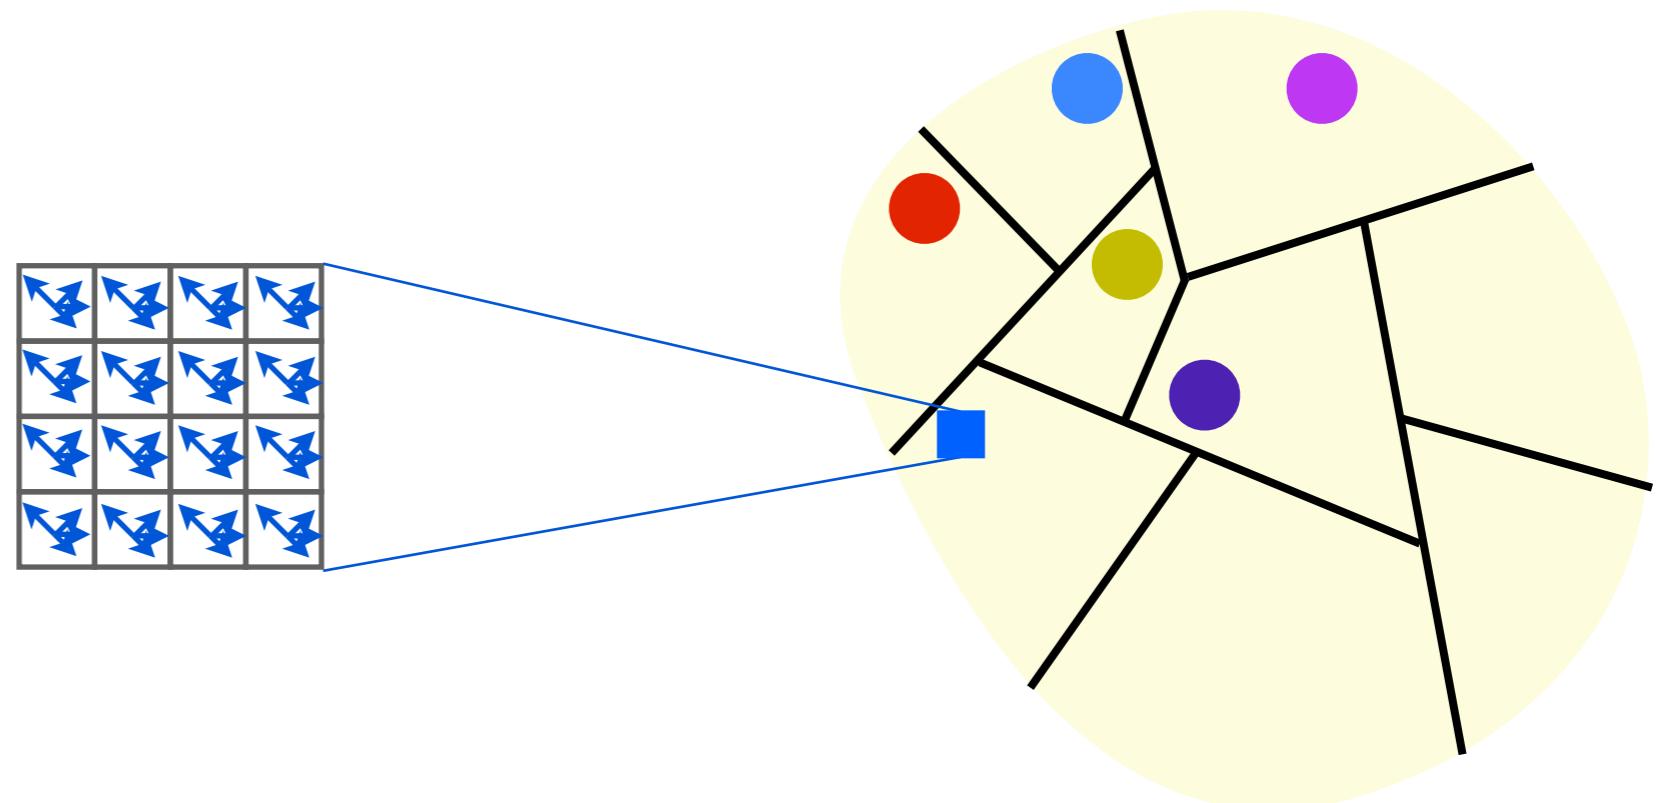
- Visual words



- Encoding = clustering

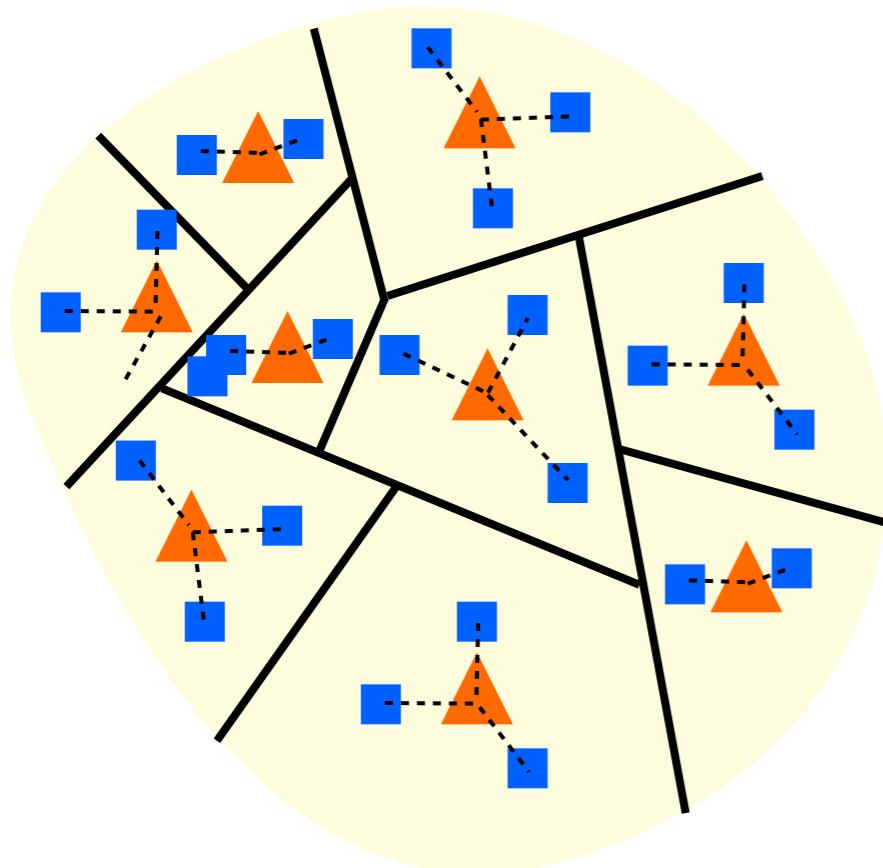
- vector quantization (k-means)
[Lloyd 1982]
- agglomerative clustering
[Leibe et al. 2006]
- affinity propagation
[Frey and Dueck 2007]
- ...

[Sivic and Zisserman 2003]



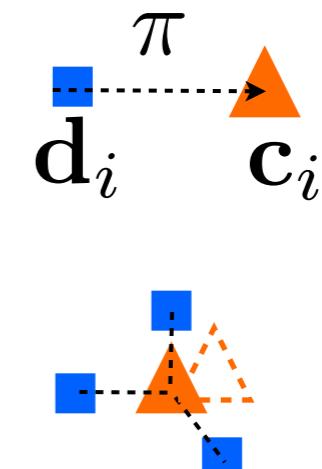
k-means

- Optimize reconstruction cost



$$E(\pi, \mathbf{c}_1, \dots, \mathbf{c}_K) = \sum_{i=1}^M \|\mathbf{d}_i - \mathbf{c}_{\pi(\mathbf{d}_i)}\|^2$$

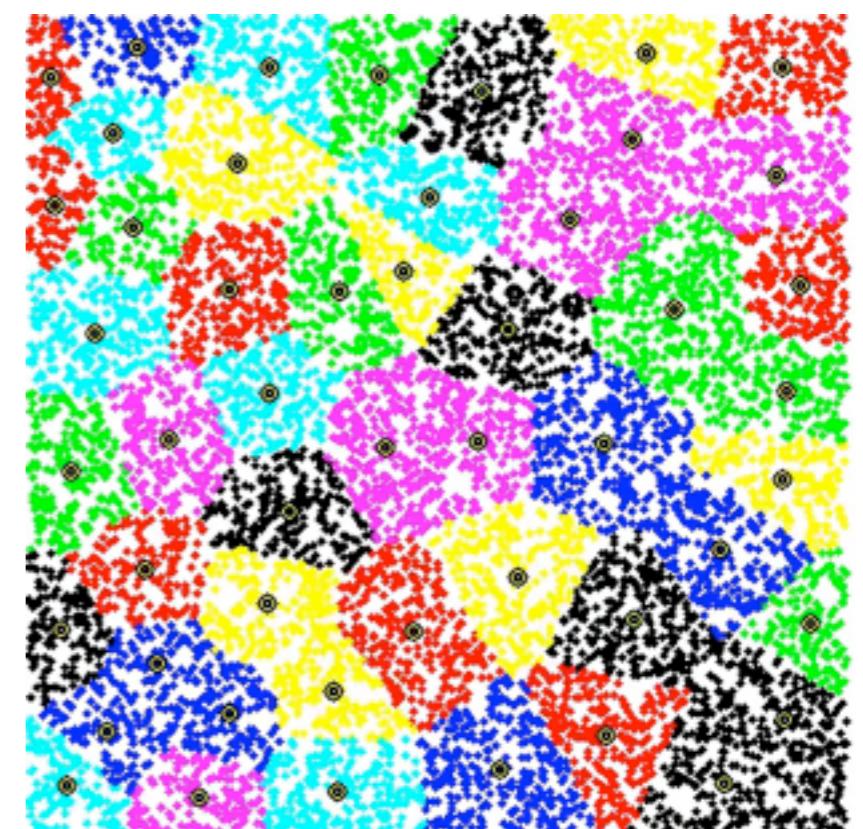
- ① reassign points to closest centers
- ② fit centers to assigned points



- As simple as you expect

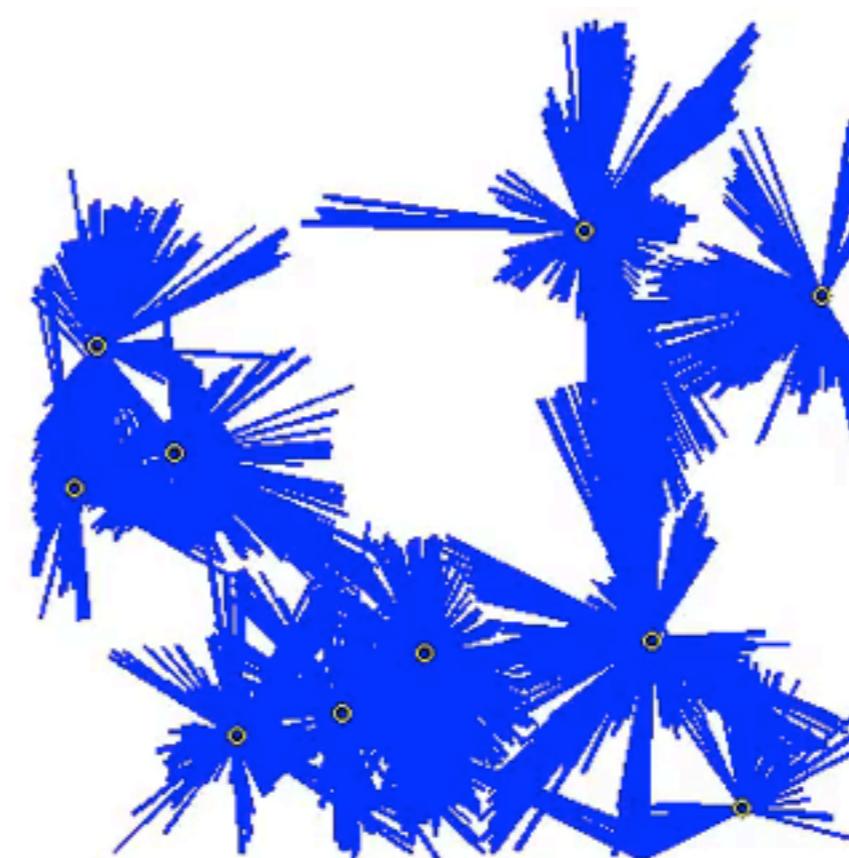
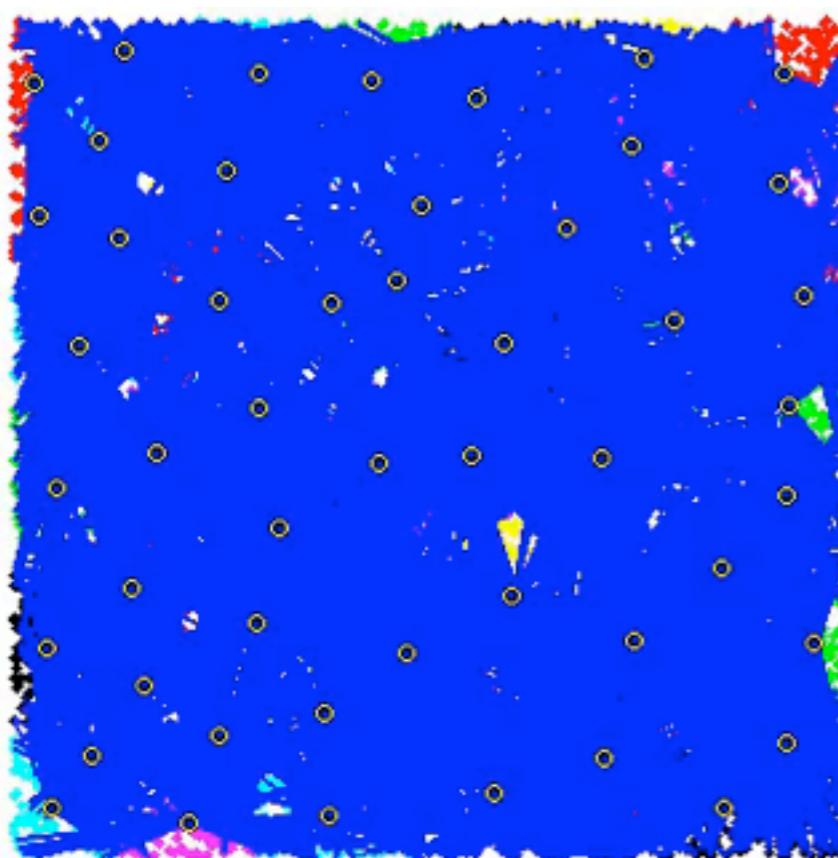
```
centers = vl_kmeans(descrs, K);
```

- Much faster than MATLAB builtin



Faster k-means

- **Bottleneck**
 - assign point to closest center
 - needs to compare each point to each center
- **Speedup**
 - if centers change little, then most assignment don't change
 - can keep track with triangle inequality
 - [Elkan 2003]
 - `centers = vl_kmeans(descrs, K, 'Algorithm', 'Elkan');`

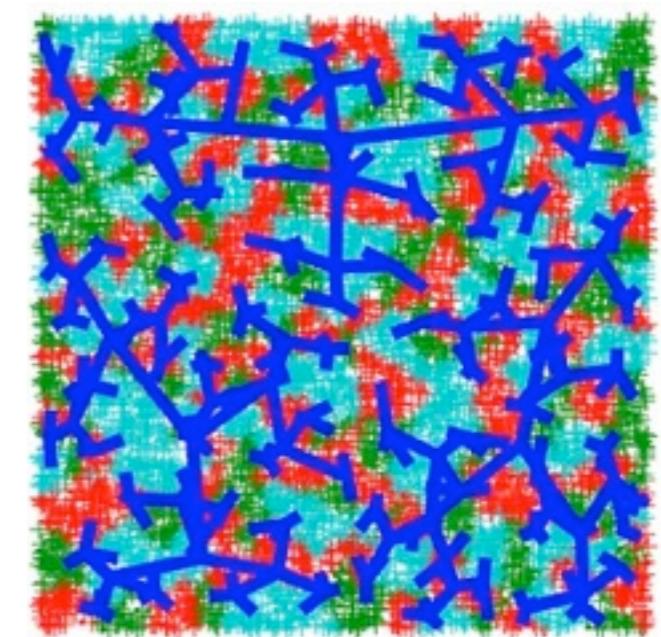
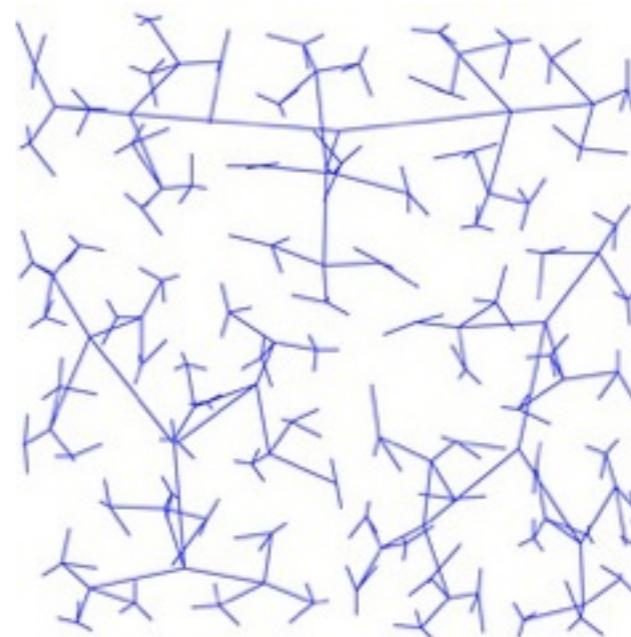


— Vector-to-center
comparison

Speeding-up encoding

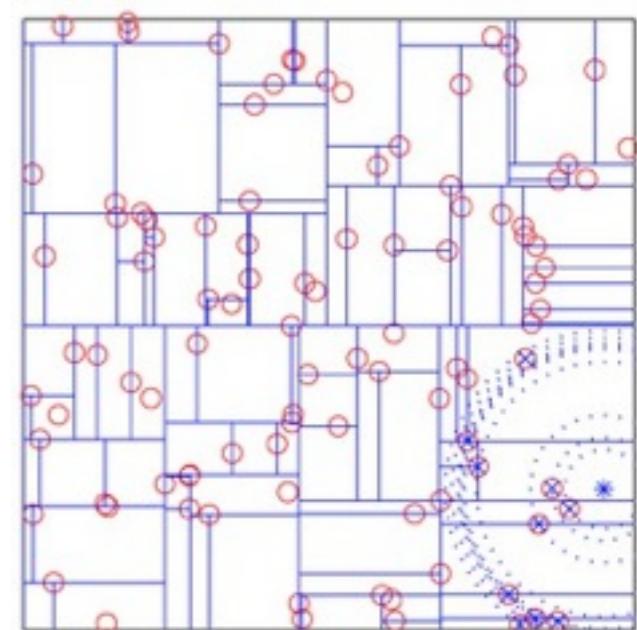
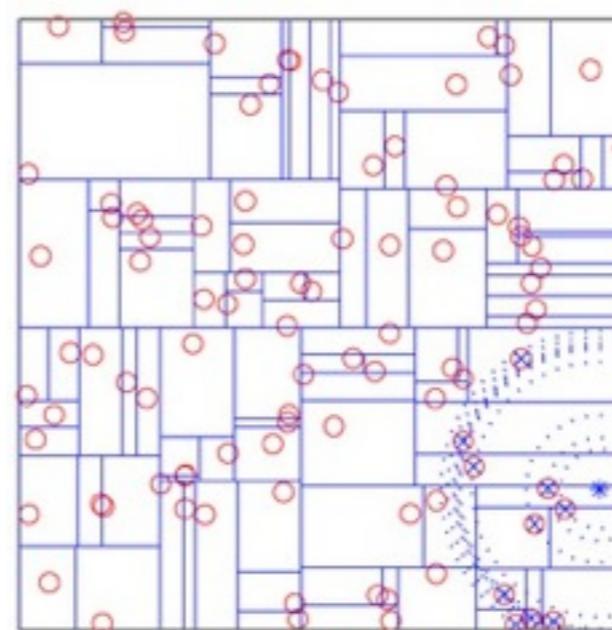
- **Hierarchical k-means clustering**

- [Nistér and Stewénius 2006]
[Hastie et al. 2001]
- apply k-means recursively
- logarithmic encoding
- **`vl_hikmeans`**



- **kd-trees**

- multi-dimensional logarithmic search
[Friedman et. al 1977]
- best-bin search and random forests
[Muja and Lowe 2009]



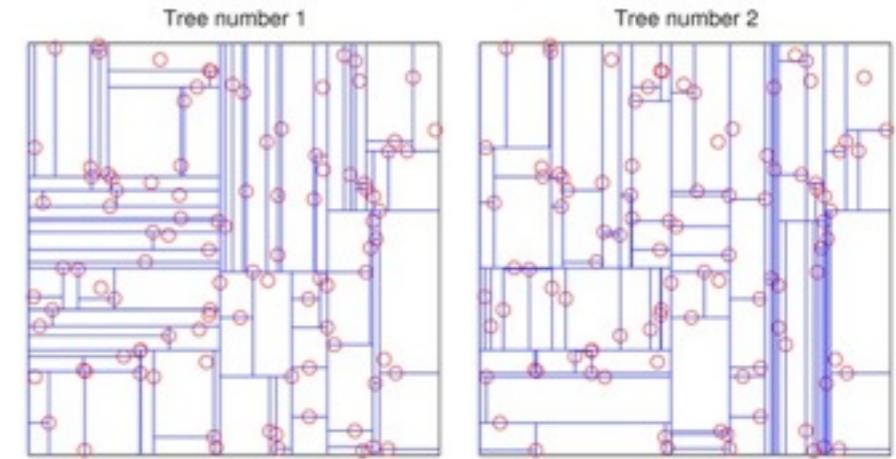
Randomized kd-tree forests

- **Build the tree**

```
kdtree = vl_kdtreebuild(X) ;
```

- random forests

```
kdtree = vl_kdtreebuild(X, 'NumTrees', 2) ;
```



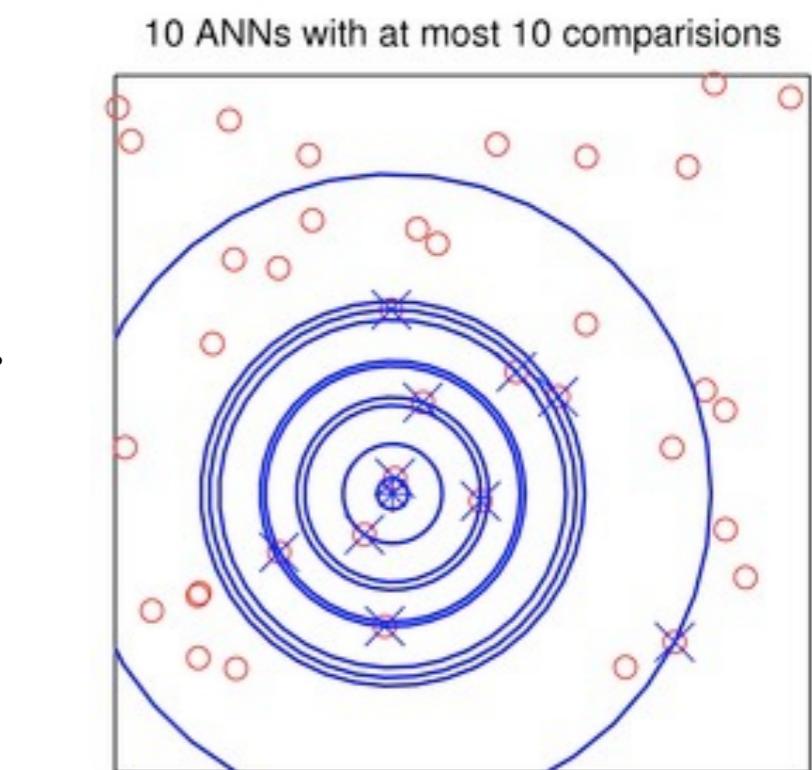
- **Query the tree**

- Exact nearest neighbors

```
[ind, dist] = vl_kdtreequery(kdtree, X, Q, ...
    'NumNeighbors' 10) ;
```

- Approximate nearest neighbors

```
[ind, dist] = vl_kdtreequery(kdtree, X, Q, ...
    'NumNeighbors' 10, ...
    'MaxComparisons', 10)
```



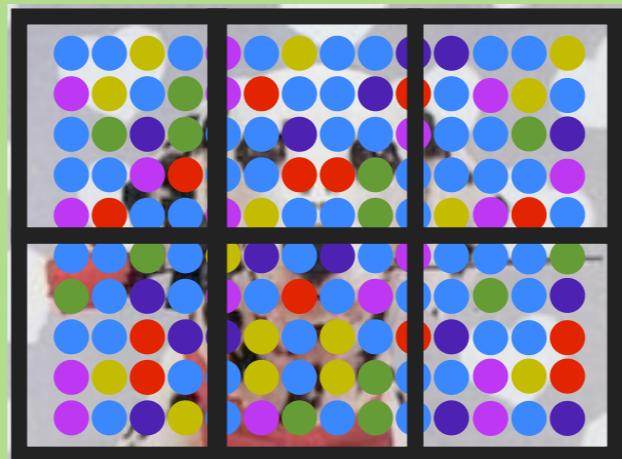
- Random kd-tree forest implementation

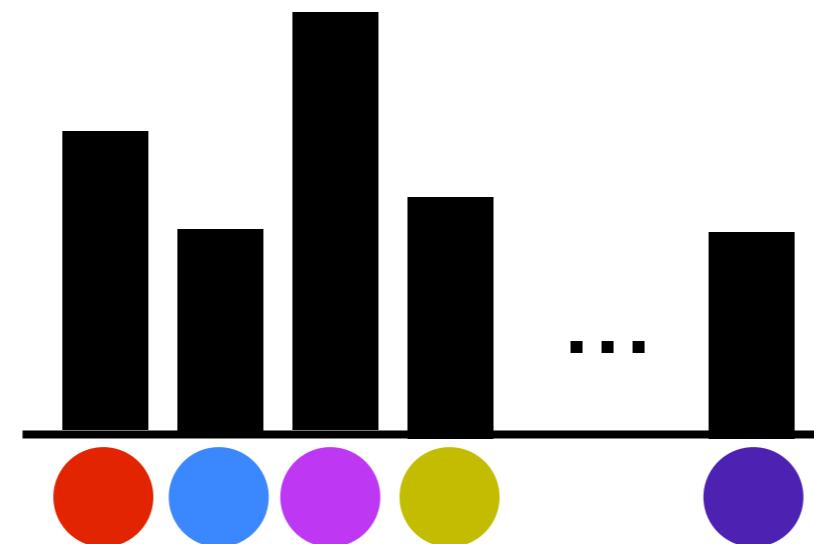
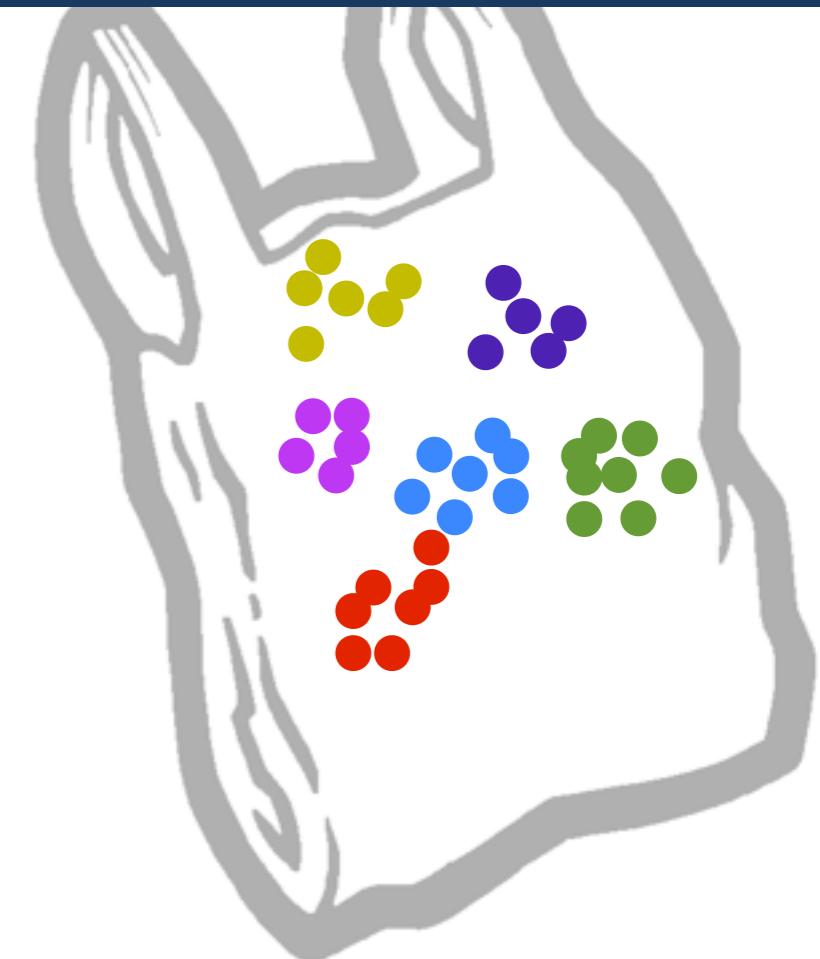
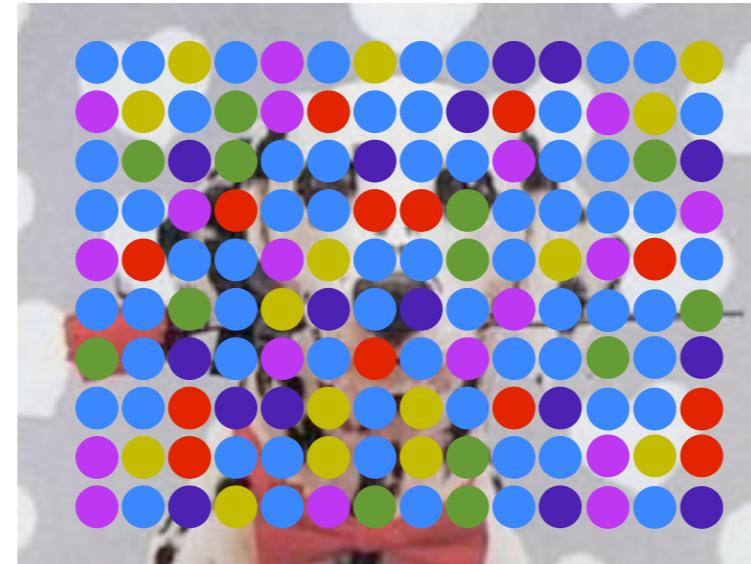
- Equivalent to FLANN

- [Muja and Lowe 2009]

- <http://www.cs.ubc.ca/~mariusm/index.php/FLANN/FLANN>

Spatial Histograms



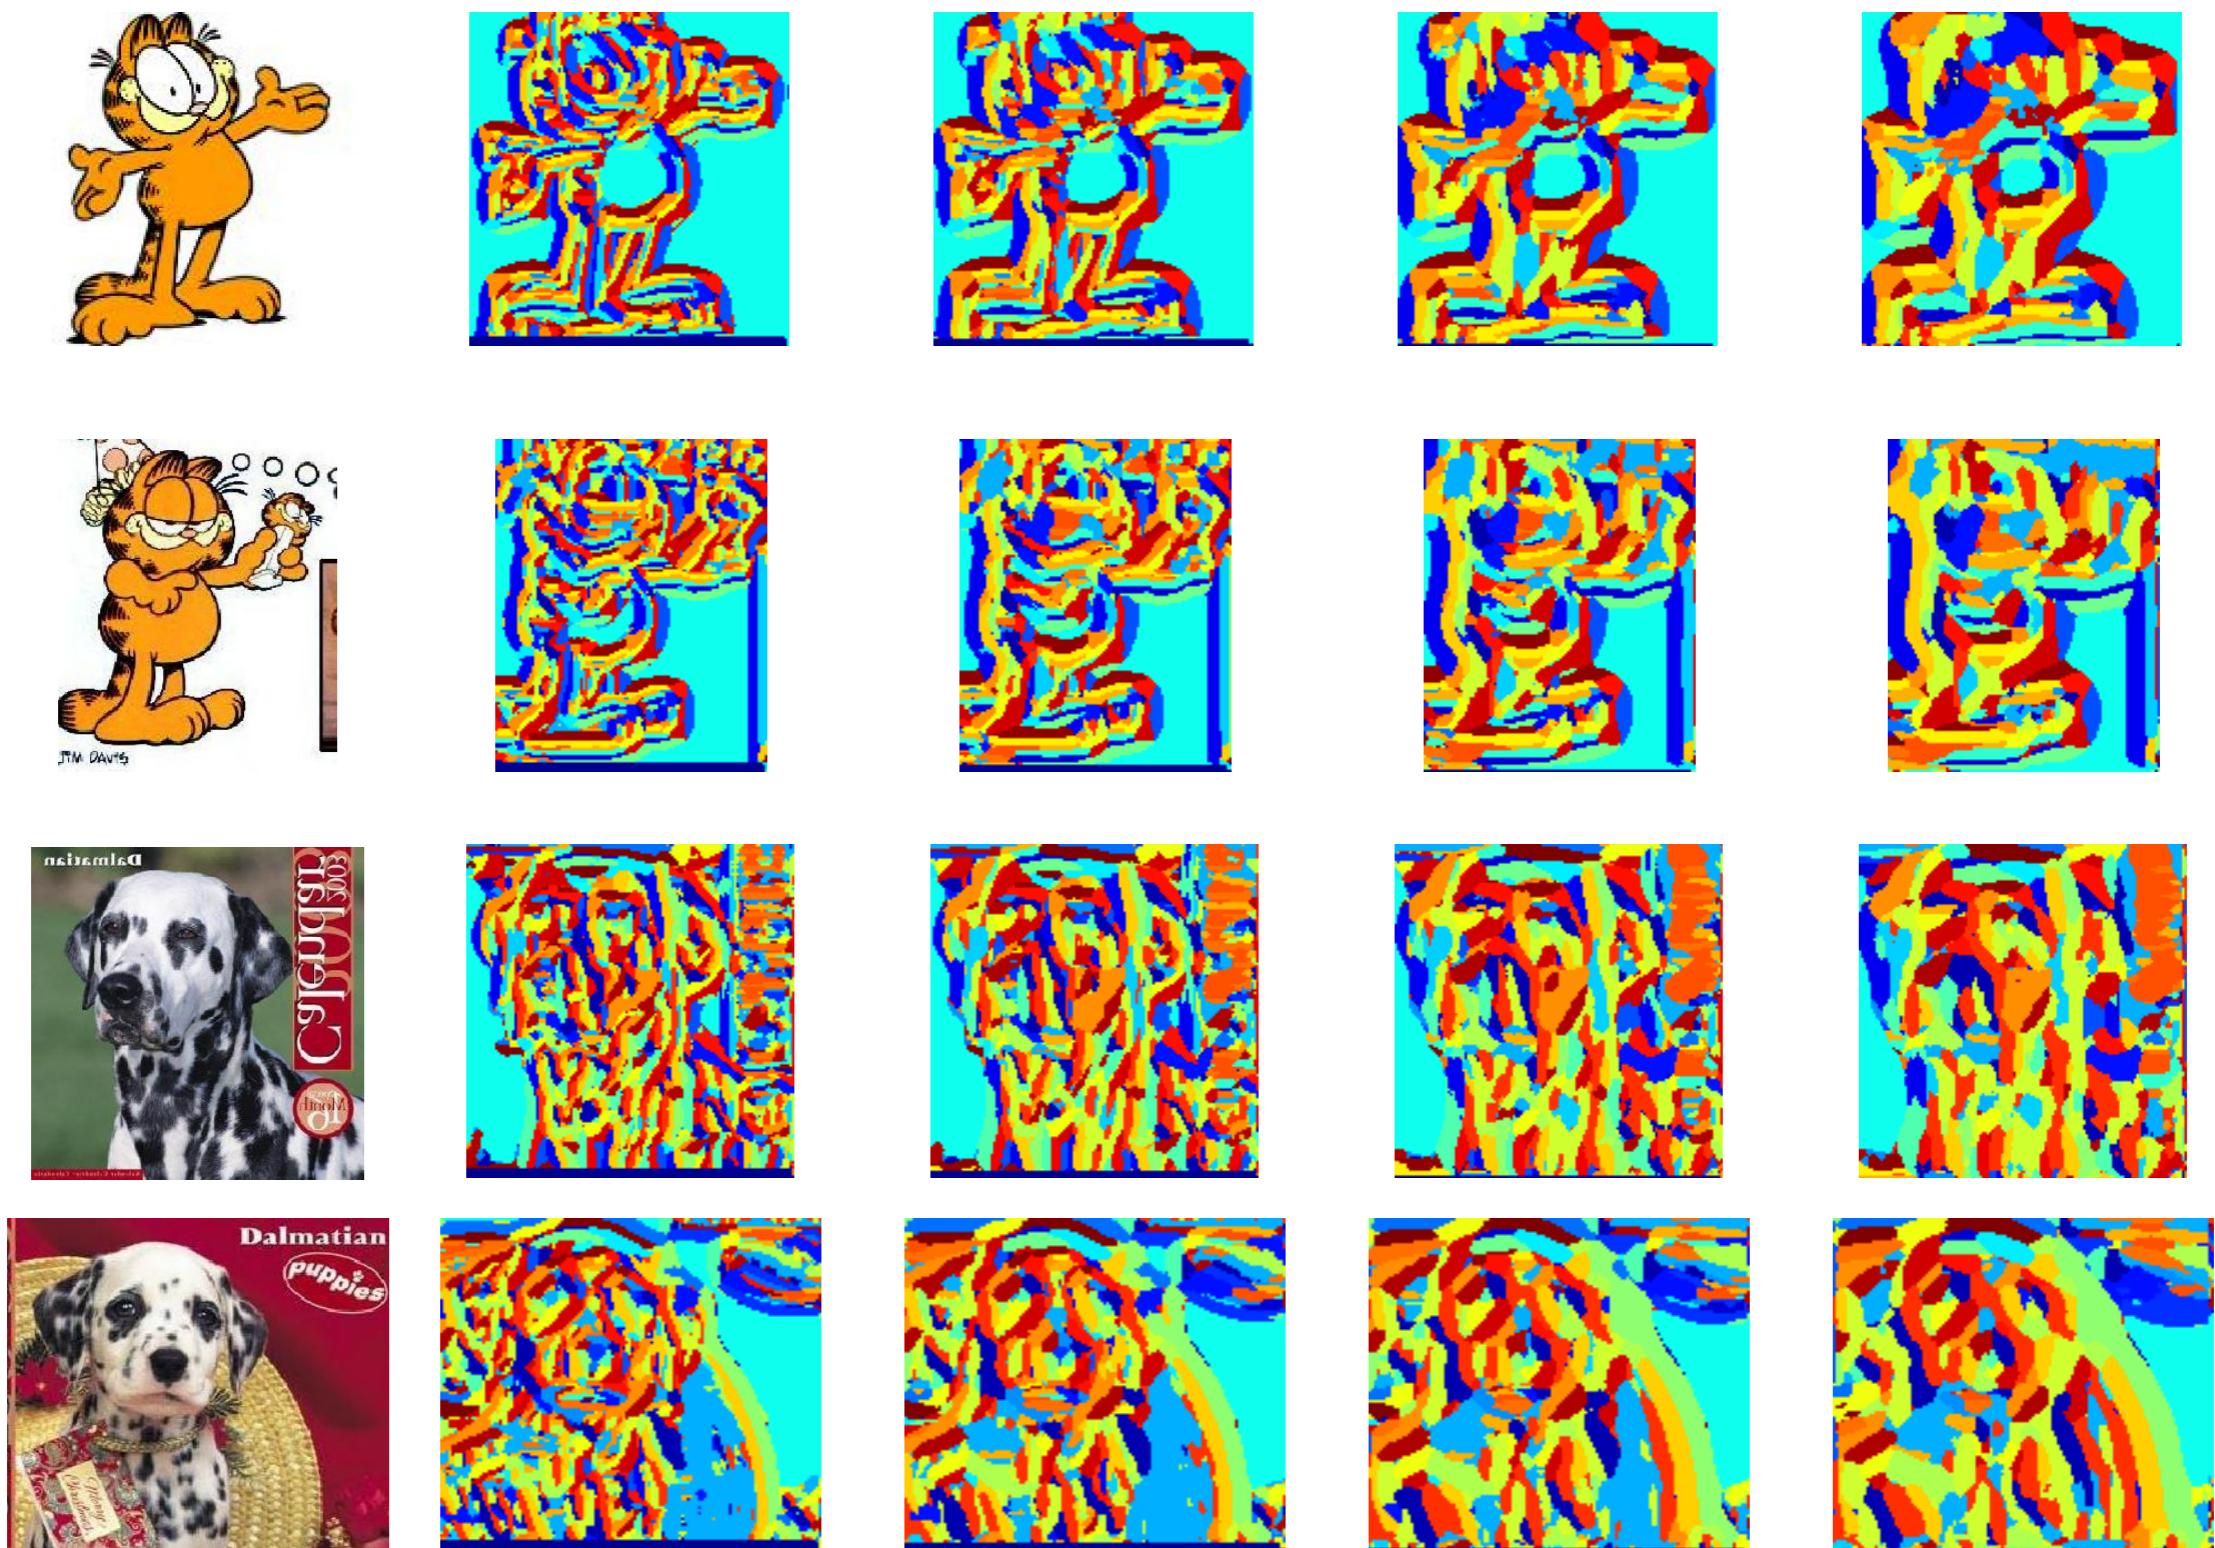


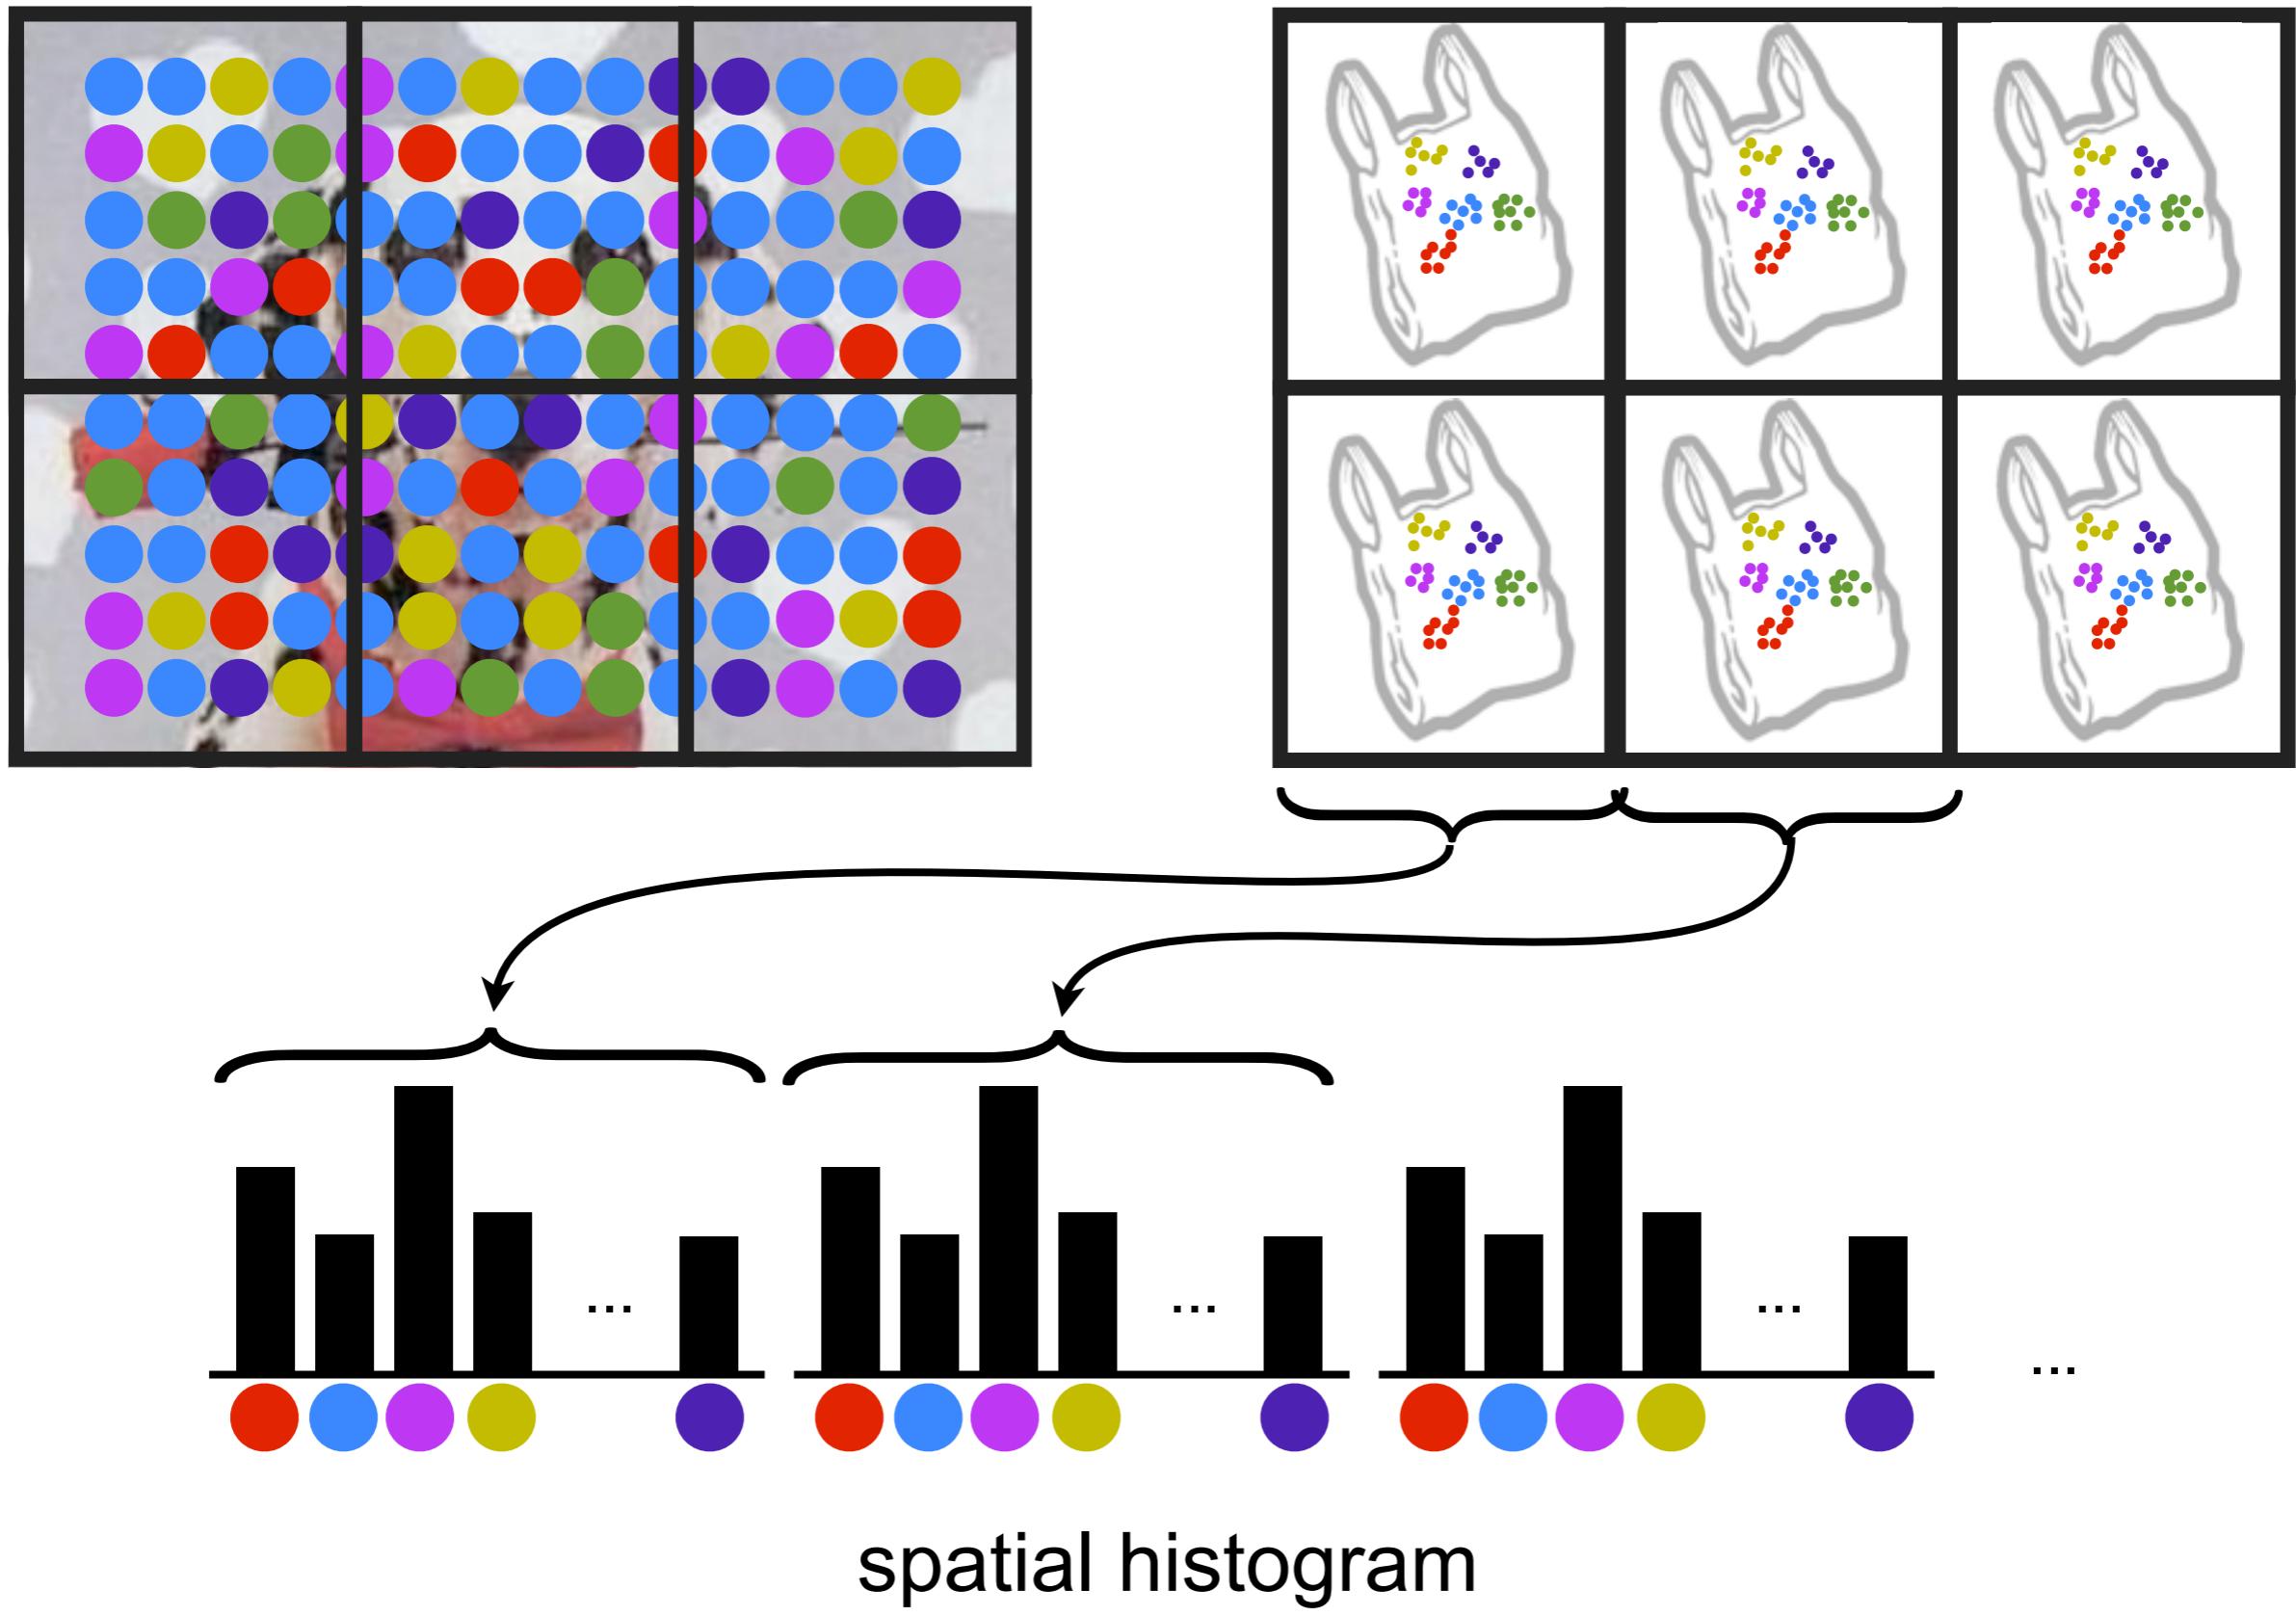
[Csurka et al. 2004]

histogram (bag) of visual words

Example encodings

SIFT descriptor scale





[Lazebnik et al. 2004]

VLFeat utility functions

- MEX-accelerated primitives
- Speed-up MATLAB in key operations
- Examples:

- **Binning**

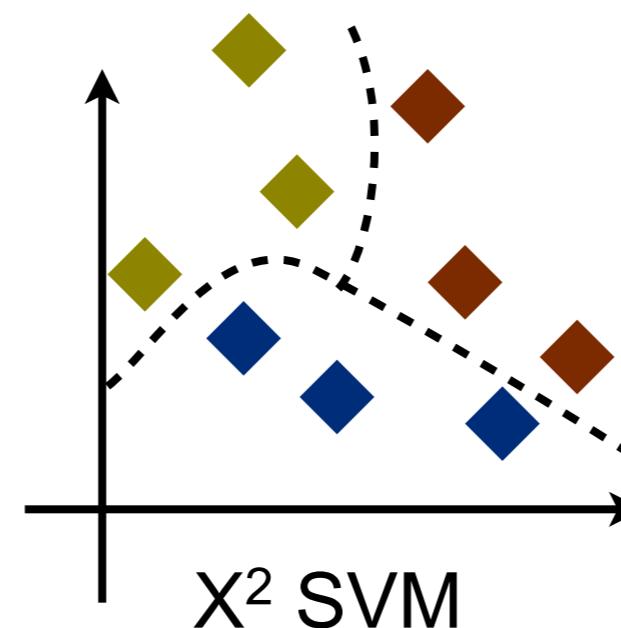
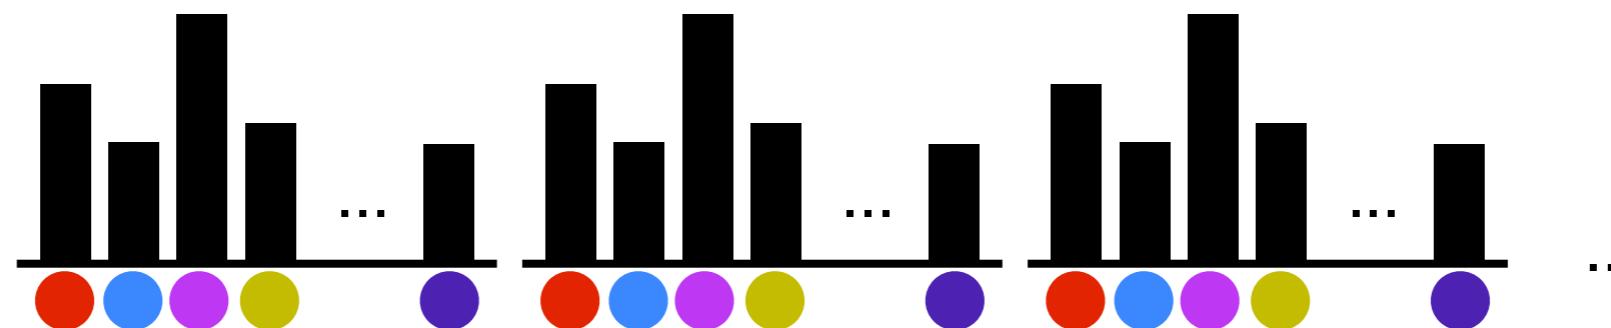
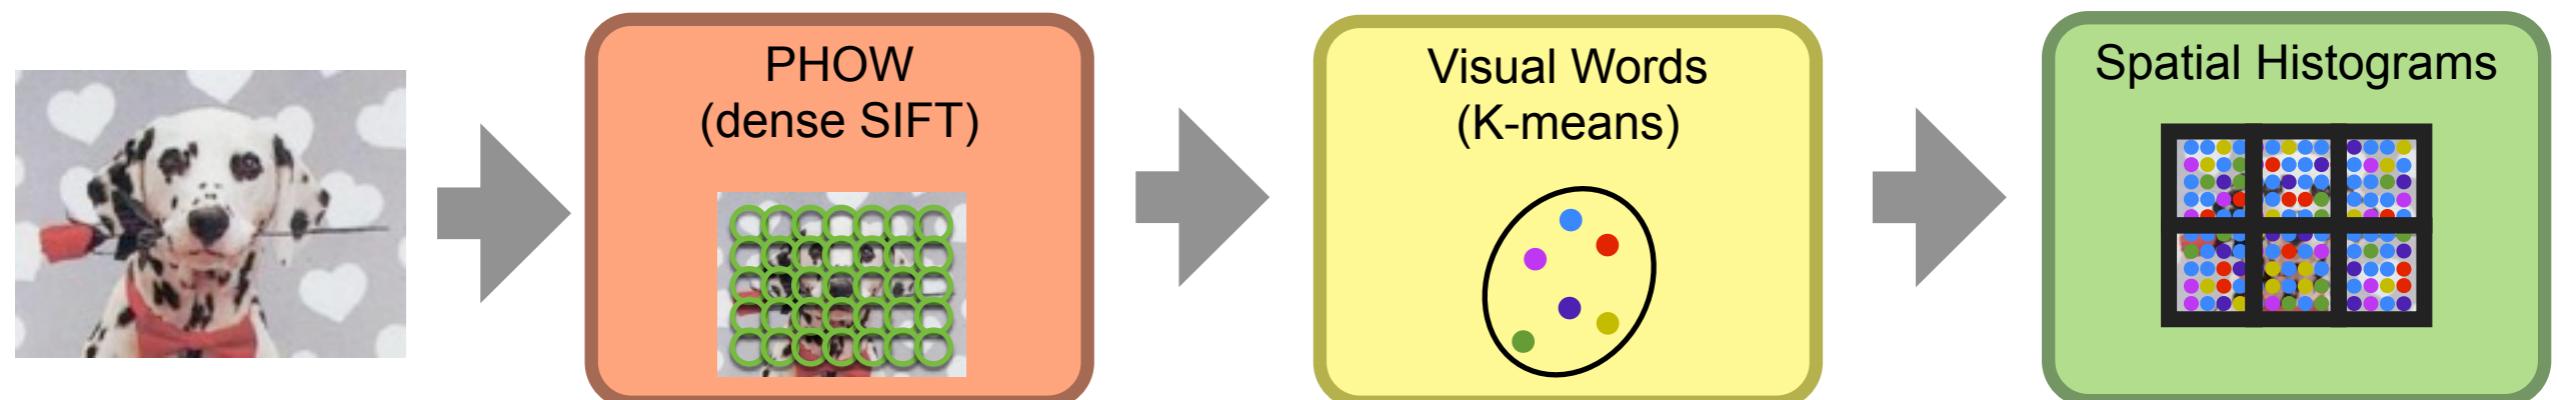
```
% quantize location  
binsx = vl_binsearch(linspace(1,width,conf.numSpatialX+1), frames(1,:)) ;  
binsy = vl_binsearch(linspace(1,height,conf.numSpatialY+1), frames(2,:)) ;
```

- **Binned summations**

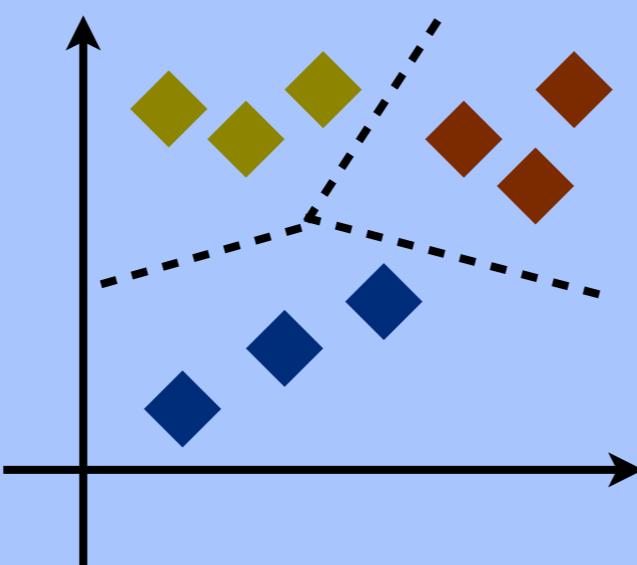
```
% histogram computation  
bins = sub2ind([conf.numSpatialY, conf.numSpatialX, conf.numWords], ...  
               binsy,binsx,binsa) ;  
hist = zeros(conf.numSpatialY * conf.numSpatialX * conf.numWords, 1) ;  
hist = vl_binsum(hist, ones(size(bins)), bins) ;
```

Summary: image descriptors

34



SVM



Fisher discriminant	LDA	Dirichelet processes
	kernel density estimation	separating hyperplane
kernel methods	logistic regression	
nearest neighbors		topic models

AdaBoost



random forests

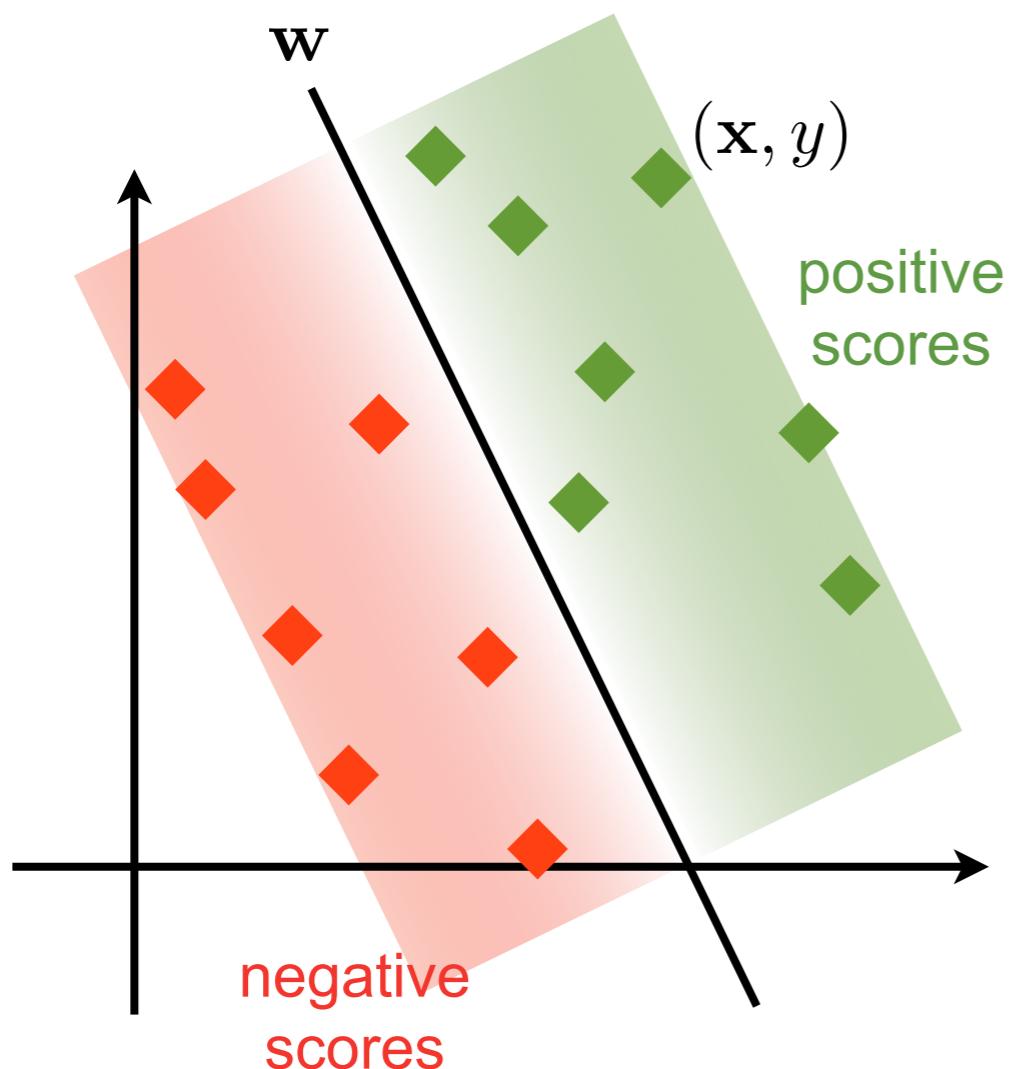
- **SVMs**

- often *state-of-the-art*
- *simple* to use and tune
- *efficient* with the latest improvements

[Shawe-Taylor and Cristianini 2000]

[Schölkopf and Smola 2002]

[Hastie et al. 2003]



Discriminant score

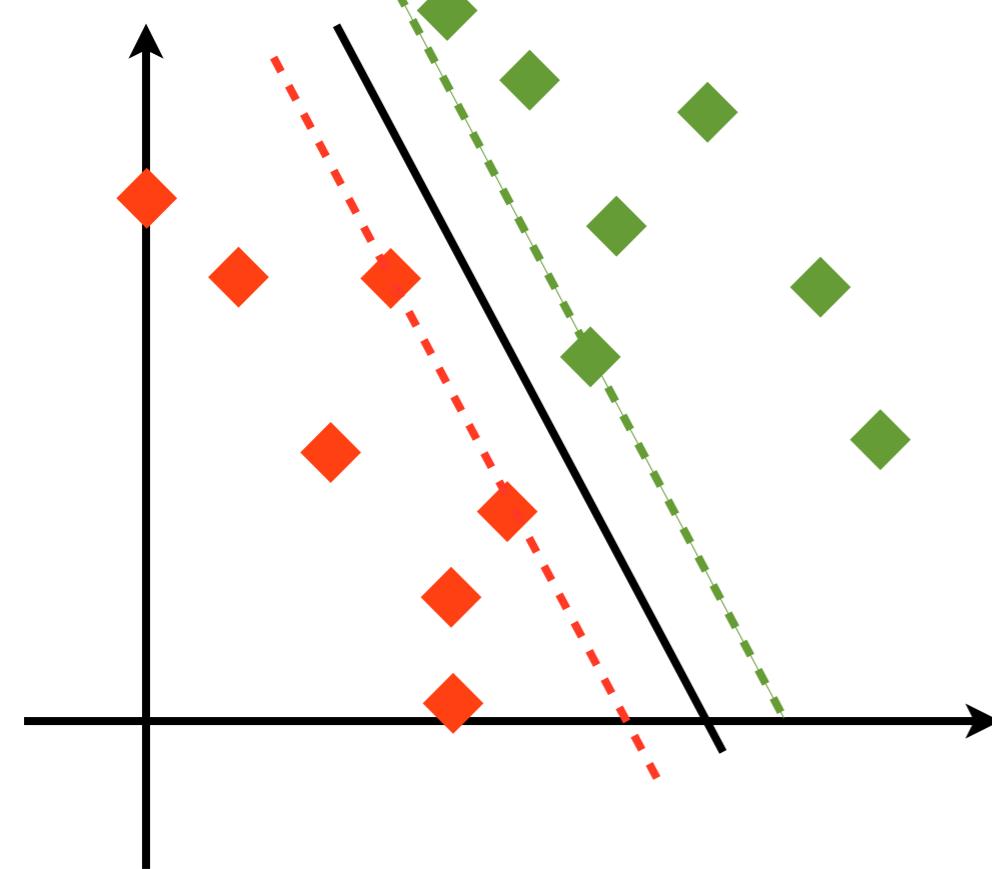
$$f(\mathbf{x}; \mathbf{w}) = \langle \mathbf{w}, \mathbf{x} \rangle$$

$$\min_{\mathbf{w}} \frac{\lambda}{2} \|\mathbf{w}\|^2 + \frac{1}{N} \sum_{i=1}^N \ell(\mathbf{w}; (\mathbf{x}_i, y_i))$$

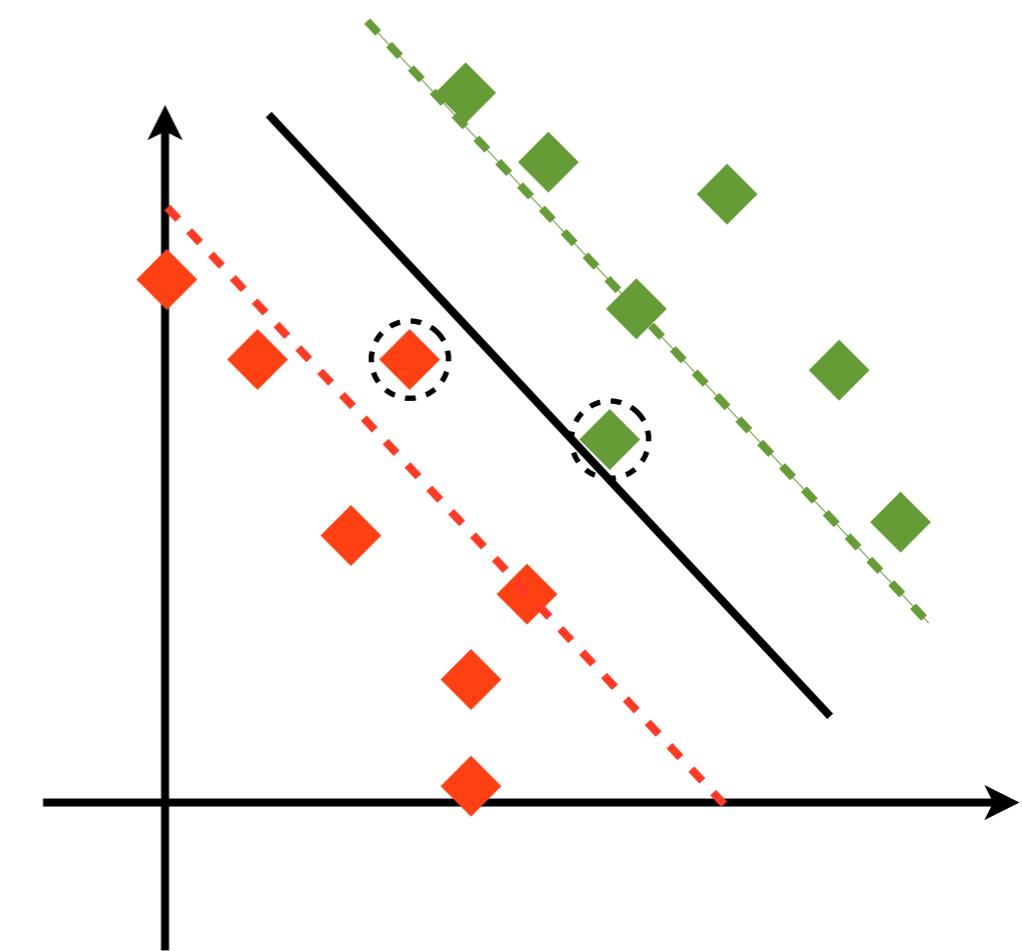
$$\min_{\mathbf{w}} \frac{\lambda}{2} \|\mathbf{w}\|^2 + \frac{1}{N} \sum_{i=1}^N \ell(\mathbf{w}; (\mathbf{x}_i, y_i))$$

hard loss

$$\ell(\mathbf{w}; (\mathbf{x}, y)) = \begin{cases} 0, & y\langle \mathbf{w}, \mathbf{x} \rangle > 1, \\ +\infty, & \text{otherwise.} \end{cases}$$

**hinge loss**

$$\ell(\mathbf{w}; (\mathbf{x}, y)) = \begin{cases} 0, & y\langle \mathbf{w}, \mathbf{x} \rangle > 1, \\ 1 - y\langle \mathbf{w}, \mathbf{x} \rangle, & \text{otherwise.} \end{cases}$$



Linear SVM: Fast algorithms

- **SVM^{perf}**

- [Joachims 2006]
- http://www.cs.cornell.edu/People/tj/svm_light/svm_perf.html
- cutting plane + one slack formulation
- **from superlinear to linear complexity**

- **LIBLINEAR**

- [R.-E. Fan et al. 2008]
- <http://www.csie.ntu.edu.tw/~cjlin/liblinear/>

I1 loss

$$\ell(\mathbf{w}; (\mathbf{x}, y)) = \max\{0, 1 - y\langle \mathbf{w}, \mathbf{x} \rangle\}$$

I2 loss

$$\ell(\mathbf{w}; (\mathbf{x}, y)) = (\max\{0, 1 - y\langle \mathbf{w}, \mathbf{x} \rangle\})^2$$

logistic loss

$$\ell(\mathbf{w}; (\mathbf{x}, y)) = \log(1 + \exp(-y\langle \mathbf{w}, \mathbf{x} \rangle))$$

I² regularizer

$$\frac{\lambda}{2} \sum_{i=1}^D w_i^2$$

I¹ regularizer

$$\frac{\lambda}{2} \sum_{i=1}^D |w_i|$$

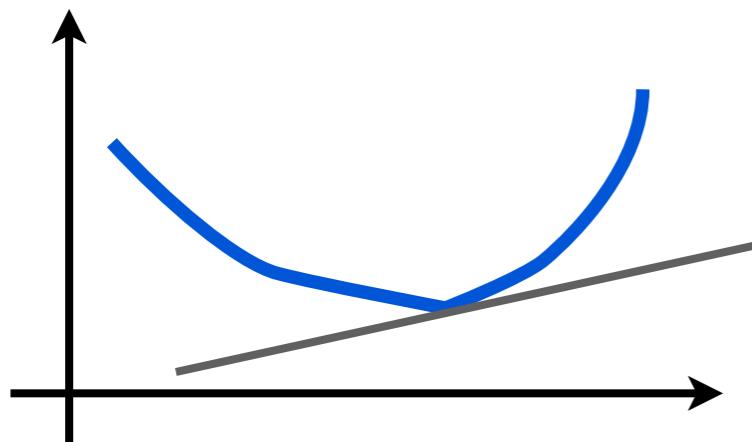
Linear SVM: PEGASOS

- [S. Shalev-Shwartz et al. 2010]
- <http://www.cs.huji.ac.il/~shais/code/index.html>

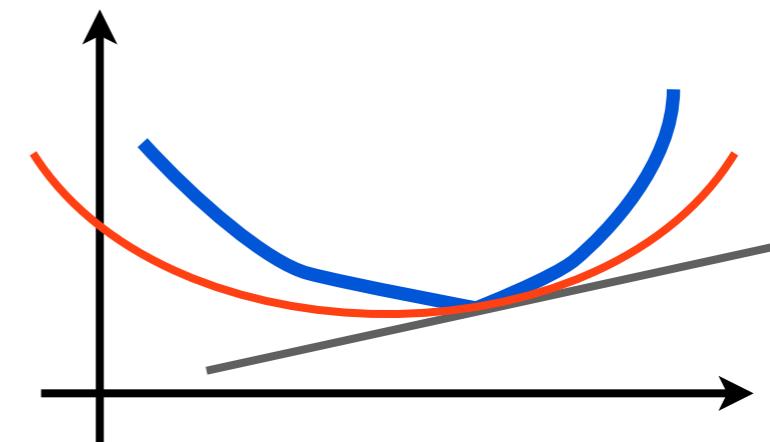
1 Pick a random data point (\mathbf{x}_i, y_i)

$$2 \quad \nabla E_t(\mathbf{w}) = \begin{cases} \lambda \mathbf{w}_t, & y_i \langle \mathbf{w}, \mathbf{x}_i \rangle > 1, \\ \lambda \mathbf{w}_t - y_i \mathbf{x}_i, & \text{otherwise.} \end{cases}$$

$$3 \quad \mathbf{w}_{t+1} = \mathbf{w}_t - \frac{1}{\lambda t} \nabla E_t(\mathbf{w})$$



non differentiable objective
⇒ **sub-gradient**



strongly-convex objective
⇒ **optimal schedule**

VLFeat PEGASOS

- **training**

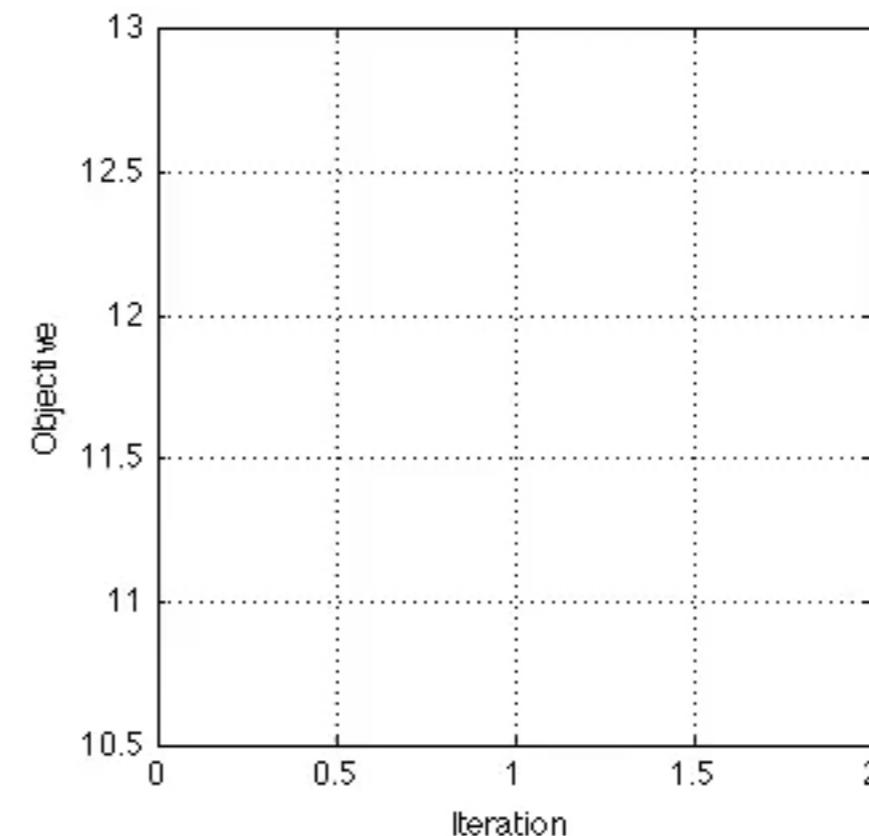
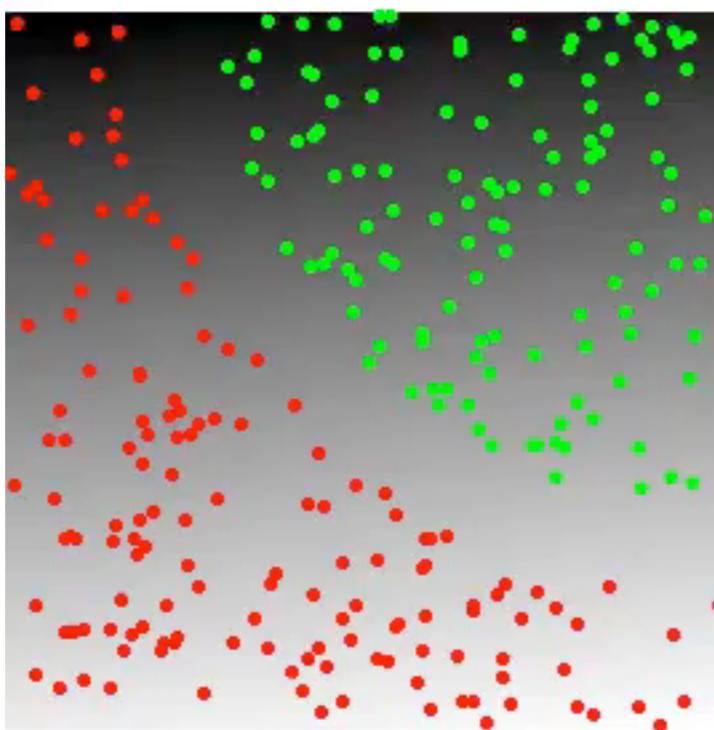
```
w = vl_pegasos(x, y, lambda)
```

- **testing**

```
w' * x
```

- Observations:

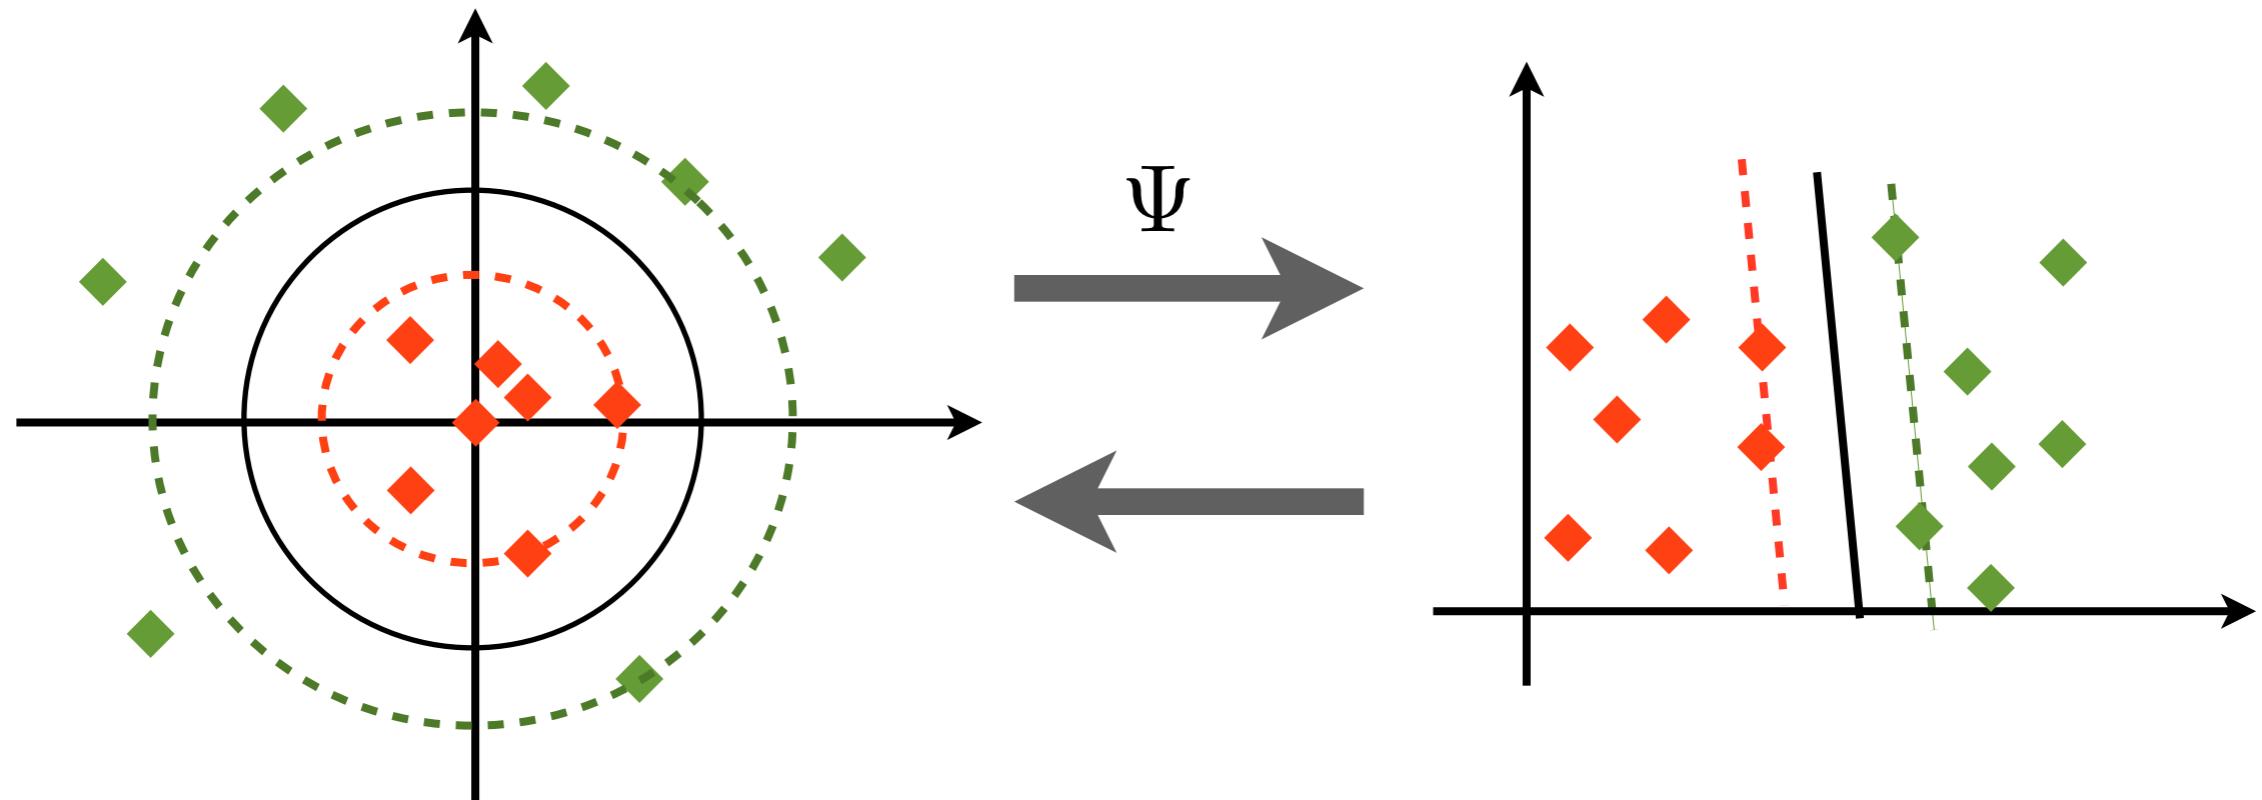
- gets a rough solution very quickly
- many solutions almost as good
- testing dominated by generalization error



Kernel Map

Ψ

- Generalize SVM through a *feature map*



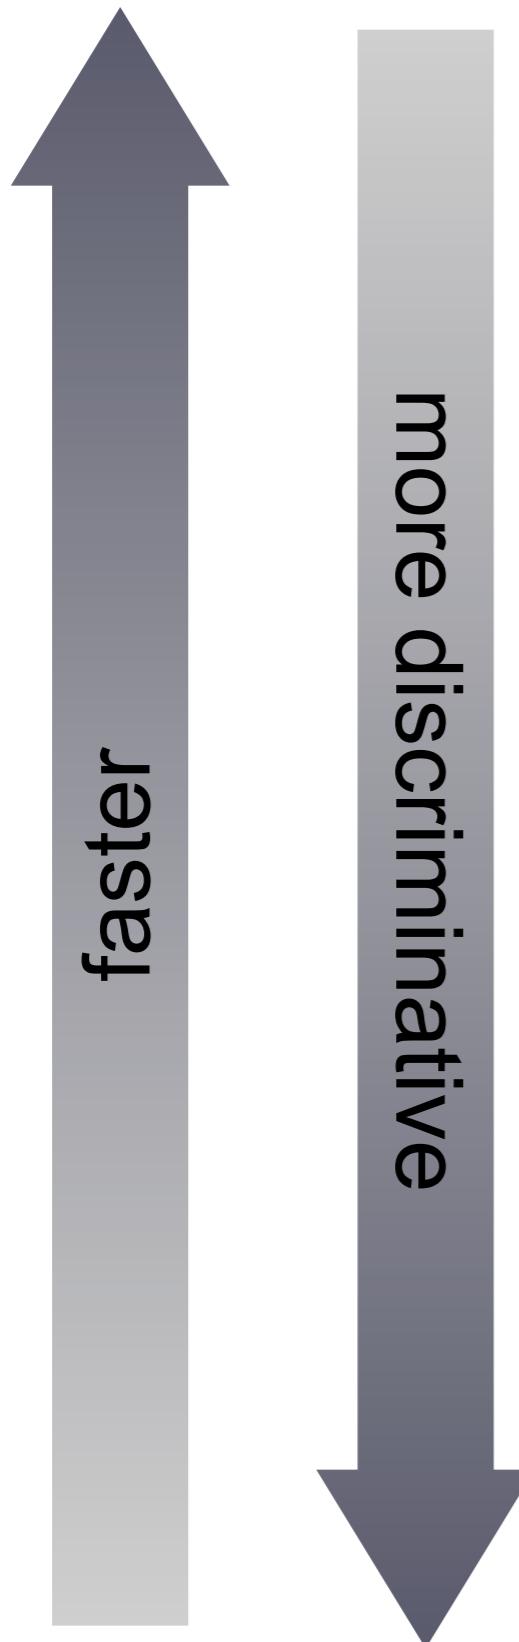
- The feature map may be defined implicitly by a **kernel function**

$$\mathbf{w} = \sum_{i=1}^N \beta_i \Psi(\mathbf{x}_i)$$

$$\langle \mathbf{w}, \Psi(\mathbf{x}) \rangle = \sum_{i=1}^N \beta_i \langle \Psi(\mathbf{x}_i), \Psi(\mathbf{x}) \rangle$$

$$K(\mathbf{x}_i, \mathbf{x}) = \langle \Psi(\mathbf{x}_i), \Psi(\mathbf{x}) \rangle$$

Common kernels



linear

$$K(\mathbf{x}, \mathbf{y}) = \sum_{l=1}^D \mathbf{x}_l \mathbf{y}_l$$

[Zhang et al. 2007]
[Vedaldi et. al 2010]

Fast kernel and
distance
matrices with
vl_alldist

additive

$$K(\mathbf{x}, \mathbf{y}) = \sum_{l=1}^D k(\mathbf{x}_l, \mathbf{y}_l)$$

additive RBF

$$K(\mathbf{x}, \mathbf{y}) = \exp\left(-\frac{1}{2\sigma^2} \sum_{l=1}^D d^2(\mathbf{x}_l, \mathbf{y}_l)\right)$$

Common kernels

faster

e discriminative

$l=1$

additive

$$K(\mathbf{x}, \mathbf{y}) = \sum_{l=1}^D k(\mathbf{x}_l, \mathbf{y}_l)$$

additive RBF

$$K(\mathbf{x}, \mathbf{y}) = \exp\left(-\frac{1}{2\sigma^2} \sum_{l=1}^D d^2(\mathbf{x}_l, \mathbf{y}_l)\right)$$

Hellinger's

$$k(x, y) = \sqrt{xy}$$

$$d^2(x, y) = (\sqrt{x} - \sqrt{y})^2$$

X²

$$k(x, y) = \frac{2xy}{x + y}$$

$$d^2(x, y) = \frac{(x - y)^2}{x + y}$$

intersection

$$k(x, y) = \min\{x, y\} \quad d^2(x, y) = |x - y|$$

Non-linear SVMs

- **LIBSVM**
 - [C.-C. Chang and C.-J. Lin 2001]
 - <http://www.csie.ntu.edu.tw/~cjlin/libsvm/>
- **SVM^{light}**
 - [T. Joachims 1999]
 - <http://svmlight.joachims.org>
- Many other open source implementations ...



bottleneck
support vector expansion

Speeding-up non-linear SVMs

- Avoid support vector expansion
- Actually compute the feature map $\Psi(\mathbf{x}) \quad K(\mathbf{x}, \mathbf{y}) = \langle \Psi(\mathbf{x}), \Psi(\mathbf{y}) \rangle$
- Exact feature map is usually:

hard to compute

high dimensional

- Seek for approximations $\widehat{\Psi}(\mathbf{x}) \quad K(\mathbf{x}, \mathbf{y}) \approx \langle \widehat{\Psi}(\mathbf{x}), \widehat{\Psi}(\mathbf{y}) \rangle$
 - low dimensional
 - efficient computation
 - good approximation quality

reduced expansions

[C. K. I. Williams and M. Seeger 2001]
 [Bach and Jordan 2006]
 [Bo and Sminchiescu 2009]

direct approximations

VLFeat homogeneous kernel map

- [Vedaldi and Zisserman 2010]
- Closed feature maps by Fourier transform in log domain

Hellinger's

$$k(x, y) = \sqrt{xy}$$

$$\Psi(x) = \sqrt{x}$$

Intersection

$$k(x, y) = \min\{x, y\}$$

$$\Psi(x) = \sqrt{\frac{2x}{\pi(1 + 4\omega^2)}} e^{i\omega \log x}$$

X²

$$k(x, y) = \frac{2xy}{x + y}$$

$$\Psi(x) = \sqrt{x \operatorname{sech}(\pi\omega)} e^{i\omega \log x}$$

- Approximate by uniform sampling and truncation

X² 3x approximation

$$\widehat{\Psi}(x) = \begin{bmatrix} 0.8 \sqrt{x} \\ 0.6 \sqrt{x} \cos(0.6 \log x) \\ 0.6 \sqrt{x} \sin(0.6 \log x) \end{bmatrix}$$

VLFeat homogeneous kernel map

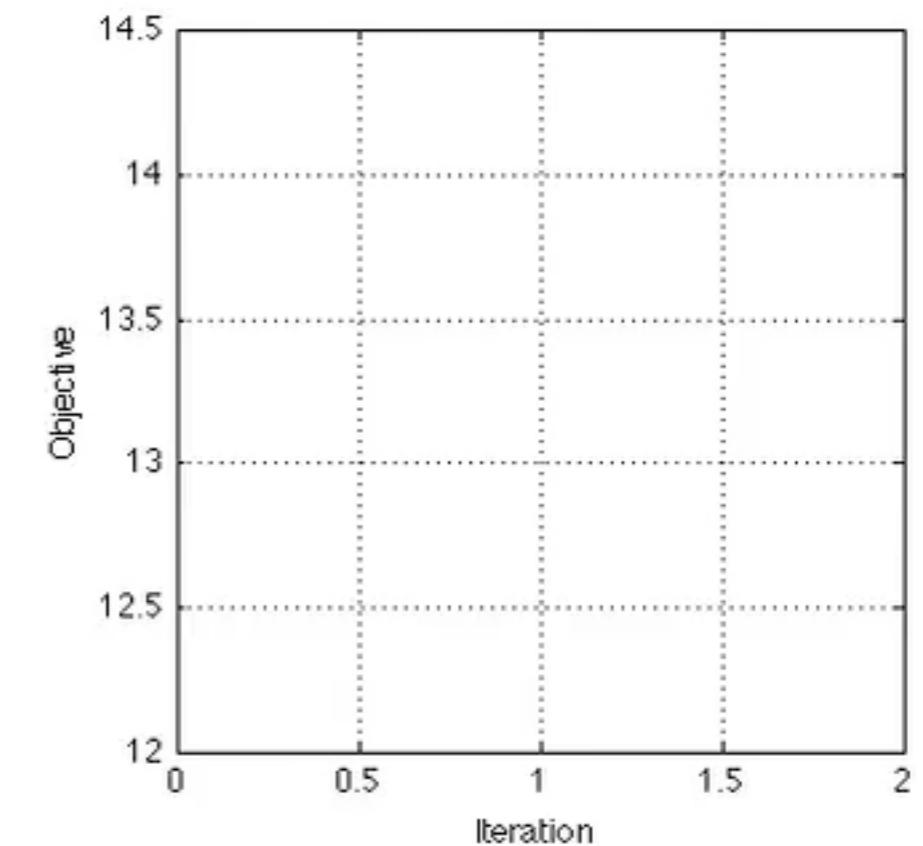
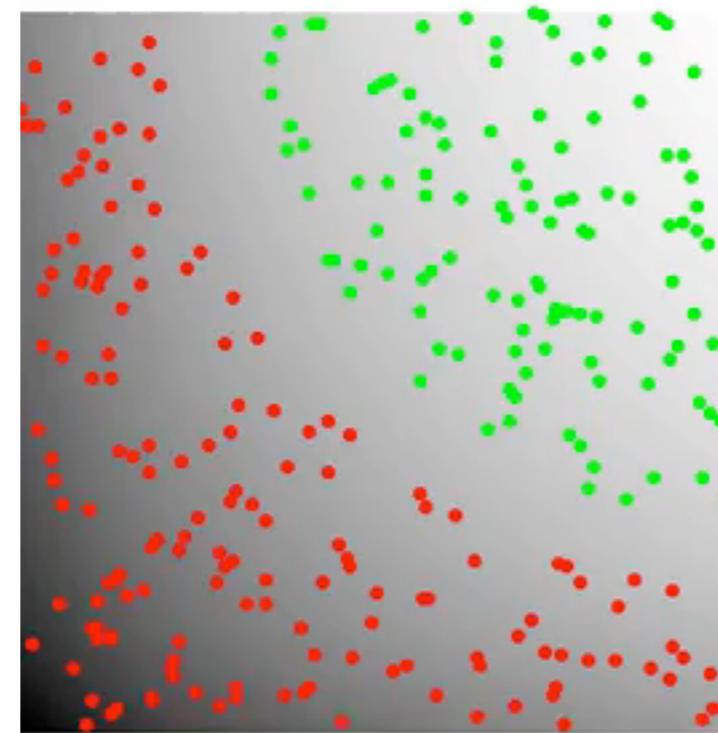
- **x** data (along columns)
- **Linear SVM**

```
w = vl_pegasos(x, y, lambda)
```

- **χ^2 SVM**

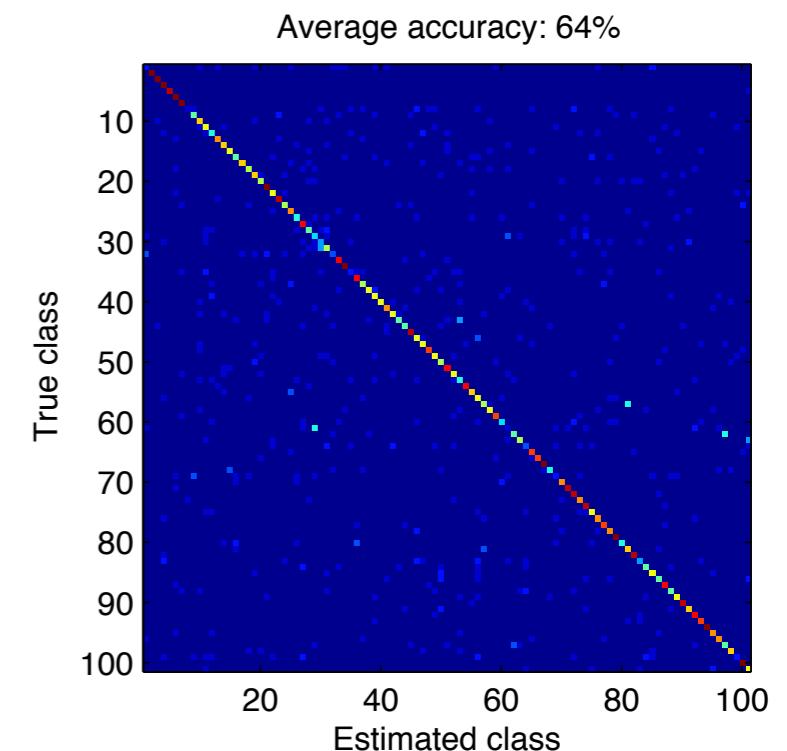
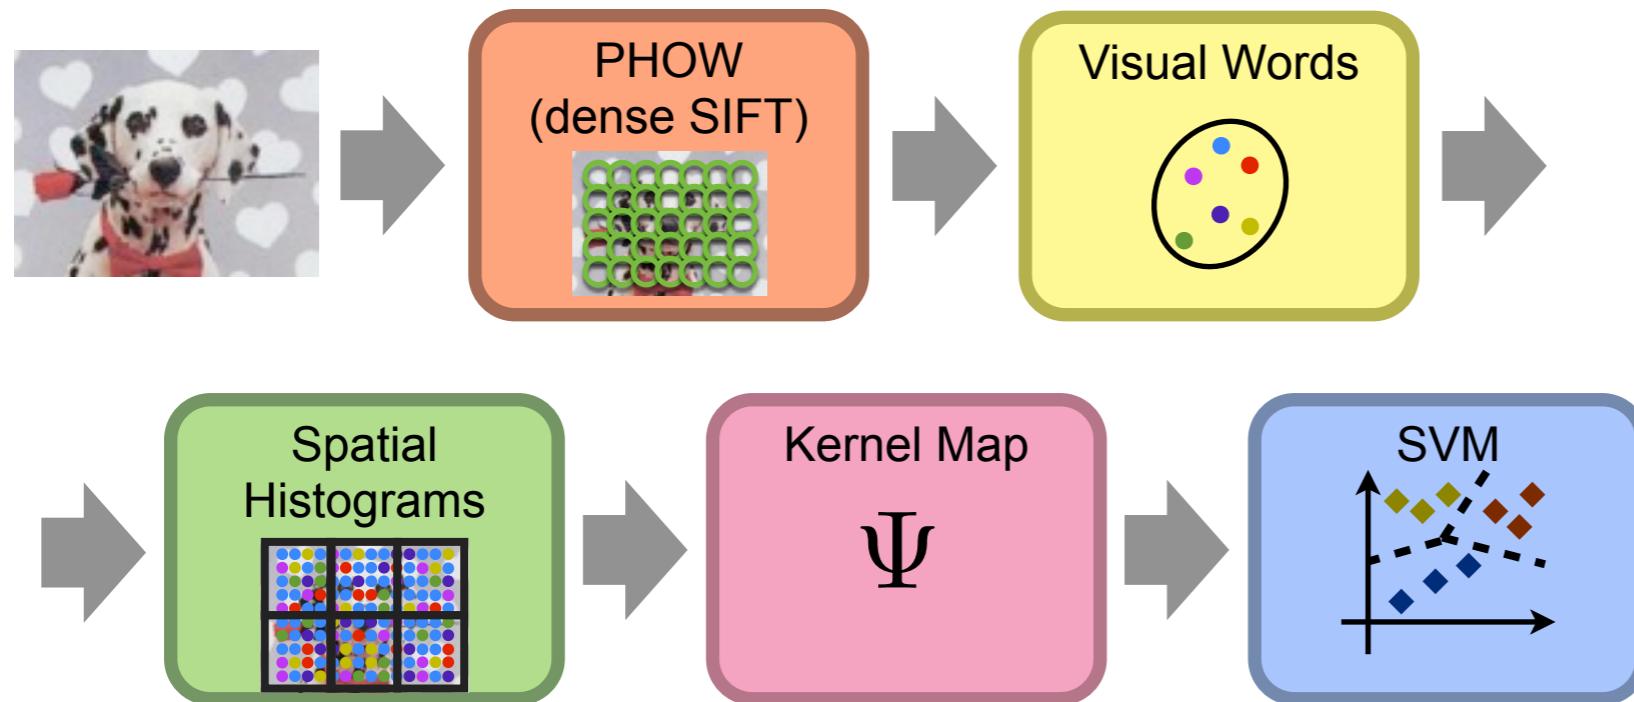
```
psix = vl_homkermap(x, 1, .6, 'kchi2') ;
w = vl_pegasos(psix, y, lambda) ;
```

- Advantages
 - universal (no training)
 - fast computation
 - high accuracy
 - on-line training and HOG like detectors



Other direct feature maps

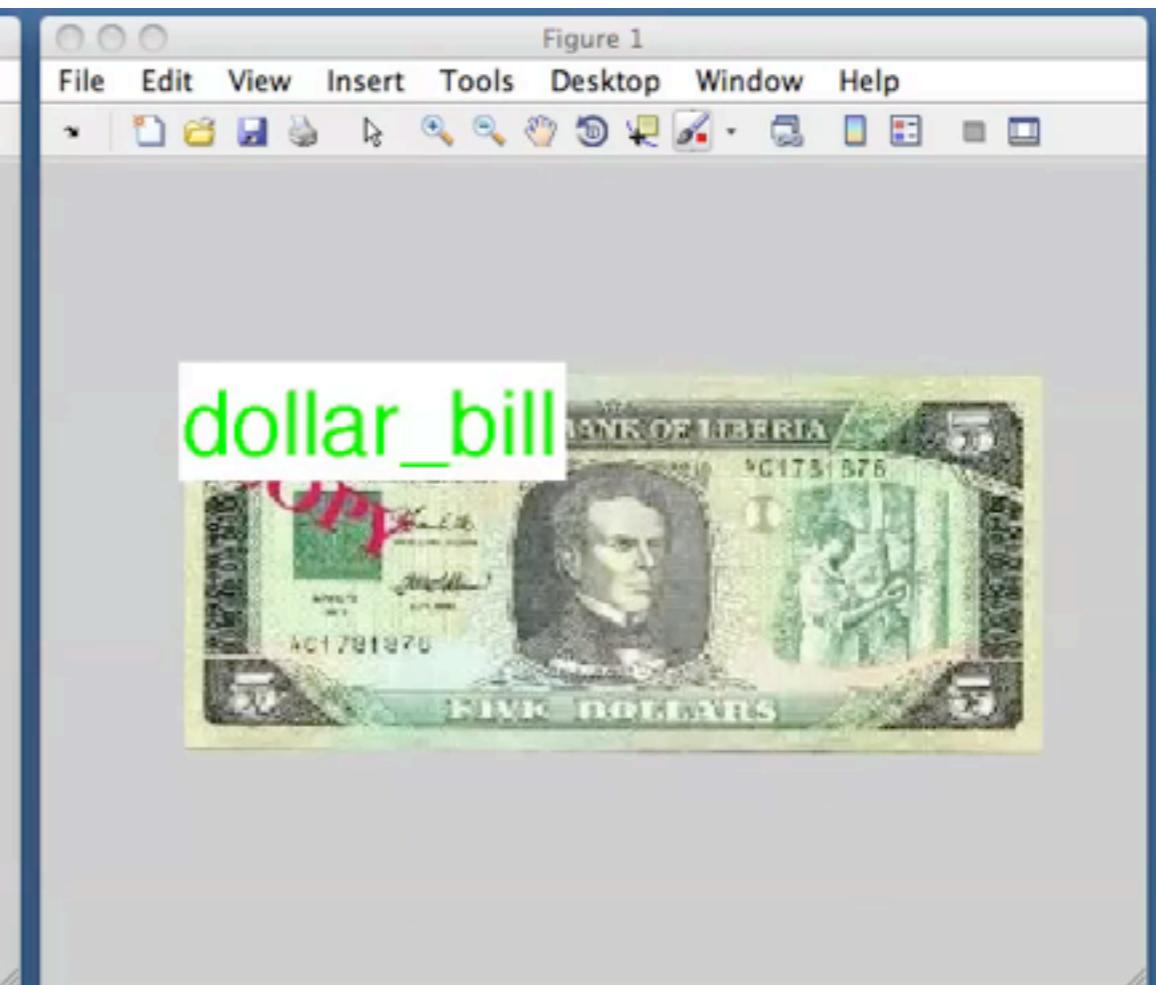
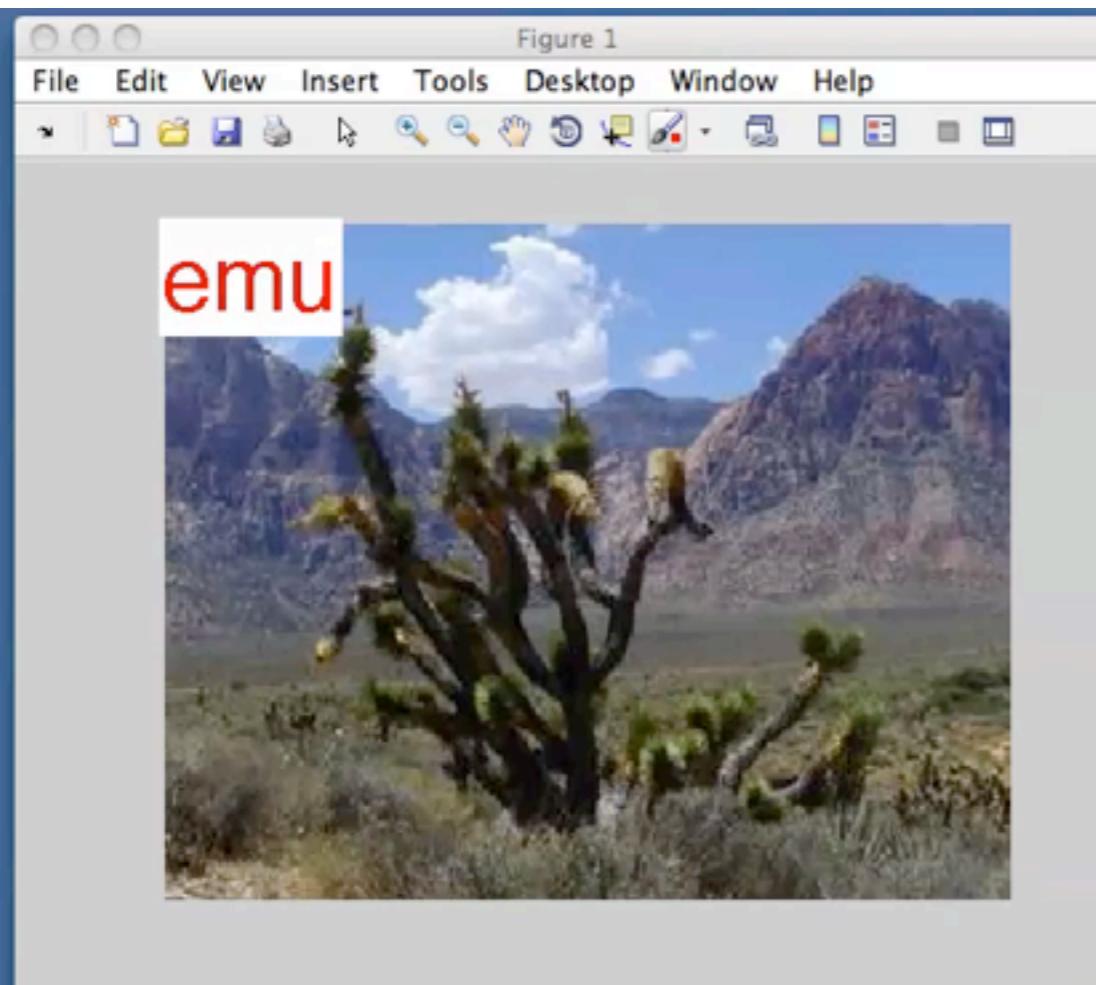
- **Explicit map for intersection kernel**
 - [Maji et al. 2008] and [Maji and Berg 2009]
 - <http://www.cs.berkeley.edu/~smaji/projects/ped-detector/>
 - <http://www.cs.berkeley.edu/~smaji/projects/add-models/>
- **Additive kernel-PCA** for additive kernels
 - [Perronnin et al. 2010]
- **Random Fourier features** for RBF kernels
 - [Rahimi and Recht 2007]
 - <http://people.csail.mit.edu/rahimi/random-features/>



- **PHOW features**
 - Fast dense SIFT ([vl_dsift](#))
- **Visual Words**
 - Elkan k-means ([vl_kmeans](#))
- **Spatial Histograms**
 - Convenience functions
([vl_binsum](#), [vl_binsearch](#))
- **Linear SVM**
 - PEGASOS ([vl_pegasos](#))
- **X² SVM**
 - Homogeneous kernel map
([vl_homkermap](#))

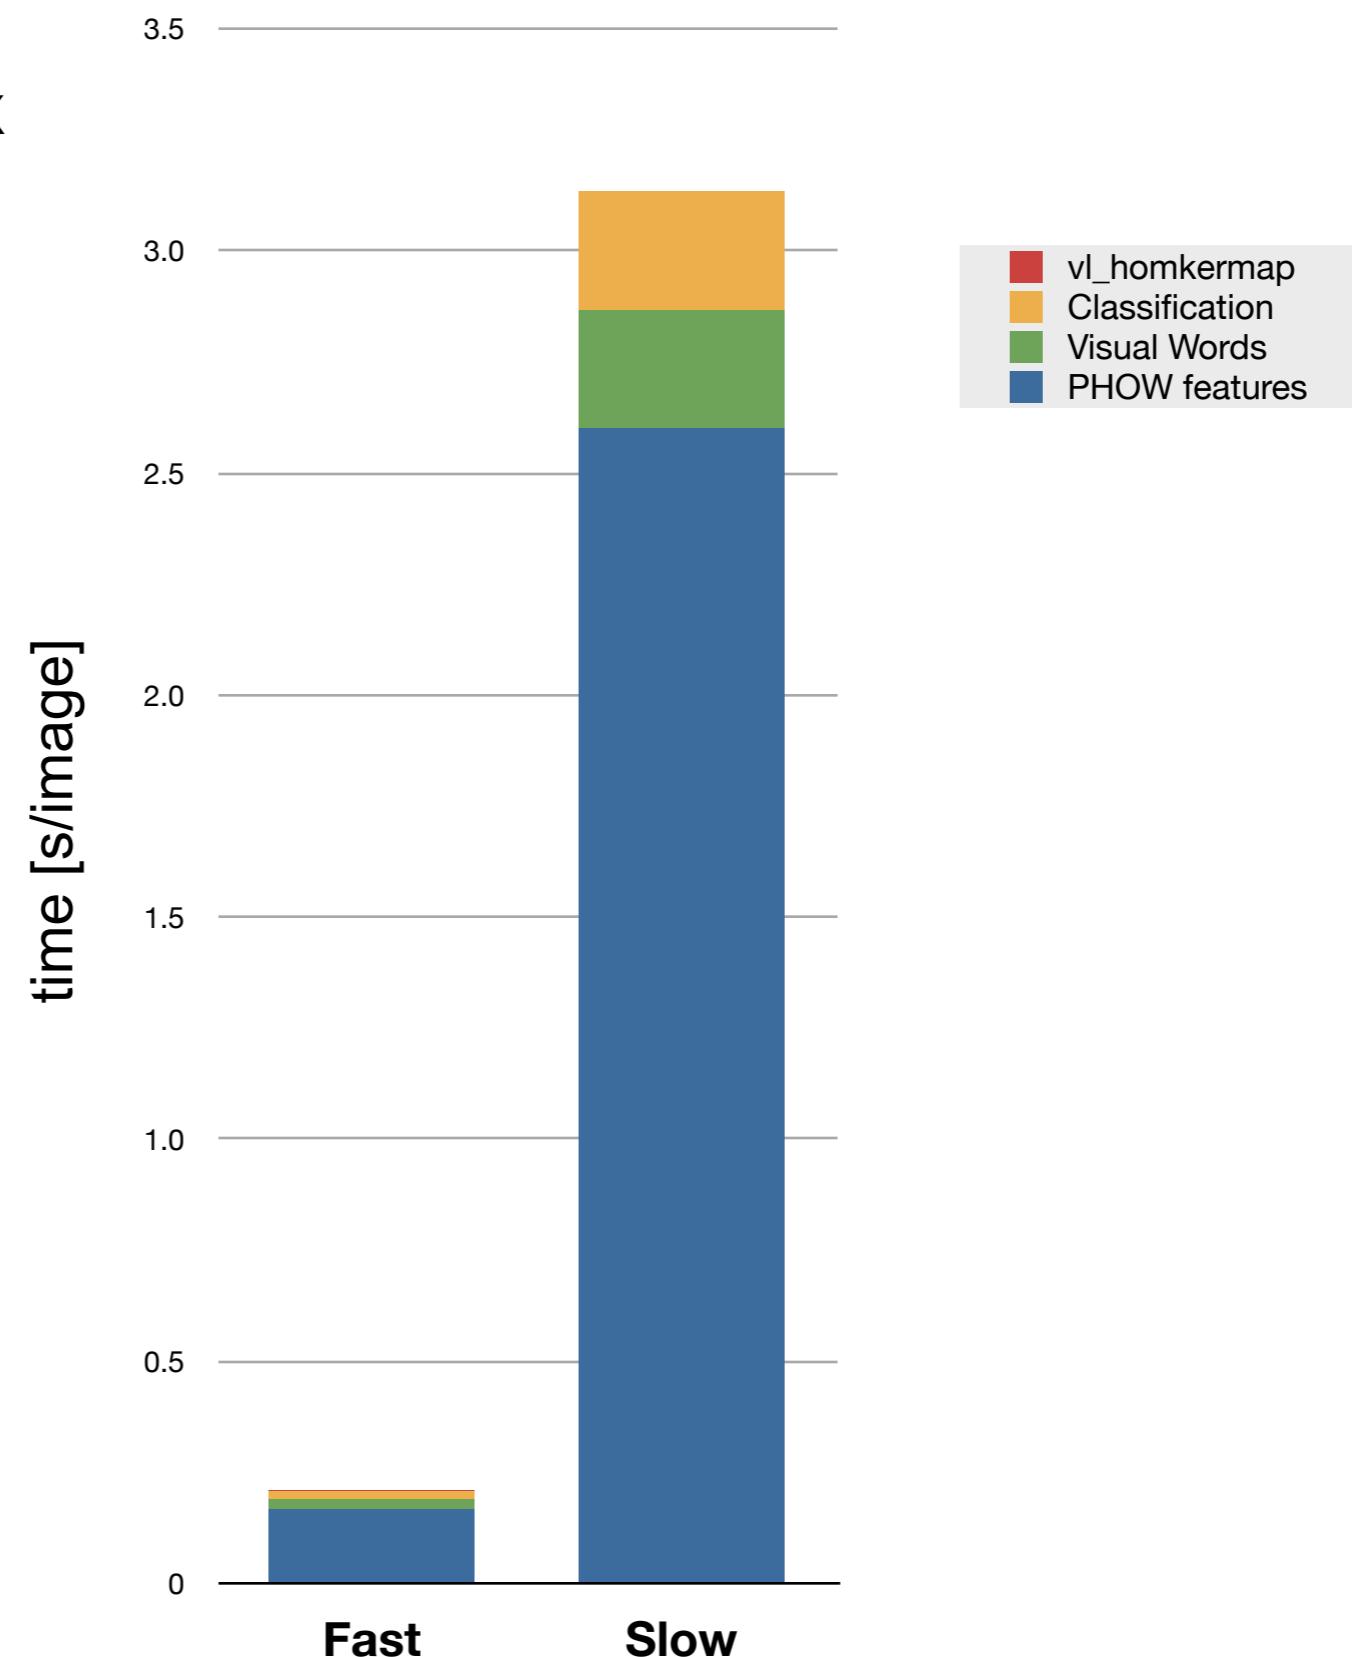
Demo

- Real time classification, 15 training images, 102 classes, 64% accuracy
- **Left**
standard dense SIFT, kernelized SVM
- **Right**
fast dense SIFT, kdtree, homogeneous kernel map SVM

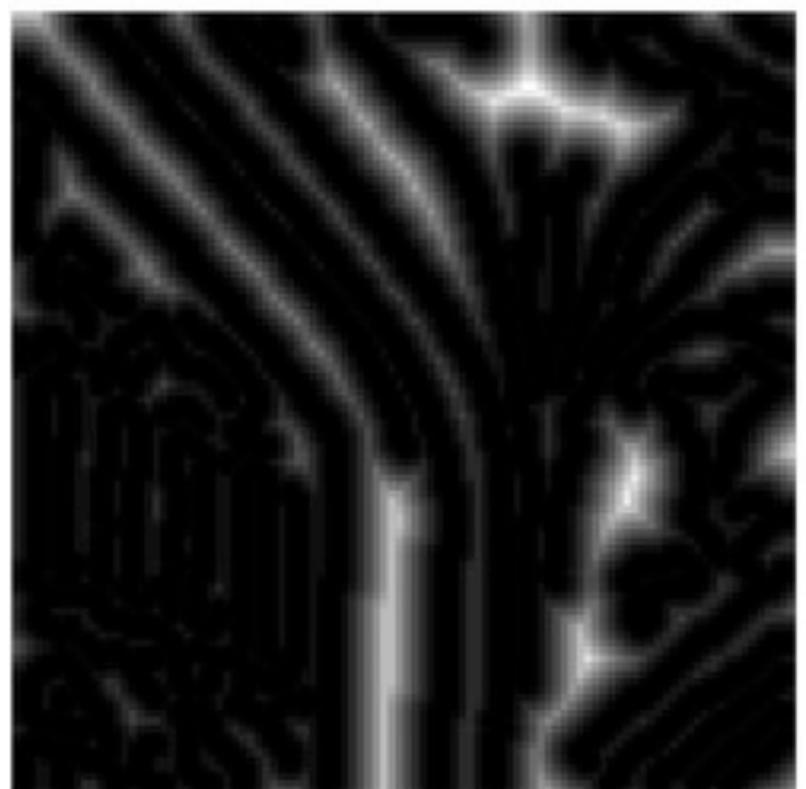
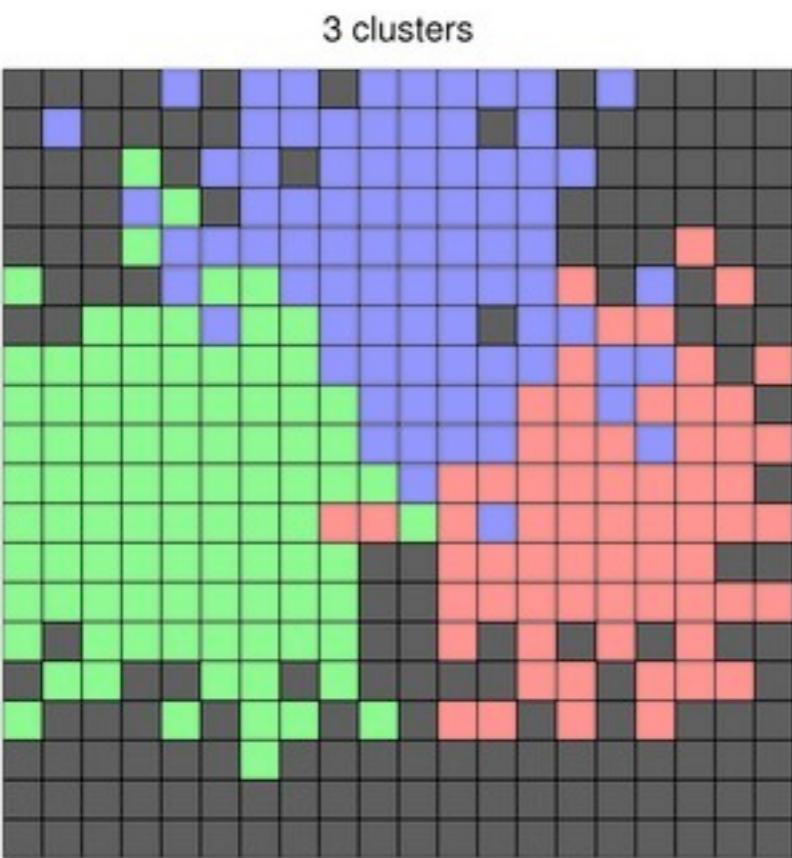
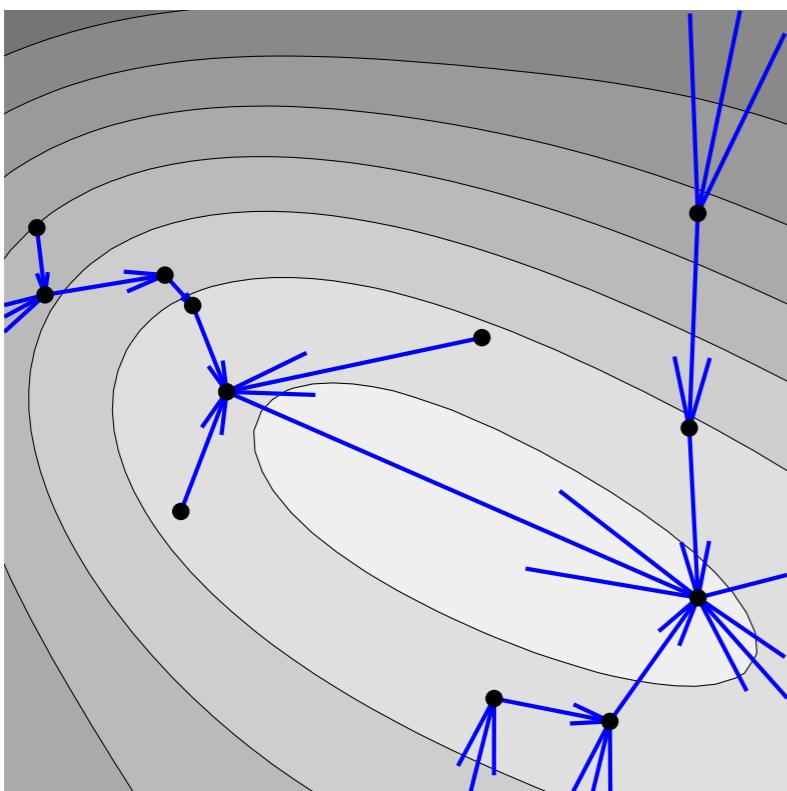


Demo: Speedups

- Feature extraction: 15x
- Feature quantization: 12x
- Classification: 13x
- Overall: 15x



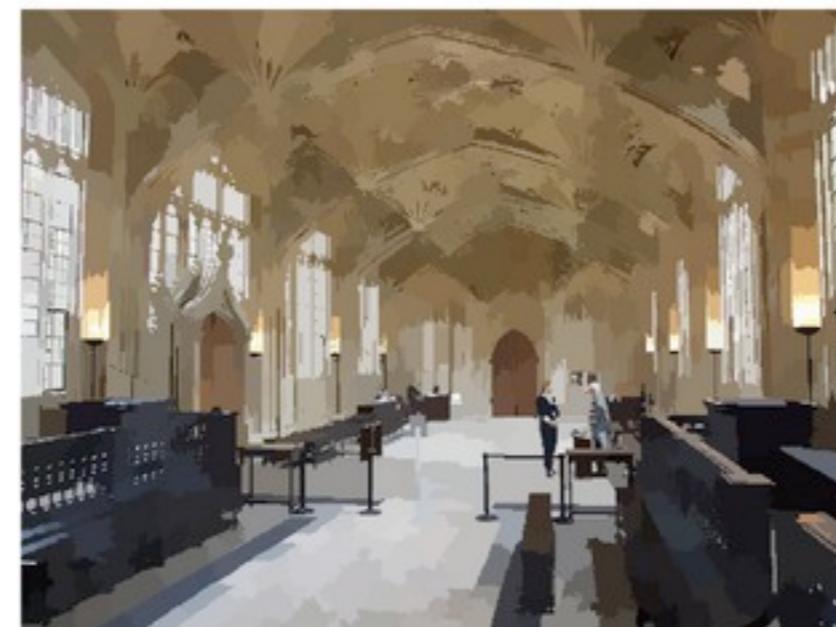
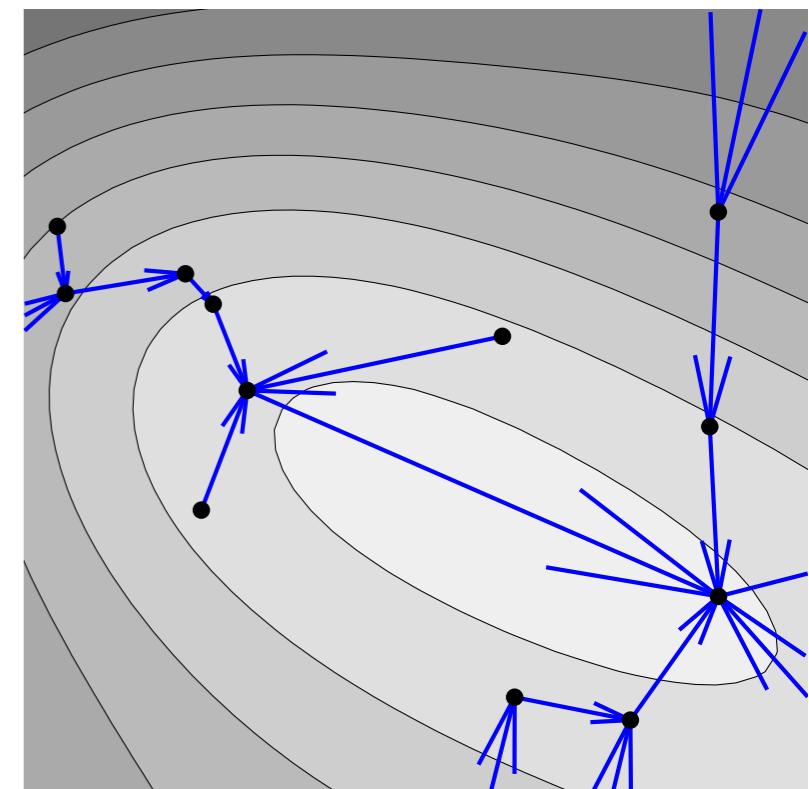
Other features



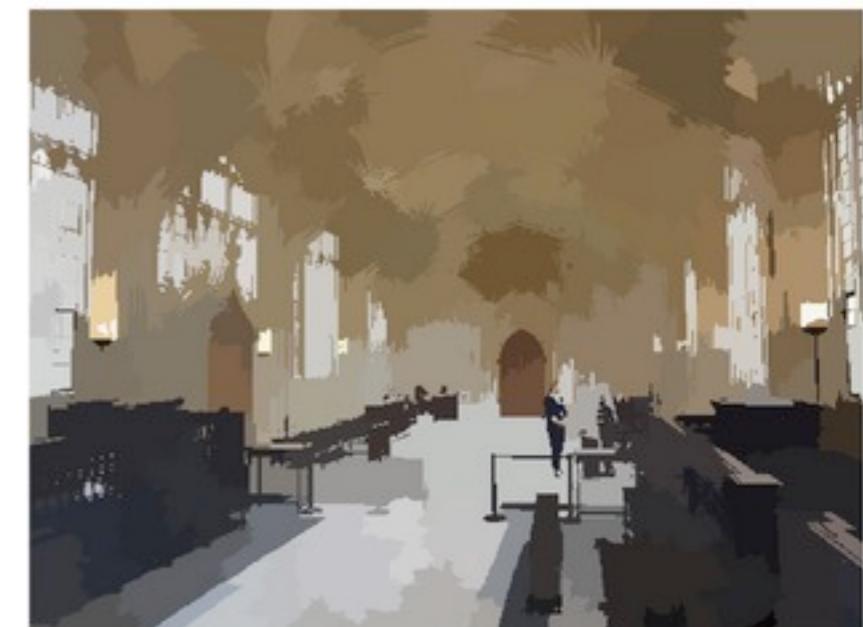
Quick shift

- Mode seeking algorithm (like mean shift)
 - [Comaniciu and Meer 2002]
[Vedaldi et. al 2008]
 - Connect to the nearest point with higher energy.
 - Break links > maxdist to form clusters.

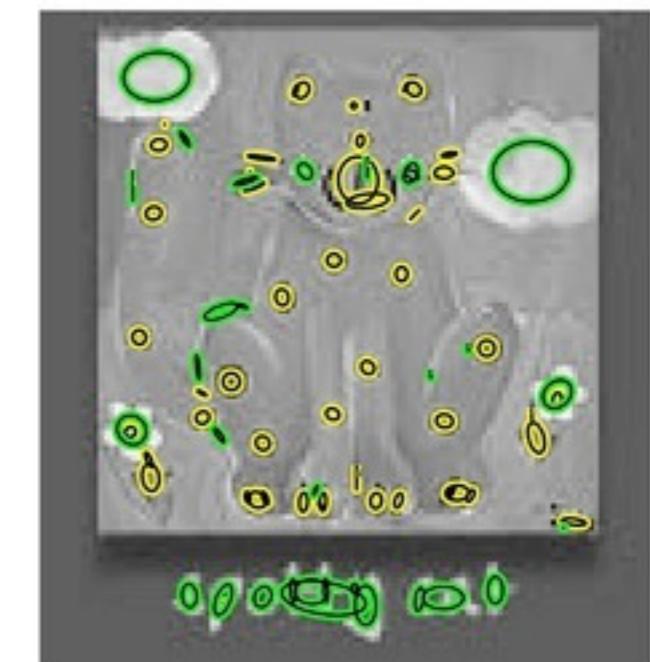
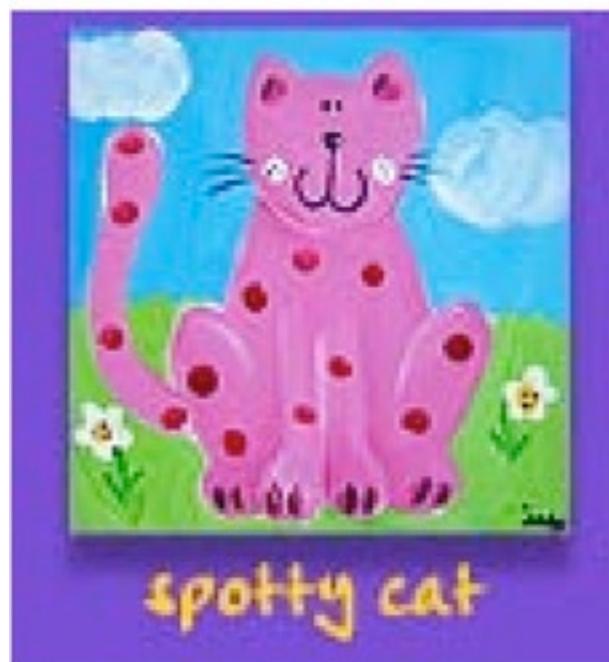
```
ratio = 0.5;
kernelsize = 2;
Iseg = vl_quickseg(I, ratio, kernelsize, maxdist) ;
```



maxdist = 10;



maxdist = 20;



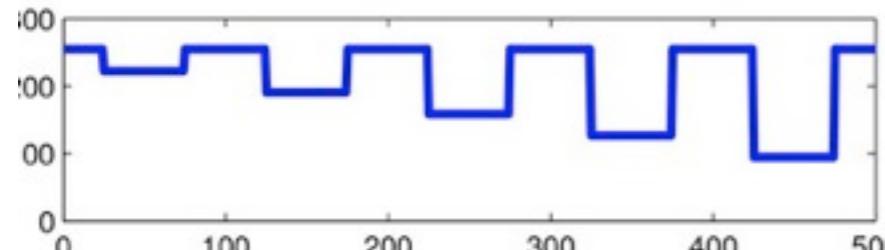
- Extracts MSER regions **regions** and fitted ellipses **frames**:

```
Ig = uint8(rgb2gray(I)) ;  
[regions, frames] = vl_mserv(Ig) ;
```

[Matas et al. 2002]

- Control region stability:

```
[regions, frames] = vl_mserv(Ig, ...  
    'MinDiversity', 0.7, ...  
    'MaxVariation', 0.2, ...  
    'Delta', 10) ;
```



delta = 1

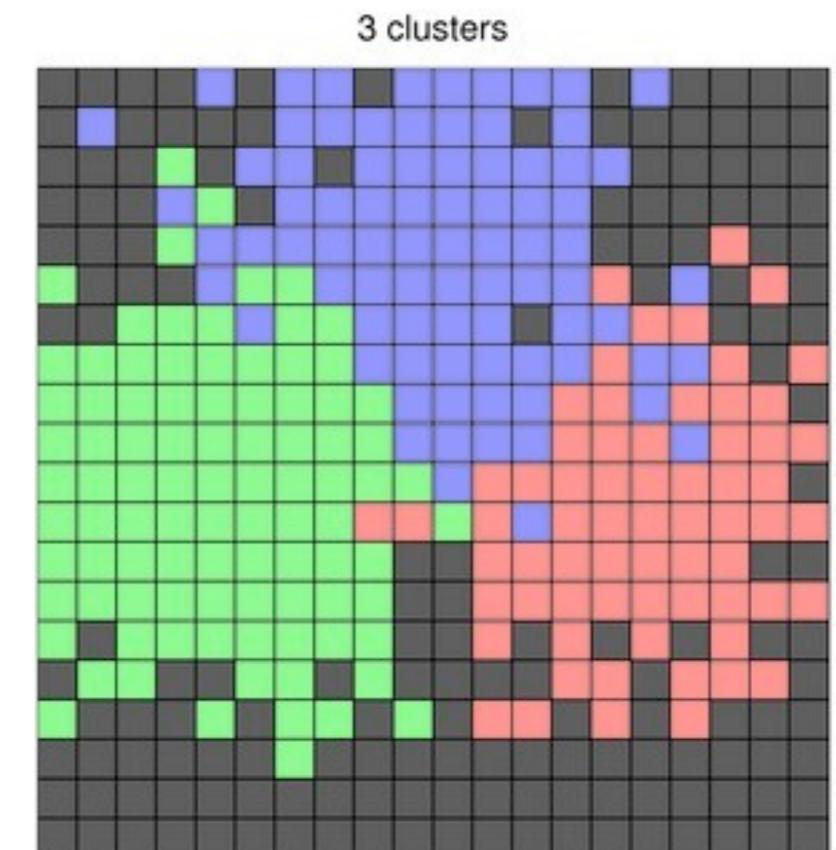
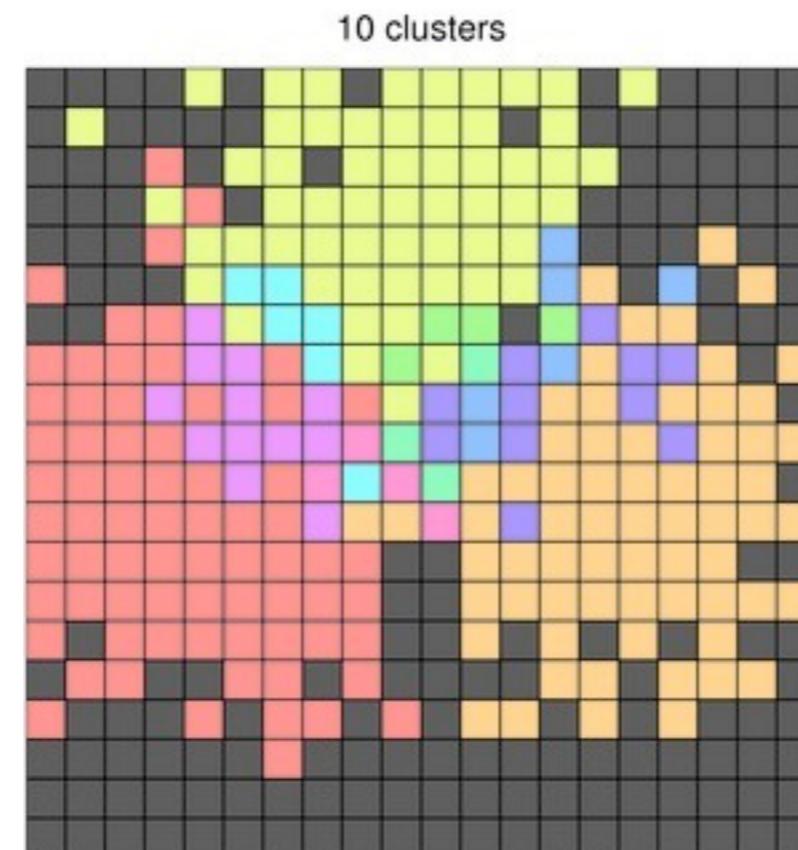
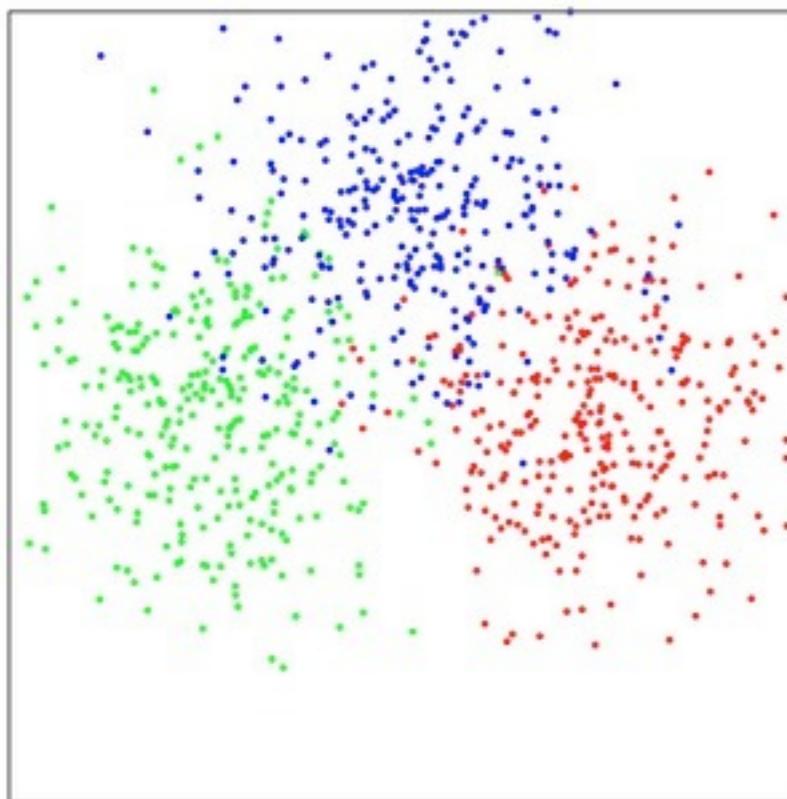


delta = 32



- Merge (visual) words while preserving information
 - given word-class co-occurrence matrix P_{cx}
`[parents, cost] = vl_aib(Pcx);`
 - produces a tree of merged words
 - tree cut
 - = partition of words
 - = simplified dictionary

[Slonim and Tishby 1999]
 [Fulkerson et al. 2008]

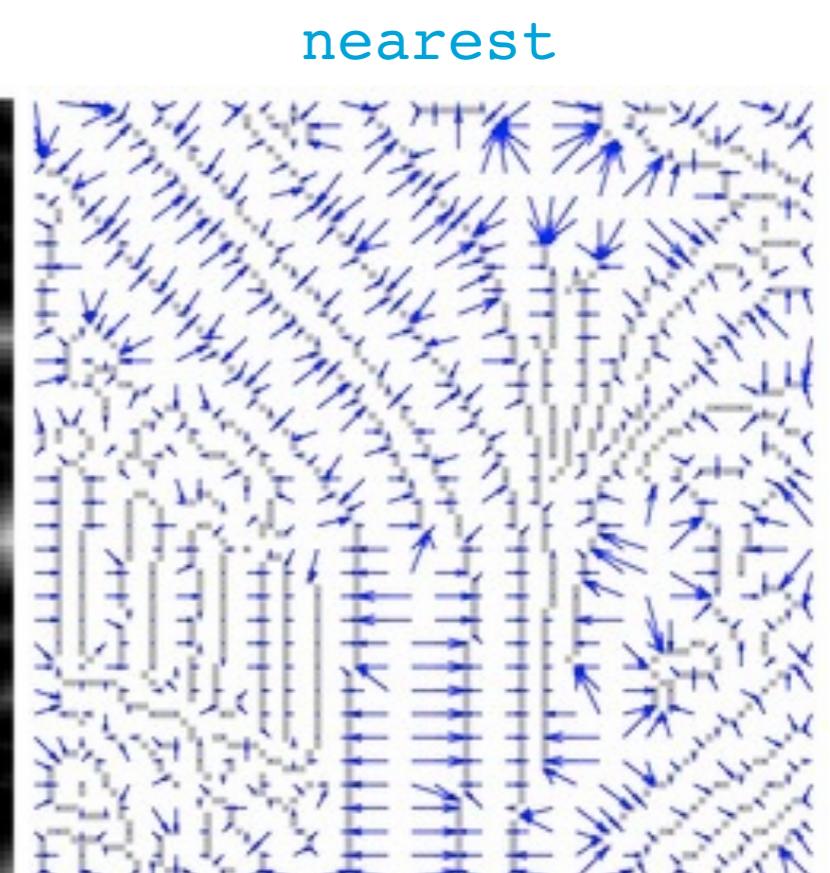
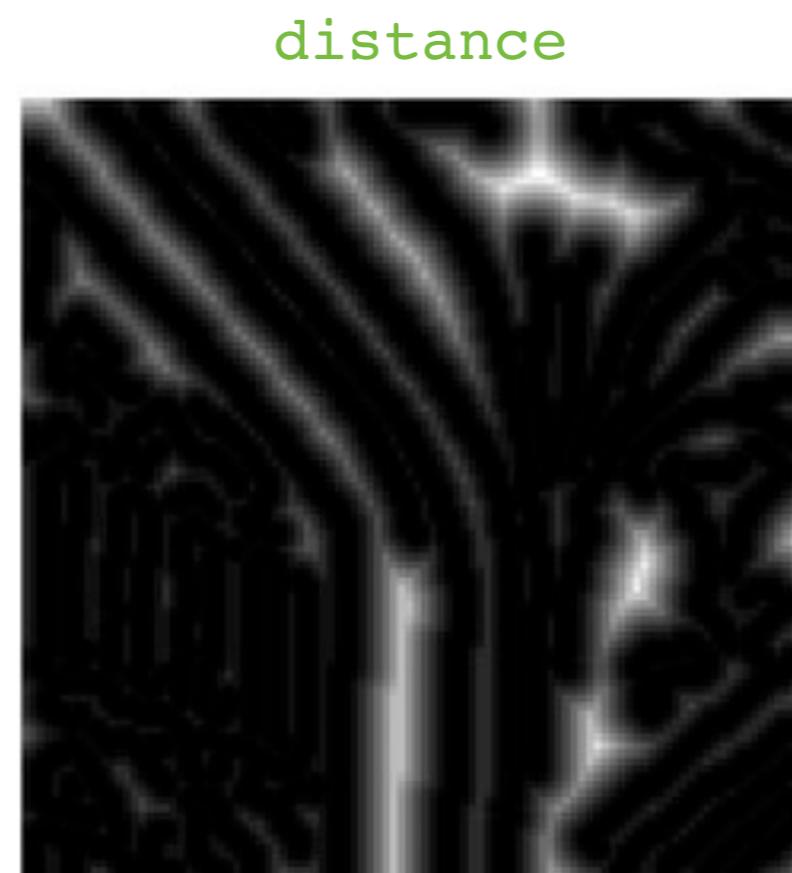
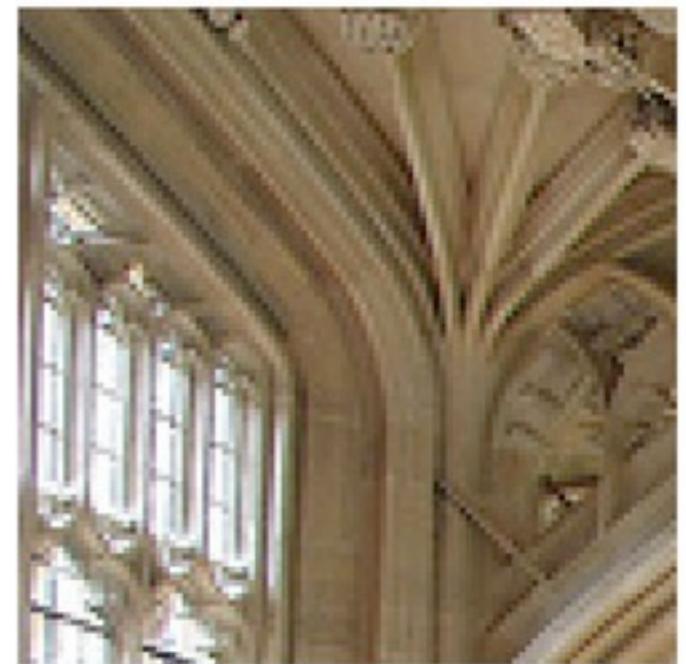


Fast image distance transform

- [Felzenszwalb and Huttenlocher 2004]
- Compute feature responses (e.g. edges):


```
edges = zeros(imsiz) + inf ;
edges(edge(rgb2gray(im), 'canny')) = 0 ;
```
- For each point, find `nearest` edge and its `distance`:


```
[distance, nearest] = vl_imdisttf(single(edges));
```



- Built-in MATLAB help

help `vl_sift`

VL_SIFT Scale-Invariant Feature Transform
`F = VL_SIFT(I)` computes the SIFT frames [1] (keypoints) `F` of the image `I`. `I` is a gray-scale image in single precision. Each column of `F` is a feature frame and has the format `[X;Y;S;TH]`, where `X,Y` is the (fractional) center of the frame, `S` is the scale and `TH` is the orientation (in radians).

`[F,D] = VL_SIFT(I)` computes the SIFT descriptors [1] as well. Each column of `D` is the descriptor of the corresponding frame in `F`. A descriptor is a 128-dimensional vector of class `UINT8`.

`VL_SIFT()` accepts the following options:

Octaves:: [maximum possible]
Set the number of octave of the DoG scale space.

Levels:: [3]
Set the number of levels per octave of the DoG scale space.

FirstOctave:: [0]
Set the index of the first octave of the DoG scale space.

PeakThresh:: [0]
Set the peak selection threshold.

EdgeThresh:: [10]
Set the non-edge selection threshold.

NormThresh:: [-inf]
Set the minimum l2-norm of the descriptors before normalization. Descriptors below the threshold are set to zero.

Magnif:: [3]
Set the descriptor magnification factor. The scale of the keypoint is multiplied by this factor to obtain the width (in pixels) of the spatial bins. For instance, if there are 4 spatial bins along each spatial direction, the ``side'' of the descriptor is approximatively `4 * MAGNIF`.

WindowSize:: [2]
Set the variance of the Gaussian window that determines the descriptor support. It is expressend in units of spatial bins.

- Also available as HTML

http://www.vlfeat.org/doc/mdoc/VL_SIFT.html

[Index](#) [Prev](#) [Next](#)

`F = VL_SIFT(I)` computes the SIFT frames [1] (keypoints) `F` of the image `I`. `I` is a gray-scale image in single precision. Each column of `F` is a feature frame and has the format `[X;Y;S;TH]`, where `X,Y` is the (fractional) center of the frame, `S` is the scale and `TH` is the orientation (in radians).

`[F,D] = VL_SIFT(I)` computes the SIFT descriptors [1] as well. Each column of `D` is the descriptor of the corresponding frame in `F`. A descriptor is a 128-dimensional vector of class `UINT8`.

VL_SIFT() accepts the following options:

Octaves [maximum possible]
Set the number of octave of the DoG scale space.

Levels [3]
Set the number of levels per octave of the DoG scale space.

FirstOctave [0]
Set the index of the first octave of the DoG scale space.

PeakThresh [0]
Set the peak selection threshold.

EdgeThresh [10]
Set the non-edge selection threshold.

NormThresh [-inf]
Set the minimum l2-norm of the descriptors before normalization. Descriptors below the threshold are set to zero.

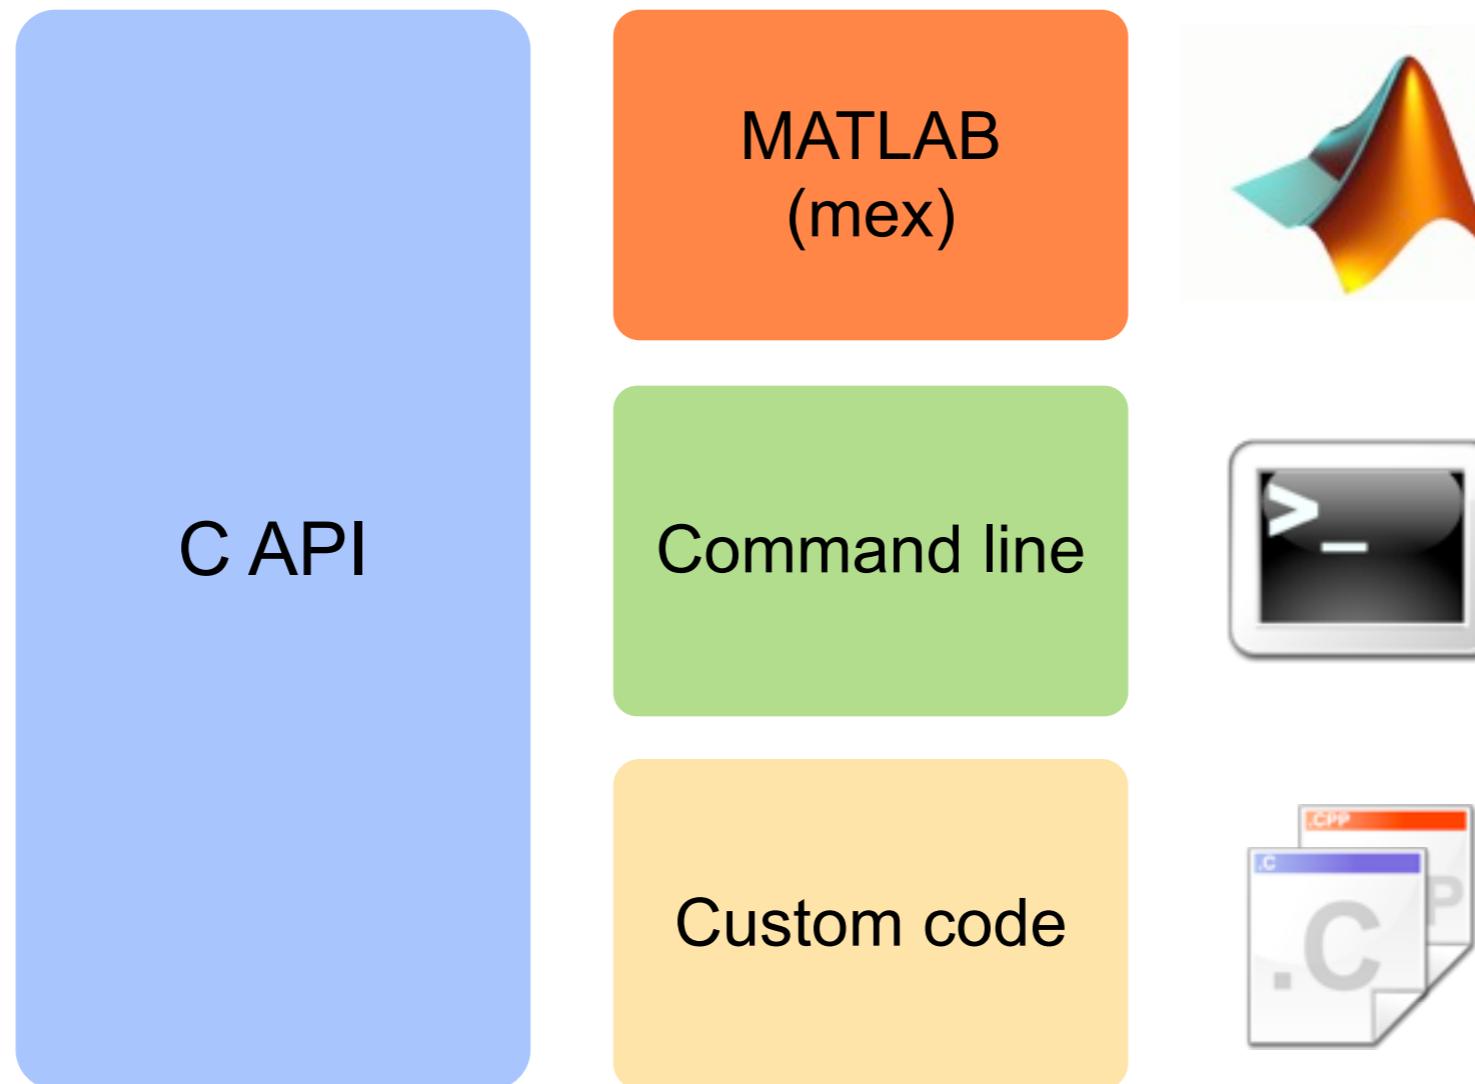
Magnif [3]
Set the descriptor magnification factor. The scale of the keypoint is multiplied by this factor to obtain the width (in pixels) of the spatial bins. For instance, if there are 4 spatial bins along each spatial direction, the ``side'' of the descriptor is approximatively `4 * MAGNIF`.

WindowSize [2]
Set the variance of the Gaussian window that determines the descriptor support. It is expressend in units of spatial bins.

Frames [not specified]
If specified, set the frames to use (bypass the detector). If frames are not passed in order of increasing scale, they are re-ordered.

Orientations

- Components



- No dependencies
- C implementation

VLFeat C API documentation

- Full algorithm details

- embedded in source code
- Doxygen format <http://www.doxygen.org>

number of scale levels per octave	detector	vl_sift_new	features
	SIFT detector	vl_sift_new	can affect the number of extracted keypoints
edge threshold	SIFT detector	vl_sift_set_edge_thresh	decrease to eliminate more keypoints
peak threshold	SIFT detector	vl_sift_set_peak_thresh	increase to eliminate more keypoints

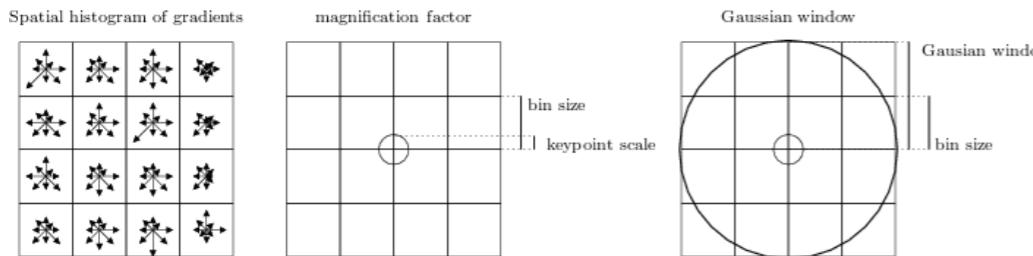
Summary of the parameters influencing the SIFT detector.

SIFT Descriptor

See also:

[Descriptor technical details](#)

A SIFT descriptor is a 3-D spatial histogram of the image gradients in characterizing the appearance of a keypoint. The gradient at each pixel is regarded as a sample of a three-dimensional elementary feature vector, formed by the pixel location and the gradient orientation. Samples are weighed by the gradient norm and accumulated in a 3-D histogram h , which (up to normalization and clamping) forms the SIFT descriptor of the region. An additional Gaussian weighting function is applied to give less importance to gradients farther away from the keypoint center. Orientations are quantized into eight bins and the spatial coordinates into four each, as follows:

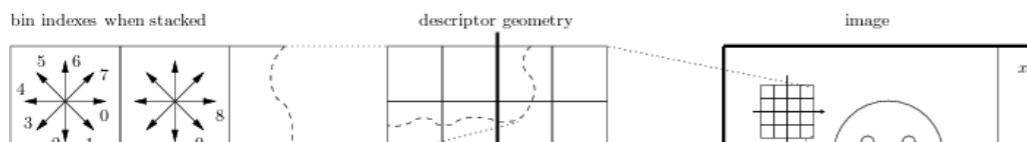


The SIFT descriptor is a spatial histogram of the image gradient.

SIFT descriptors are computed by either calling [vl_sift_calc_keypoint_descriptor](#) or [vl_sift_cal_keypoint_descriptor_raw](#). They accept as input a keypoint frame, which specifies the descriptor center, its size, and its orientation on the image plane. The following parameters influence the descriptor calculation:

- **magnification factor**. The descriptor size is determined by multiplying the keypoint scale by this factor. It is set by [vl_sift_set_magnif](#).
- **Gaussian window size**. The descriptor support is determined by a Gaussian window, which discounts gradient contributions farther away from the descriptor center. The standard deviation of this window is set by [vl_sift_set_window_size](#) and expressed in unit of bins.

VLFeat SIFT descriptor uses the following convention. The y axis points downwards and angles are measured clockwise (to be consistent with the standard image convention). The 3-D histogram (consisting of $8 \times 4 \times 4 = 128$ bins) is stacked as a single 128-dimensional vector, where the fastest varying dimension is the orientation and the slowest the y spatial coordinate. This is illustrated by the following figure.



[vl_kmeans_get_centers](#) to obtain the `numCluster` cluster centers. Use [vl_kmeans_push](#) to quantize new data points.

Initialization algorithms

[kmeans.h](#) supports the following cluster initialization algorithms:

- **Random data points** ([vl_kmeans_init_centers_with_rand_data](#)) initialize the centers from a random selection of the training data.
- **k-means++** ([vl_kmeans_init_centers_plus_plus](#)) initialize the centers from a random selection of the training data while attempting to obtain a good coverage of the dataset. This is the strategy from [1].

Optimization algorithms

[kmeans.h](#) supports the following optimization algorithms:

- **Lloyd** [2] ([VIKMeansLloyd](#)). This is the standard k-means algorithm, alternating the estimation of the point-to-cluster membership and of the cluster centers (means in the Euclidean case). Estimating membership requires computing the distance of each point to all cluster centers, which can be extremely slow.
- **Elkan** [3] ([VIKMeansElkan](#)). This is a variation of [2] that uses the triangular inequality to avoid many distance calculations when assigning points to clusters and is typically much faster than [2]. However, it uses storage proportional to the square of the number of clusters, which makes it unpractical for a very large number of clusters.

Technical details

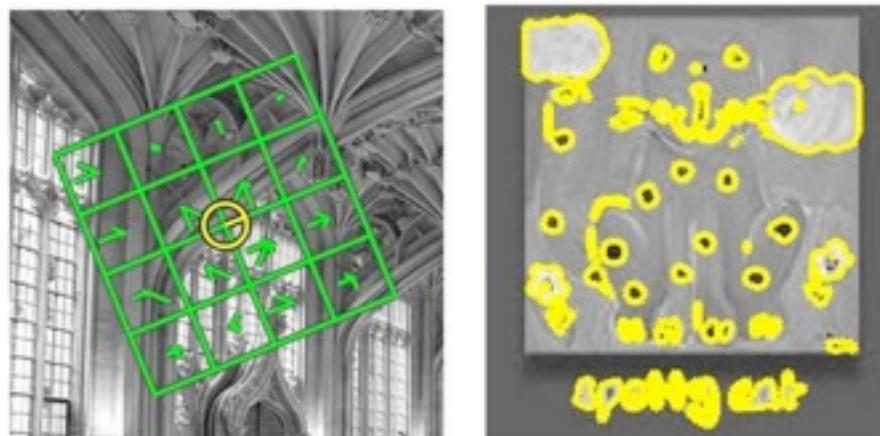
Given data points $x_1, \dots, x_n \in R^d$, k-means searches for k vectors $c_1, \dots, c_n \in R^d$ (cluster centers) and a function $\pi : \{1, \dots, n\} \rightarrow \{1, \dots, k\}$ (cluster memberships) that minimize the objective:

$$E(c_1, \dots, c_n, \pi) = \sum_{i=1}^n d^2(x_i, c_{\pi(i)})$$

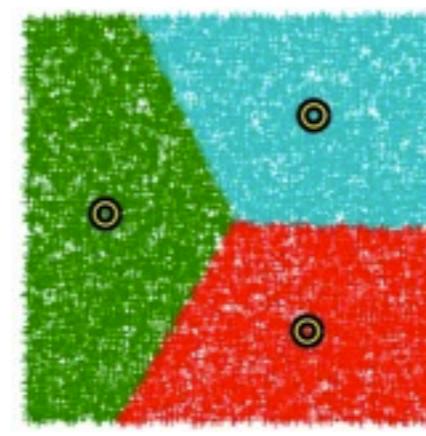
A simple procedure due to Lloyd [2] to locally optimize this objective alternates estimating the cluster centers and the membership function. Specifically, given the membership function π , the objective can be minimized independently for each c_k by minimizing

$$\sum_{i:\pi(i)=k} d^2(x_i, c_k)$$

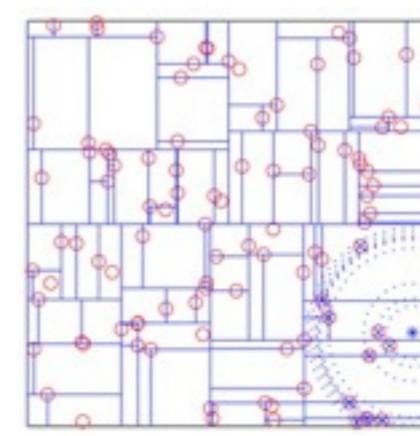
For the Euclidean distance, the minimizer is simply the mean of the points assigned to that cluster. For other distances, the minimizer is a generalized average. For instance, for the l^1 distance, this is the median. Assuming that computing the average is linear in the number of points and the data dimension, this step requires $O(nd)$ operations.



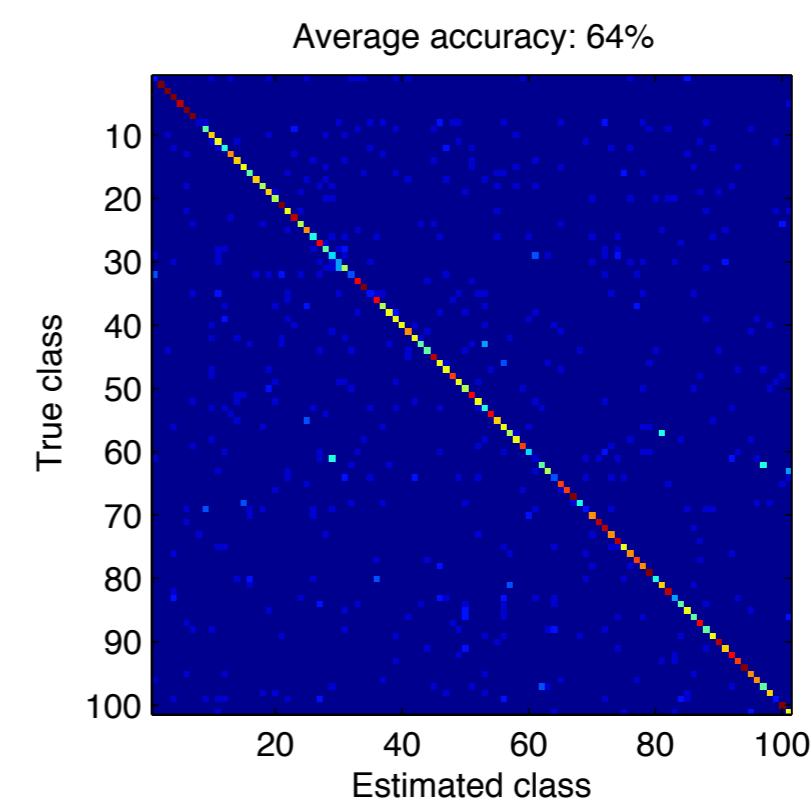
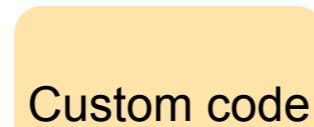
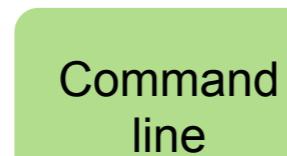
features



clustering



matching



References

- F. R. Bach and M. I. Jordan. Predictive low-rank decomposition for kernel methods. In *ICML*, 2005.
- H. Bay, A. Ess, T. Tuytelaars, and L. van Gool. Speeded-up robust features (SURF). *Computer Vision and Image Understanding*, 2008.
- L. Bo and C. Sminchisescu. Efficient match kernels between sets of features for visual recognition. In *Proc. NIPS*, 2009.
- A. Bosch, A. Zisserman, and X. Muñoz. Scene classification via pLSA. In *Proc. ECCV*, 2006.
- A. Bosch, A. Zisserman, and X. Muñoz. Image classification using random forests and ferns. In *Proc. ICCV*, 2007.
- C.-C. Chang and C.-J. Lin. *LIBSVM: a library for support vector machines*, 2001. Software available at <http://www.csie.ntu.edu.tw/~cjlin/libsvm>.
- D. Comaniciu and P. Meer. Mean shift: A robust approach toward feature space analysis. *PAMI*, 24(5), 2002.
- G. Csurka, C. R. Dance, L. Dan, J. Willamowski, and C. Bray. Visual categorization with bags of keypoints. In *Proc. ECCV Workshop on Stat. Learn. in Comp. Vision*, 2004.
- C. Elkan. Using the triangle inequality to accelerate k -means. In *Proc. ICML*, 2003.
- R.-E. Fan, K.-W. Chang, C.-J. Hsieh, X.-R. Wang, and C.-J. Lin. LIBLINEAR: A library for large linear classification. *Journal of Machine Learning Research*, 9, 2008.
- L. Fei-Fei, R. Fergus, and P. Perona. A Bayesian approach to unsupervised one-shot learning of object categories. In *Proc. ICCV*, 2003.
- P. F. Felzenszwalb and D. P. Huttenlocher. Distance transforms of sampled functions. Technical report, Cornell University, 2004.
- B. J. Frey and D. Dueck. Clustering by passing messages between data points. *Science*, 315, 2007.
- J. H. Friedman, J. L. Bentley, and R. A. Finkel. An algorithm for finding best matches in logarithmic expected time. *ACM Transactions on Mathematical Software*, 1977.
- B. Fulkerson, A. Vedaldi, and S. Soatto. Localizing objects with smart dictionaries. In *Proc. ECCV*, 2008.
- T. Hastie. Support vector machines, kernel logistic regression, and boosting. Lecture Slides, 2003.
- T. Joachims. Making large-scale support vector machine learning practical. In *Advances in kernel methods: support vector learning*, pages 169–184. MIT Press, Cambridge, MA, USA, 1999.
- T. Joachims. Training linear SVMs in linear time. In *Proc. KDD*, 2006.
- S. Lazebnik, C. Schmid, and J. Ponce. Beyond bag of features: Spatial pyramid matching for recognizing natural scene categories. In *Proc. CVPR*, 2006.
- B. Leibe, K. Mikolajczyk, and B. Schiele. Efficient clustering and matching for object class recognition. In *Proc. BMVC*, 2006.
- D. G. Lowe. Distinctive image features from scale-invariant keypoints. *IJCV*, 2(60):91–110, 2004.
- S. Maji and A. C. Berg. Max-margin additive classifiers for detection. In *Proc. ICCV*, 2009.
- D. Nistér and H. Stewénius. Scalable recognition with a vocabulary tree. In *Proc. CVPR*, 2006.
- F. Perronnin, J. Sánchez, and Y. Liu. Large-scale image categorization with explicit data embedding. In *Proc. CVPR*, 2010.
- A. Rahimi and B. Recht. Random features for large-scale kernel machines. In *Proc. NIPS*, 2007.
- B. Schölkopf and A. J. Smola. *Learning with Kernels*. MIT Press, 2002.
- S. Shalev-Shwartz, Y. Singer, N. Srebro, and A. Cotter. Pegasos: Primal Estimated sub-GrAdient SOLver for SVM. *MBP*, 2010.
- J. Shawe-Taylor and N. Cristianini. *Support Vector Machines and other kernel-based learning methods*. Cambridge University Press, 2000.
- S. N. Sinha, J.-M. Frahm, M. Pollefeys, and Y. Gen. Gpu-based video feature tracking and matching. In *Workshop on Edge Computing Using New Commodity Architectures*, 2006.
- J. Sivic and A. Zisserman. Video Google: A text retrieval approach to object matching in videos. In *Proc. ICCV*, 2003.
- N. Slonim and N. Tishby. Agglomerative information bottleneck. In *Proc. NIPS*, 1999.
- E. Tola, V. Lepetit, and P. Fua. DAISY: An efficient dense descriptor applied to wide-baseline stereo. *PAMI*, 2010.
- A. Vedaldi and S. Soatto. Quick shift and kernel methods for mode seeking. In *Proc. ECCV*, 2008.
- C. K. I. Williams and M. Seeger. Using the Nyström method to speed up kernel machines. In *Proc. NIPS*, 2001.